



# Hidden Markov Model

Il-Chul Moon

Department of Industrial and Systems Engineering



KAIST

icmoon@kaist.ac.kr

# HIDDEN MARKOV MODEL

# Times Series Data for GMM

- Imagine the following case
  - Data points on the plane
  - Have a temporal trace of data points
  - Now, any broken assumption in the analysis?
- Any real world applications
  - Many, many, many...
  - Stock market analysis, text mining...



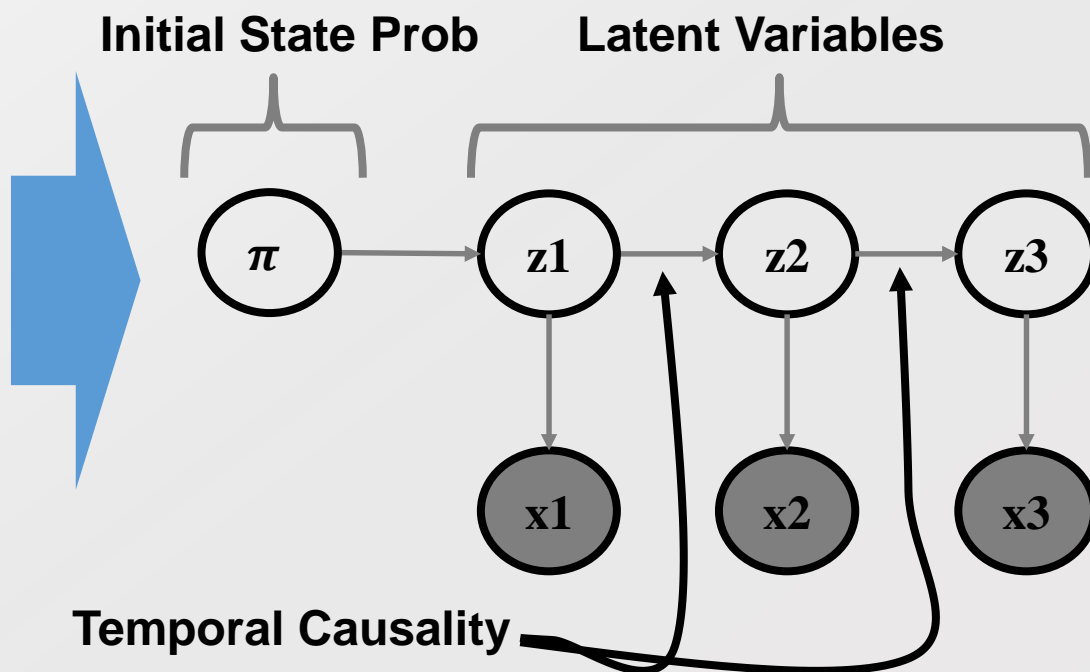


DJIA

If they make the ballot in November, an array of proposals will be among the first in the nation to ask a state's voters to sharply



Related...

# What to Model and How to Model

- Previously, all data points are independent trials
  - Now, they are not any further
- Temporal relation: causality from time  $t$  to time  $t+1$
- Overall trend: latent state variables



- Observation,  $x$ 
  - Can be either discrete or continuous
    - Just a difference in probability distributions
    - Will only handle discrete case in this course
  - $x_1 \dots x_T$ : Observation from time 1 to time  $T$
  - $x_i \in \{c_1, \dots, c_m\}$ :  $m$  types of observation values
- Latent state,  $z$ 
  - Vector variable with  $K$  elements
    - Let's say that there are  $K$  types of state values corresponding to each element
  - Can be either discrete or continuous
    - If continuous  $\rightarrow$  Kalman filter, and this is a continuous version of HMM
    - Out of scope of this course
- Initial State probabilities
  - $P(z_1) \sim \text{Mult}(\pi_1, \dots, \pi_k)$
- Transition probabilities
  - $P(z_t | z_{t-1}^i = 1) \sim \text{Mult}(a_{i,1}, \dots, a_{i,k})$
  - Or,  $P(z_t^j = 1 | z_{t-1}^i = 1) = a_{i,j}$
- Emission probabilities
  - $P(x_t | z_t^i = 1) \sim \text{Mult}(b_{i,1}, \dots, b_{i,m}) \sim f(x_t | \theta_i)$
  - Or,  $P(x_t^j = 1 | z_t^i = 1) = b_{i,j}$
- **A stochastic generative model**



# Main Questions on HMM

- Given the topology of the Bayesian network, HMM, or M
- Evaluation question
  - Given  $\pi, a, b, X$
  - Find  $P(X|M, \pi, a, b)$
  - How much is X likely to be observed in the trained model?
- Decoding question
  - Given  $\pi, a, b, X$
  - Find  $\operatorname{argmax}_z P(Z|X, M, \pi, a, b)$
  - What would be the most probable sequences of latent states?
- Learning question
  - Given X
  - Find  $\operatorname{argmax}_{\pi, a, b} P(X|M, \pi, a, b)$
  - What would be the underlying parameters of the HMM given the observations?
- Decoding questions and learning questions are very similar to
  - Supervised and unsupervised learning
- Anyhow, we often need to find  $\pi, a, b$  prior to the supervised learning with X

Initial State probabilities

$$P(z_1) \sim \text{Mult}(\pi_1, \dots, \pi_k)$$

Transition probabilities

$$P(z_t | z_{t-1}^i = 1) \sim \text{Mult}(a_{i,1}, \dots, a_{i,k})$$

$$\text{Or, } P(z_t^j = 1 | z_{t-1}^i = 1) = a_{i,j}$$

Emission probabilities

$$P(x_t | z_t^i = 1) \sim \text{Mult}(b_{i,1}, \dots, b_{i,m}) \sim f(x_t | \theta_i)$$

$$\text{Or, } P(x_t^j = 1 | z_t^i = 1) = b_{i,j}$$

# Obtaining $\pi$ , $a$ , $b$ given $X$ and $M$

Initial State probabilities

$$P(z_1) \sim \text{Mult}(\pi_1, \dots, \pi_k)$$

Transition probabilities

$$P(z_t | z_{t-1}^i = 1) \sim \text{Mult}(a_{i,1}, \dots, a_{i,k})$$

$$\text{Or, } P(z_t^j = 1 | z_{t-1}^i = 1) = a_{i,j}$$

Emission probabilities

$$P(x_t | z_t^i = 1) \sim \text{Mult}(b_{i,1}, \dots, b_{i,m}) \sim f(x_t | \theta_i)$$

$$\text{Or, } P(x_t^j = 1 | z_t^i = 1) = b_{i,j}$$



## $M_i$ observations for $i$ -th sequence

- Finding  $\pi$ ,  $a$ ,  $b$  from the data in the supervised learning approach requires  $X$  as well as  $Z$
- Example scenario
  - Loaded dice and fair dice
  - Two dices yield different probability distributions from one to six
  - Dealer changes the dice as he wishes
- Probability estimation
  - Use MLE, MAP and counting...
  - Find out
    - Dealer starts with a certain dice type:  $P(z_1^L = 1) = 1/2$
    - Dealer switches the dice:  $P(z_t^L = 1 | z_{t-1}^L = 1) = 0.7, P(z_t^L = 1 | z_{t-1}^F = 1) = 0.5$
    - Loaded dice:  $P(X = 1) = P(X = 2) = P(X = 3) = P(X = 4) = P(X = 5) = 1/10, P(X = 6) = 1/2$
    - Fair dice:  $P(X = 1) = P(X = 2) = P(X = 3) = P(X = 4) = P(X = 5) = P(X = 6) = 1/6$
- What if the  $X$  is continuous? Use a known distribution, and estimate its parameters

# Joint Probability

- Let's assume that we have a training dataset with  $X$  and  $Z$

**Initial State probabilities**

$$P(z_1) \sim \text{Mult}(\pi_1, \dots, \pi_k)$$

**Transition probabilities**

$$P(z_t | z_{t-1} = 1) \sim \text{Mult}(a_{i,1}, \dots, a_{i,k})$$

$$\text{Or, } P(z_t^j = 1 | z_{t-1}^i = 1) = a_{i,j}$$

**Emission probabilities**

$$P(x_t | z_t = 1) \sim \text{Mult}(b_{i,1}, \dots, b_{i,m}) \sim f(x_t | \theta_i)$$


$$\text{Or, } P(x_t^j = 1 | z_t^i = 1) = b_{i,j}$$

N  
seq.

```
LLLLLLLLLLLLLLLL
12342531643242
.....
FFFLFLFLFLFLFL
3526152436152436152
```

$M_i$  observations for  $i$ -th sequence

- Can we compute the joint probability,  $P(X, Z)$ 
  - Yes. Easily by the virtue of the network structure
- Anyway, a Bayesian network, so...
  - Factorize
  - $P(X, Z) = P(x_1, \dots, x_t, z_1, \dots, z_t)$
  - $= P(z_1)P(x_1|z_1)P(z_2|z_1)P(x_2|z_2)P(z_3|z_2)P(x_3|z_3)$ 
    - Nothing but a combination of initial, transition, and emission probabilities
  - $= \pi_{idx(z_1=1)} b_{idx(x_1=1), idx(z_1=1)} a_{idx(z_1=1), idx(z_2=1)} \dots$



- Assume that we have 166 as  $X$ 
  - Let's check  $Z=LLL$  and  $FFF$
  - $P(166, LLL) = \frac{1}{2} \times \frac{1}{10} \times \frac{7}{10} \times \frac{1}{2} \times \frac{7}{10} \times \frac{1}{2} = 0.0061$
  - $P(166, FFF) = \frac{1}{2} \times \frac{1}{6} \times \frac{1}{2} \times \frac{1}{6} \times \frac{1}{2} \times \frac{1}{6} = 5.7870e - 04$
  - What about  $FLL, FFL, FLF, \dots$ ? Exponential combination to check



# Marginal Probability

- Eventually, we only want to use  $X$  and marginalize  $Z$

- Just like GMM,  

$$P(X|\theta) = \sum_Z P(X, Z|\theta)$$

- In HMM,  

$$P(X|\pi, a, b) = \sum_Z P(X, Z|\pi, a, b)$$

- $P(X) = \sum_Z P(X, Z) = \sum_{z_1} \dots \sum_{z_t} P(x_1, \dots, x_t, z_1, \dots, z_t)$

- $= \sum_{z_1} \dots \sum_{z_t} \pi_{z_1} \prod_{t=2}^T a_{z_{t-1}, z_t} \prod_{t=1}^T b_{z_t, x_t}$

- Many summations yield an exponential number of combinations

- Need to avoid a repetitive computing

- Compute only necessary terms for a single time

- Let's work on the formula

- $P(A, B, C) = P(A)P(B|A)P(C|A, B)$

- $P(x_1, \dots, x_t, z_t^k = 1) = \sum_{z_{t-1}} P(x_1, \dots, x_{t-1}, x_t, z_{t-1}, z_t^k = 1)$

- $= \sum_{z_{t-1}} P(x_1, \dots, x_{t-1}, z_{t-1}) P(z_t^k = 1 | x_1, \dots, x_{t-1}, z_{t-1}) P(x_t | z_t^k = 1, x_1, \dots, x_{t-1}, z_{t-1})$

- By the virtue of the structure

- $= \sum_{z_{t-1}} P(x_1, \dots, x_{t-1}, z_{t-1}) P(z_t^k = 1 | z_{t-1}) P(x_t | z_t^k = 1)$

- $= P(x_t | z_t^k = 1) \sum_{z_{t-1}} P(x_1, \dots, x_{t-1}, z_{t-1}) P(z_t^k = 1 | z_{t-1})$

- $= b_{z_t^k, x_t} \sum_{z_{t-1}} P(x_1, \dots, x_{t-1}, z_{t-1}) a_{z_{t-1}, z_t^k}$

- Now, we see a repeating structure of terms

- $P(x_1, \dots, x_t, z_t^k = 1) = \alpha_t^k = b_{k, x_t} \sum_i \alpha_{t-1}^i a_{i, k}$

**Initial State probabilities**

$$P(z_1) \sim \text{Mult}(\pi_1, \dots, \pi_k)$$

**Transition probabilities**

$$P(z_t | z_{t-1}^i = 1) \sim \text{Mult}(a_{i,1}, \dots, a_{i,k})$$

$$\text{Or, } P(z_t^j = 1 | z_{t-1}^i = 1) = a_{i,j}$$

**Emission probabilities**


$$P(x_t | z_t^i = 1) \sim \text{Mult}(b_{i,1}, \dots, b_{i,m}) \sim f(x_t | \theta_i)$$

$$\text{Or, } P(x_t^j = 1 | z_t^i = 1) = b_{i,j}$$


N  
seq.

```
LLLLLLLLLLLLLLLL
12342531643242
.....
FFFFLLFLFLFLLLFL
3526152436152436152
```

$M_i$  observations for  $i$ -th sequence



- Dynamic programming:
  - A general algorithm design technique for solving problems defined by or formulated as **recurrences with overlapping sub-instances**
  - In this context, Programming == Planning
- Main storyline
  - Setting up a recurrence
    - Relating a solution of a larger instance to solutions of some smaller instances
    - Solve small instances once
    - Record solutions in a table
    - Extract a solution of a larger instance from the table



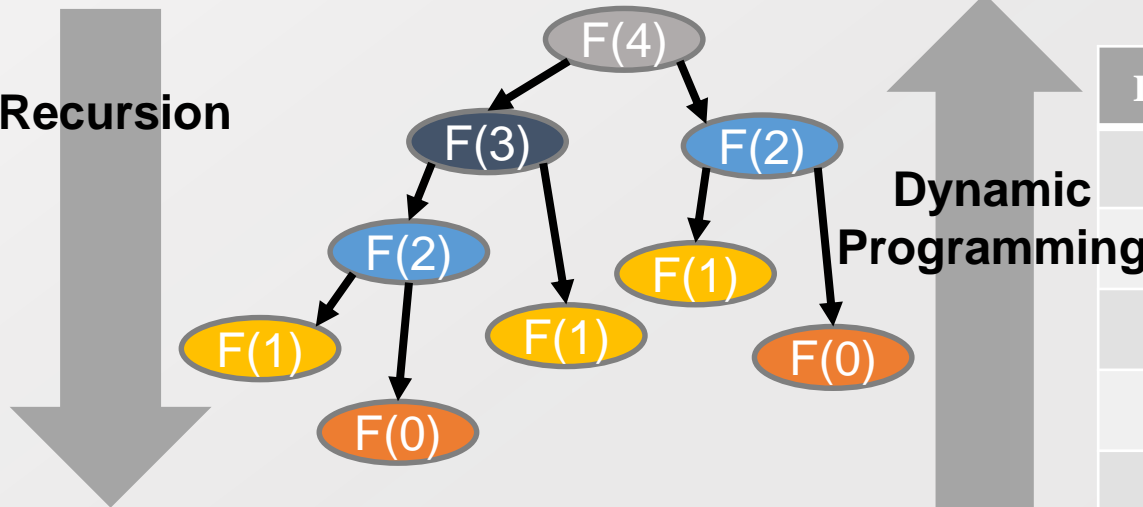
Instance	Solution
F(0)	0
F(1)	1
F(2)	1
F(3)	2
F(4)	?

- Key technique of dynamic programming
  - Simply put
    - Storing the results of previous function calls to reuse the results again in the future
  - More philosophical sense
    - Bottom-up approach for problem-solving
      - Recursion: Top-down of divide and conquer
      - Dynamic programming: Bottom-up of storing and building

## Stackframe

n	2
R.A	
n	3
R.A	
n	4
R.A	

Recursion




Dynamic  
Programming

## Memoization

Instance	Solution
F(0)	0
F(1)	1
F(2)	1
F(3)	2
F(4)	3

- Need to know  $\alpha_t^k$ 
  - Time X States
  - When we know  $\alpha_t^k$  with X, then we know the value of  $P(X)$ 
    - Answering the evaluation question without Z
- ForwardAlgorithm
  - Initialize
    - $\alpha_1^k = b_{k,x_1} \pi_k$
  - Iterate until time T
    - $\alpha_t^k = b_{k,x_t} \sum_i \alpha_{t-1}^i a_{i,k}$
  - Return  $\sum_i \alpha_T^i$
- Proof of correctness
  - $\sum_i \alpha_T^i = \sum_i P(x_1, \dots, x_T, z_T^i = 1) = P(x_1, \dots, x_T)$
- Where to use the memoization table?
  - $\alpha_t^k$
- Limitation of the forward probability
  - Only takes the input sequence of X before time  $t$
  - $P(x_1, \dots, x_t, z_t^k = 1) = \alpha_t^k$  and  $t \neq T$
  - Need to see a probability distribution of a latent variable at time  $t$  given the whole X
  - Recall the Bayes ball algorithm



- We need  $P(z_t^k = 1|X)$  instead of  $P(x_1, \dots, x_t, z_t^k = 1)$

- Let's derive from the joint probability

- $P(z_t^k = 1, X) = P(x_1, \dots, x_t, z_t^k = 1, x_{t+1}, \dots, x_T)$
- $= P(x_1, \dots, x_t, z_t^k = 1)P(x_{t+1}, \dots, x_T | x_1, \dots, x_t, z_t^k = 1)$

- By the virtue of the structure

- $= P(x_1, \dots, x_t, z_t^k = 1)P(x_{t+1}, \dots, x_T | z_t^k = 1)$

- We already handled  $P(x_1, \dots, x_t, z_t^k = 1)$ 
  - $P(x_1, \dots, x_t, z_t^k = 1) = \alpha_t^k$
- So, we need to compute  $P(x_{t+1}, \dots, x_T | z_t^k = 1)$ 
  - $P(x_{t+1}, \dots, x_T | z_t^k = 1) = \beta_t^k$

- $P(x_{t+1}, \dots, x_T | z_t^k = 1)$

- $= \sum_{z_{t+1}} P(z_{t+1}, x_{t+1}, \dots, x_T | z_t^k = 1)$


- $= \sum_i P(z_{t+1}^i = 1 | z_t^k = 1) P(x_{t+1} | z_{t+1}^i = 1, z_t^k = 1) P(x_{t+2}, \dots, x_T | x_{t+1}, z_{t+1}^i = 1, z_t^k = 1)$

- $= \sum_i P(z_{t+1}^i = 1 | z_t^k = 1) P(x_{t+1} | z_{t+1}^i = 1) P(x_{t+2}, \dots, x_T | z_{t+1}^i = 1)$

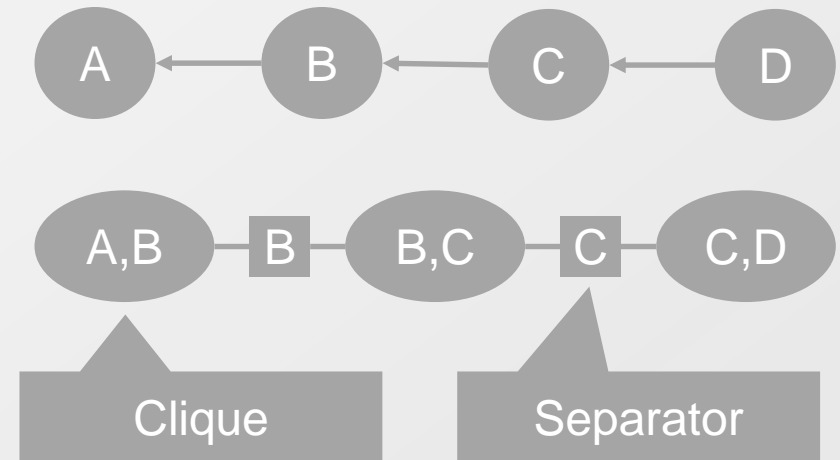
- $= \sum_i a_{k,i} b_{i,x_t} \beta_{t+1}^i$

- Again, recursive structure. How to calculate this efficiently?

- $P(z_t^k = 1, X) = \alpha_t^k \beta_t^k = (b_{k,x_t} \sum_i \alpha_{t-1}^i a_{i,k}) \times (\sum_i a_{k,i} b_{i,x_t} \beta_{t+1}^i)$



- $P(A, B, C, D) = P(A|B)P(B|C)P(C|D)P(D)$
- Let's define a potential function
  - Potential function:  
a function which is not a probability function yet, but once normalized it can be a probability distribution function
  - Potential function on nodes
    - $\psi(a, b), \psi(b, c), \psi(c, d)$
  - Potential function on links
    - $\phi(b), \phi(c)$
- How to setup the function?




- $P(A, B, C, D) = P(U) = \frac{\prod_N \psi(N)}{\prod_L \phi(L)} = \frac{\psi(a,b)\psi(b,c)\psi(c,d)}{\phi(b)\phi(c)}$ 
  - $\psi(a, b) = P(A|B), \psi(b, c) = P(B|C), \psi(c, d) = P(C|D)P(D)$
  - $\phi(b) = 1, \phi(c) = 1$
- $P(A, B, C, D) = P(U) = \frac{\prod_N \psi(N)}{\prod_L \phi(L)} = \frac{\psi^*(a,b)\psi^*(b,c)\psi^*(c,d)}{\phi^*(b)\phi^*(c)}$ 
  - $\psi^*(a, b) = P(A, B), \psi^*(b, c) = P(B, C), \psi^*(c, d) = P(C, D)$
  - $\phi^*(b) = P(B), \phi^*(c) = P(C)$

Marginalization is also applicable:

$$\psi(w) = \sum_{v-w} \psi(v)$$

Constructing a potential of a subset (w) of all variables (v)

- Only applicable to the tree structure of clique graph
- Let's assume
  - $P(B) = \sum_A \psi(A, B)$
  - $P(B) = \sum_C \psi(B, C)$
  - $P(B) = \phi(B)$
  - How to find out the  $\psi$ s and the  $\phi$ s?
    - When the  $\psi$ s change by the observations:  $P(A, B) \rightarrow P(A=1, B)$
    - A single  $\psi$  change can result in the change of multiple  $\psi$ s
    - The effect of the observation propagates through the clique graph
    - Belief propagation!
- How to propagate the belief?
  - Absorption (update) rule
  - Assume  $\psi^*(A, B), \psi(B, C)$ , and  $\phi(B)$
  - Define the update rule for separators
    - $\phi^*(B) = \sum_A \psi^*(A, B)$
  - Define the update rule for cliques
    - $\psi^*(B, C) = \psi(B, C) \frac{\phi^*(B)}{\phi(B)}$



Why does this work?

$$\begin{aligned} \sum_C \psi^*(B, C) &= \sum_C \psi(B, C) \frac{\phi^*(B)}{\phi(B)} \\ &= \frac{\phi^*(B)}{\phi(B)} \sum_C \psi(B, C) = \frac{\phi^*(B)}{\phi(B)} \phi(B) = \sum_A \psi^*(A, B) \end{aligned}$$

Guarantees the local consistency  
→ Global consistency after iterations

- Initialized the potential functions


- $\psi(a, b) = P(a|b), \psi(b, c) = P(b|c)P(c)$
- $\phi(b) = 1$

- Example 1.  $P(b) = ?$


- $\phi^*(b) = \sum_a \psi(a, b) = 1$
- $\psi^*(b, c) = \psi(b, c) \frac{\phi^*(b)}{\phi(b)} = P(b|c)P(c) = P(b, c)$
- $\phi^{**}(b) = \sum_c \psi(b, c) = \sum_c P(b, c) = P(b)$
- $\psi^*(a, b) = \psi(a, b) \frac{\phi^{**}(b)}{\phi^*(b)} = \frac{P(a|b)P(b)}{1} = P(a, b)$
- $\phi^{***}(b) = \sum_a \psi^*(a, b) = P(b)$

- Example 2.  $P(b|a = 1, c = 1) = ?$

- $\phi^*(b) = \sum_a \psi(a, b) \delta(a = 1) = P(a = 1|b)$
- $\psi^*(b, c) = \psi(b, c) \frac{\phi^*(b)}{\phi(b)} = P(b|c = 1)P(c = 1) \frac{P(a=1|b)}{1}$
- $\phi^{**}(b) = \sum_c \psi(b, c) \delta(c = 1) = P(b|c = 1)P(c = 1)P(a = 1|b)$
- $\psi^*(a, b) = \psi(a, b) \frac{\phi^{**}(b)}{\phi^*(b)} = P(a = 1|b) \frac{P(b|c = 1)P(c=1)P(a = 1|b)}{P(a = 1|b)} = P(b|c = 1)P(c = 1)P(a = 1|b)$
- $\phi^{***}(b) = \sum_a \psi^*(a, b) \delta(a = 1) = P(b|c = 1)P(c = 1)P(a = 1|b)$





**Bayesian Network**



**Clique Graph**





- $$P(z_1, z_2, z_3, x_1, x_2, x_3) = \frac{\prod_N \psi(N)}{\prod_L \phi(L)}$$
  - $$= \frac{\psi(z_1, x_1) \psi(z_1, z_2) \psi(z_2, z_3) \psi(z_2, x_2) \psi(z_3, x_3)}{\phi(z_1) \phi_1(z_2) \phi_2(z_2) \phi(z_3)}$$

- $$P(z_1, z_2, z_3, x_1, x_2, x_3)$$
  - $$= P(z_1) P(x_1|z_1) P(z_2|z_1) P(x_2|z_2) P(z_3|z_2) P(x_3|z_3)$$

- Initialized the potential functions

- $$\psi(z_1, x_1) = P(z_1) P(x_1|z_1), \psi(z_1, z_2) = P(z_2|z_1), \psi(z_2, z_3) = P(z_3|z_2), \psi(z_2, x_2) = P(x_2|z_2), \psi(z_3, x_3) = P(x_3|z_3)$$
- $$\phi(z_1) = \phi_1(z_2) = \phi_2(z_2) = \phi(z_3) = 1$$

- Start absorbing and updating

- $$\phi_2^*(z_2) = \sum_{z_1} \psi^*(z_1, z_2) = \sum_{z_1} \psi(z_1, z_2) \phi^*(z_1) \phi_1^*(z_2) = \sum_{z_1} \psi(z_1, z_2) \phi^*(z_1) \phi_1^*(z_2)$$
  - Because  $x_2$  is already observed, so the summation on  $x_2$  does not happen, use just fixed  $x_2$
- $$= \sum_{z_1} P(z_2|z_1) \phi^*(z_1) P(x_2|z_2) = P(x_2|z_2) \sum_{z_1} P(z_2|z_1) \phi^*(z_1) = b_{idx(z_2), x_2} \sum_{i \in z_1} \alpha_{2-1}^i a_{i, z_2}$$
- Same as the forward probability calculation
- This is the upward process, then the downward process is same as the backward probability calculation

Define the update rule for separators

$$\phi^*(B) = \sum_A \psi^*(A, B)$$

Define the update rule for cliques

$$\psi^*(B, C) = \psi(B, C) \frac{\phi^*(B)}{\phi(B)}$$

Initial State probabilities

$$P(z_1) \sim \text{Mult}(\pi_1, \dots, \pi_k)$$

Transition probabilities

$$P(z_t | z_{t-1}^i = 1) \sim \text{Mult}(a_{i,1}, \dots, a_{i,k})$$

$$\text{Or, } P(z_t^j = 1 | z_{t-1}^i = 1) = a_{i,j}$$

Emission probabilities

$$P(x_t | z_t^i = 1) \sim \text{Mult}(b_{i,1}, \dots, b_{i,m})$$

$$\text{Or, } P(x_t^j = 1 | z_t^i = 1) = b_{i,j}$$

# Viterbi Decoding

Initial State probabilities

$$P(z_1) \sim \text{Mult}(\pi_1, \dots, \pi_k)$$

Transition probabilities

$$P(z_t | z_{t-1}^i = 1) \sim \text{Mult}(a_{i,1}, \dots, a_{i,k})$$

$$\text{Or, } P(z_t^j = 1 | z_{t-1}^i = 1) = a_{i,j}$$

Emission probabilities

$$P(x_t | z_t^i = 1) \sim \text{Mult}(b_{i,1}, \dots, b_{i,m})$$

$$\text{Or, } P(x_t^j = 1 | z_t^i = 1) = b_{i,j}$$


- $P(z_t^k = 1, X) = \alpha_t^k \beta_t^k = (b_{k,x_t} \sum_i \alpha_{t-1}^i a_{i,k}) \times (\sum_i a_{k,i} b_{i,x_t} \beta_{t+1}^i)$ 
  - This dictates the most probable assignment to a single latent variable,  $z_t$ , given the whole observed sequence,  $X$
  - $k_t^* = \text{argmax}_k P(z_t^k = 1 | X) = \text{argmax}_k P(z_t^k = 1, X) = \text{argmax}_k \alpha_t^k \beta_t^k$
  - What if we want to have the most probable assignment of  $Z$  given  $X$ ?
    - Exactly the decoding question
    - Different from the most probable assignment of a single latent variable
- Viterbi decoding
  - $k^* = \text{argmax}_k P(z^k = 1 | X) = \text{argmax}_k P(z^k = 1, X)$
  - Need to model the sequence of  $Z$ .
    - Let's use the forward approach (Bottom-up)
  - $V_t^k = \max_{z_1 \dots z_{t-1}} P(x_1, \dots, x_{t-1}, z_1, \dots, z_{t-1}, x_t, z_t^k = 1)$ 
    - Most probable sequence of latent states until  $t-1$  and fixing the state  $k$  at time  $t$
  - $= \max_{z_1 \dots z_{t-1}} P(x_t, z_t^k = 1 | x_1, \dots, x_{t-1}, z_1, \dots, z_{t-1}) P(x_1, \dots, x_{t-1}, z_1, \dots, z_{t-1})$
  - $= \max_{z_1 \dots z_{t-1}} P(x_t, z_t^k = 1 | z_{t-1}) P(x_1, \dots, x_{t-2}, z_1, \dots, z_{t-2}, x_{t-1}, z_{t-1})$
  - $= \max_{z_{t-1}} P(x_t, z_t^k = 1 | z_{t-1}) \max_{z_1 \dots z_{t-2}} P(x_1, \dots, x_{t-2}, z_1, \dots, z_{t-2}, x_{t-1}, z_{t-1})$
  - $= \max_{i \in Z_{t-1}} P(x_t, z_t^k = 1 | z_{t-1}^i = 1) V_{t-1}^i = \max_{i \in Z_{t-1}} P(x_t | z_t^k = 1) P(z_t^k = 1 | z_{t-1}^i = 1) V_{t-1}^i$
  - $= P(x_t | z_t^k = 1) \max_{i \in Z_{t-1}} P(z_t^k = 1 | z_{t-1}^i = 1) V_{t-1}^i = b_{k, \text{idx}(x_t)} \max_{i \in Z_{t-1}} a_{i,k} V_{t-1}^i$
  - Keep going until time  $T$

# Detour: Assembly Line Scheduling



**Assembly Line 1**


station  $S_{1,1}$  station  $S_{1,2}$  station  $S_{1,3}$  station  $S_{1,4}$  station  $S_{1,5}$  station  $S_{1,6}$



**Goal:** Computing the fastest production route

# Detour: Tracing Assembly Line Scheduling in DP

## Assembly Line 1



## Assembly Line 2

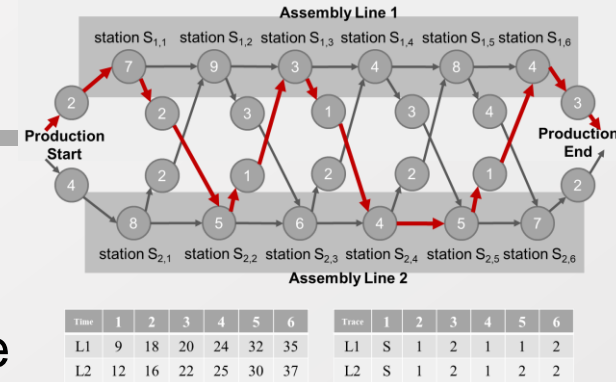
Used for retrace purpose

Time	1	2	3	4	5	6
L1	9	18	20	24	32	35
L2	12	16	22	25	30	37

Trace	1	2	3	4	5	6
L1	S	1	2	1	1	2
L2	S	1	2	1	2	2

# Viterbi Decoding Algorithm

- Need to know  $V_t^k$ 
  - Time X States
  - Store  $V_t^k$  Two variables to store the trace and the probability up to time  $t$ . Two memoization tables
  - Answering the decoding question with  $\pi, a, b, X$
- ViterbiDecodingAlgorithm
  - Initialize
    - $V_1^k = b_{k,x_1} \pi_k$
  - Iterate until time T
    - $V_t^k = b_{k,idx(x_t)} \max_{i \in Z_{t-1}} a_{i,k} V_{t-1}^i$
    - $trace_t^k = \operatorname{argmax}_{i \in Z_{t-1}} a_{i,k} V_{t-1}^i$
  - Return  $P(X, Z^*) = \max_k V_T^k, z_T^* = \operatorname{argmax}_k V_T^k, z_{t-1}^* = trace_t^{z_T^*}$
- Technical difficulties in the implementation
  - Very frequent underflow problems.
  - Turn this into the log domain  $\rightarrow$  from multiplication to summation



- Importance of  $\pi, a, b$ 
  - HMM parameters
  - Forward algorithm (evaluation) and Viterbi algorithm (decoding) depends on knowing  $\pi, a, b$
- However, knowing  $\pi, a, b$  assumes that we have observed X and Z
  - But, often Z is hard to observe. Need tagging, annotation, etc
  - Often the latent space is what we want to know, so we can't assume that we know Z
- If we don't know Z, we can assign the most probable Z to X
  - However, this is decoding problem, and this requires knowing  $\pi, a, b$
- Most likely scenario in the real world
  - You have only X
  - You don't have Z,  $\pi, a, b$ , and you need to find out Z,  $\pi, a, b$
- Strategy
  - Finding the optimized  $\bar{\pi}, \bar{a}, \bar{b}$  with X
  - Finding the most probable Z with X,  $\bar{\pi}, \bar{a}, \bar{b}$
  - How to find the unknown parameter of the latent distribution without supervision?
- **EM algorithm!**
  - **Iteratively optimizing  $\bar{\pi}, \bar{a}, \bar{b}$  and Z**

# Detour: EM Algorithm

$$\begin{aligned}l(\theta) &= \ln P(X|\theta) = \ln \left\{ \sum_Z q(Z) \frac{P(X, Z|\theta)}{q(Z)} \right\} \geq \sum_Z q(Z) \ln \frac{P(X, Z|\theta)}{q(Z)} = Q(\theta, q) \\Q(\theta, q) &= E_{q(Z)} \ln P(X, Z|\theta) + H(q) \\L(\theta, q) &= \ln P(X|\theta) - \sum_Z \left\{ q(Z) \ln \frac{q(Z)}{P(Z|X, \theta)} \right\}\end{aligned}$$

- EM algorithm
  - Finds the maximum likelihood solutions for models with latent variables
  - $P(X|\theta) = \sum_Z P(X, Z|\theta) \rightarrow \ln P(X|\theta) = \ln \{ \sum_Z P(X, Z|\theta) \}$
- EM algorithm
  - Initialize  $\theta^0$  to an arbitrary point
  - Loop until the likelihood converges
    - Expectation step
      - $q^{t+1}(z) = \operatorname{argmax}_q Q(\theta^t, q) = \operatorname{argmax}_q L(\theta^t, q) = \operatorname{argmin}_q KL(q || P(Z|X, \theta^t))$
      - $\rightarrow q^t(z) = P(Z|X, \theta) \rightarrow$  Assign Z by  $P(Z|X, \theta)$
    - Maximization step
      - $\theta^{t+1} = \operatorname{argmax}_\theta Q(\theta, q^{t+1}) = \operatorname{argmax}_\theta L(\theta, q^{t+1})$
      - $\rightarrow$  fixed Z means that there is no unobserved variables
      - $\rightarrow$  Same optimization of ordinary MLE



# EM for HMM

- EM algorithm for HMM

- Initialize  $\pi^0, a^0, b^0$  to an arbitrary point
- Loop until the likelihood converges
  - Expectation step
    - $q^{t+1}(z) = P(Z|X, \pi^t, a^t, b^t) \rightarrow$  Assign Z by  $P(Z|X, \pi^t, a^t, b^t)$
  - Maximization step
    - $\pi^{t+1}, a^{t+1}, b^{t+1} = \operatorname{argmax}_{\pi, a, b} Q(\pi, a, b, q^{t+1}) = \operatorname{argmax}_{\pi, a, b} E_{q^{t+1}(z)} \ln P(X, Z|\pi, a, b) + H(q)$

- Assign Z and optimize  $\pi, a, b$  alternatively

- Coordinated optimization
- How to optimize? Derivation of EM update formula from HMM?
- $P(X, Z|\pi, a, b) = \pi_{z_1} \prod_{t=2}^T a_{z_{t-1}, z_t} \prod_{t=1}^T b_{z_t, x_t}$
- $\ln P(X, Z|\pi, a, b) = \pi_{z_1} + \sum_{t=2}^T \ln a_{z_{t-1}, z_t} + \sum_{t=1}^T \ln b_{z_t, x_t}$
- $E_{q^{t+1}(z)} \ln P(X, Z|\pi, a, b) = \sum_Z q^{t+1}(z) \ln P(X, Z|\pi, a, b) = \sum_Z P(Z|X, \pi^t, a^t, b^t) \ln P(X, Z|\pi, a, b)$

$$\begin{aligned} l(\theta) &= \ln P(X|\theta) = \ln \left\{ \sum_Z q(Z) \frac{P(X, Z|\theta)}{q(Z)} \right\} \\ &\geq \sum_Z q(Z) \ln \frac{P(X, Z|\theta)}{q(Z)} = Q(\theta, q) \\ Q(\theta, q) &= E_{q(Z)} \ln P(X, Z|\theta) + H(q) \\ L(\theta, q) &= \ln P(X|\theta) - \sum_Z \{q(Z) \ln \frac{q(Z)}{P(Z|X, \theta)}\} \end{aligned}$$

**Initial State probabilities**

$$P(z_1) \sim \text{Mult}(\pi_1, \dots, \pi_k)$$

**Transition probabilities**

$$P(z_t | z_{t-1}^i = 1) \sim \text{Mult}(a_{i,1}, \dots, a_{i,k})$$

$$\text{Or, } P(z_t^j = 1 | z_{t-1}^i = 1) = a_{i,j}$$

**Emission probabilities**

$$P(x_t | z_t^i = 1) \sim \text{Mult}(b_{i,1}, \dots, b_{i,m})$$

$$\text{Or, } P(x_t^j = 1 | z_t^i = 1) = b_{i,j}$$



# Derivation of EM Update Formula

- $Q(\pi, a, b, q^{t+1}) = E_{q^{t+1}(z)} \ln P(X, Z | \pi, a, b)$
- $= \sum_Z q^{t+1}(z) \ln P(X, Z | \pi, a, b)$
- $= \sum_Z P(Z | X, \pi^t, a^t, b^t) \ln P(X, Z | \pi, a, b)$ 
  - $\ln P(X, Z | \pi, a, b)$
  - $= \ln \pi_{z_1} + \sum_{t=2}^T \ln a_{z_{t-1}, z_t} + \sum_{t=1}^T \ln b_{z_t, x_t}$
- $= \sum_Z P(Z | X, \pi^t, a^t, b^t) \{ \ln \pi_{z_1} + \sum_{t=2}^T \ln a_{z_{t-1}, z_t} + \sum_{t=1}^T \ln b_{z_t, x_t} \}$
- Need to optimize  $Q(\pi, a, b, q^{t+1})$  by using  $\pi, a, b$ 
  - Remember that  $\pi, a, b$  is actually probabilities.  $\sum_i \pi_i = 1$
  - Since there are constraints on  $\pi, a, b$  and  $Q$  is smooth  $\rightarrow$  Lagrange method!
- $L(\pi, a, b, q^{t+1})$
- $= Q(\pi, a, b, q^{t+1}) - \lambda_\pi (\sum_{i=1}^K \pi_i - 1) - \sum_{i=1}^K \lambda_{a_i} (\sum_{j=1}^K a_{i,j} - 1) - \sum_{i=1}^K \lambda_{b_i} (\sum_{j=1}^M b_{i,j} - 1)$
- Now, a typical optimization with partial derivative
- $\frac{dL(\pi, a, b, q^{t+1})}{d\pi_i} = \frac{d}{d\pi_i} \sum_Z P(Z | X, \pi^t, a^t, b^t) \ln \pi_{z_1} - \lambda_\pi (\sum_{i=1}^K \pi_i - 1) = 0$ 
  - Only the terms with  $z_1 = i$  survives
- $\frac{d}{d\pi_i} \{ \sum_Z P(Z | X, \pi^t, a^t, b^t) \ln \pi_{z_1} - \lambda_\pi (\sum_{i=1}^K \pi_i - 1) \} = \frac{P(z_1^i = 1 | X, \pi^t, a^t, b^t)}{\pi_i} - \lambda_\pi = 0 \rightarrow \pi_i = \frac{P(z_1^i = 1 | X, \pi^t, a^t, b^t)}{\lambda_\pi}$
- $\frac{d}{d\lambda_\pi} L(\pi, a, b, q^{t+1}) = -(\sum_{i=1}^K \pi_i - 1) = 0 \rightarrow \sum_{i=1}^K \pi_i = 1$
- Together,  $\pi^{t+1}_i = \frac{P(z_1^i = 1 | X, \pi^t, a^t, b^t)}{\sum_{j=1}^K P(z_1^j = 1 | X, \pi^t, a^t, b^t)}$ 
  - This is an update function for  $\pi_i$  at the M Step

Similarly, we can compute the update formula for  $a$  and  $b$  with the partial derivatives

$$a^{t+1}_{i,j} = \frac{\sum_{t=2}^T P(z_{t-1}^i = 1, z_t^j = 1 | X, \pi^t, a^t, b^t)}{\sum_{t=2}^T P(z_{t-1}^i = 1 | X, \pi^t, a^t, b^t)}$$

$$b^{t+1}_{i,j} = \frac{\sum_{t=1}^T P(z_t^i = 1 | X, \pi^t, a^t, b^t) \delta(\text{idx}(x_t) = j)}{\sum_{t=1}^T P(z_t^i = 1 | X, \pi^t, a^t, b^t)}$$

- Answer to the learning question of HMM
- Again, EM for HMM with more details
- EM algorithm for HMM, a.k.a. Baum-Welch, Forward-Backward...
  - Initialize  $\pi^0, a^0, b^0$  to an arbitrary point
  - Loop until the likelihood converges

- Expectation step

- $q^{t+1}(z) = P(Z|X, \pi^t, a^t, b^t) \rightarrow$  Assign Z by  $P(Z|X, \pi^t, a^t, b^t)$

- Maximization step

- $\pi^{t+1}_i = \frac{P(z_1^i = 1|X, \pi^t, a^t, b^t)}{\sum_{j=1}^K P(z_1^j = 1|X, \pi^t, a^t, b^t)}$

- $a^{t+1}_{i,j} = \frac{\sum_{t=2}^T P(z_{t-1}^i = 1, z_t^j = 1|X, \pi^t, a^t, b^t)}{\sum_{t=2}^T P(z_{t-1}^i = 1|X, \pi^t, a^t, b^t)}$

- $b^{t+1}_{i,j} = \frac{\sum_{t=1}^T P(z_t^i = 1|X, \pi^t, a^t, b^t) \delta(\text{idx}(x_t)=j)}{\sum_{t=1}^T P(z_t^i = 1|X, \pi^t, a^t, b^t)}$

$$\left[ \begin{array}{l} P(z_t^k = 1, X) = \alpha_t^k \beta_t^k \\ \alpha_t^k = b_{k,x_t} \sum_i \alpha_{t-1}^i a_{i,k} \\ \beta_t^k = \sum_i a_{k,i} b_{i,x_t} \beta_{t+1}^i \\ P(X) = \sum_i \alpha_T^i \end{array} \right]$$

- This slideset is greatly influenced
  - By Prof. Eric P. Xing at CMU

- Bishop Chapter 13