



# Diffusion Models

Il-Chul Moon

Department of Industrial and Systems Engineering

KAIST

[icmoon@kaist.ac.kr](mailto:icmoon@kaist.ac.kr)




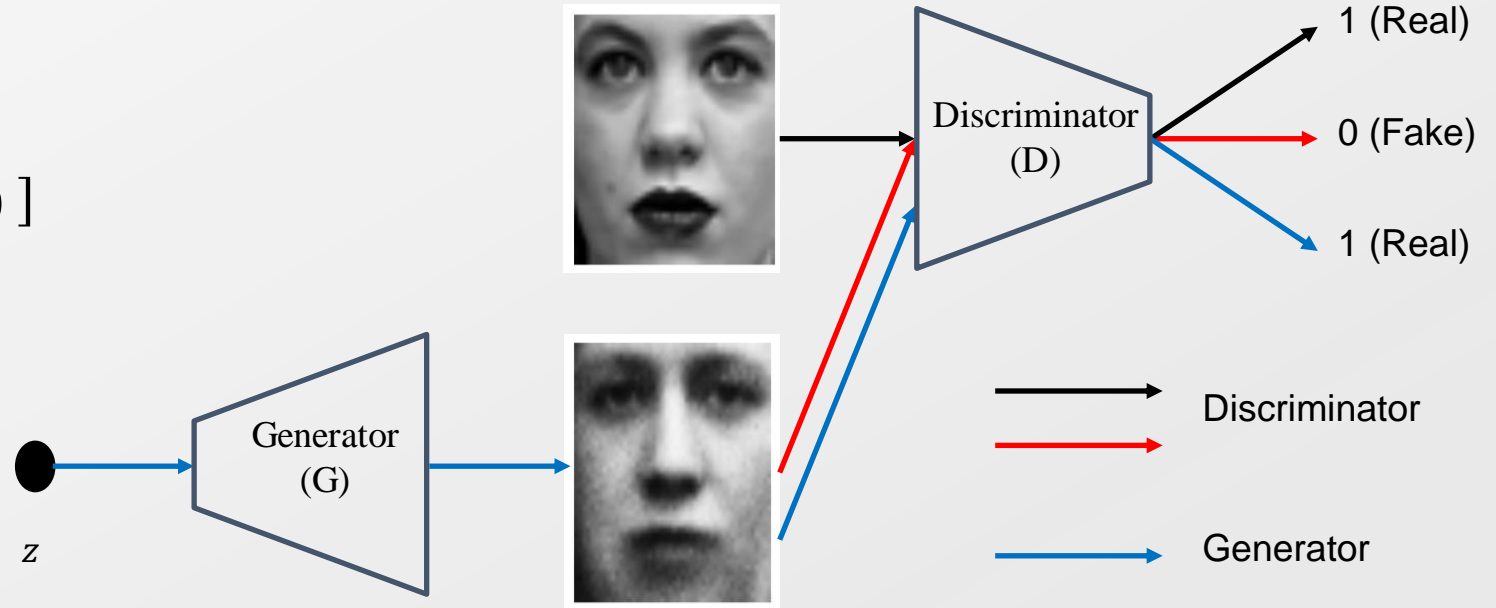
# Learning Structure of GAN

- Learning objective function

- $$\min_G \max_D E_{x \sim p_{data}(x)} [\log D(x)] + E_{z \sim p_z(z)} [\log(1 - D(G(z)))]$$

- Critical Issue of GAN

- Mode collapse, Non-convergence, Performance degradation


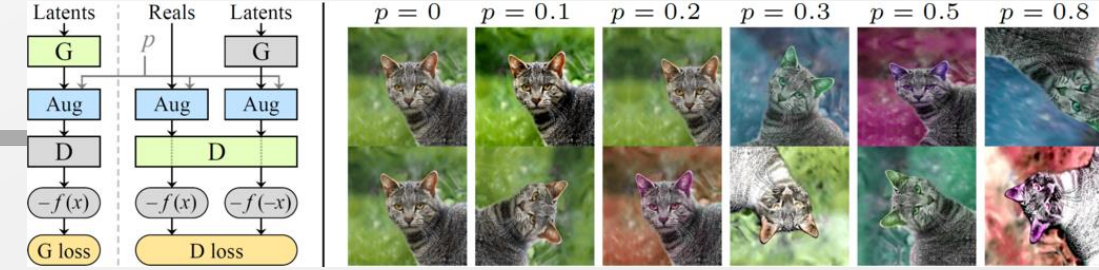


Generates only  
a subset of  
modes in  
 $p_{data}$

Metz, Luke, et al. "Unrolled generative adversarial networks." *arXiv preprint arXiv:1611.02163* (2016).

# Learning Divergence of GAN

- Learning objective function
  - $\min_G \max_D E_{x \sim p_{data}(x)} [\log D(x)] + E_{z \sim p_z(z)} [\log(1 - D(G(z)))]$
  - This can be generalized by  $f$ -divergence
    - $D_f(P||Q) = \int_x q(x) f\left(\frac{p(x)}{q(x)}\right) dx$
- Is it related to the critical Issue of GAN?
  - Yes. See the domain of  $f$  is  $\frac{p(x)}{q(x)}$
  - “No-man’s land” between the generated and the data instances?
- Method 1. Change the divergence to something else, i.e. Integral Probability Metrics
- Method 2. Fill the GAP with “augmented” data instances



Nowozin, Sebastian, Botond Cseke, and Ryota Tomioka. "f-gan: Training generative neural samplers using variational divergence minimization." *Proceedings of the 30th International Conference on Neural Information Processing Systems*. 2016.

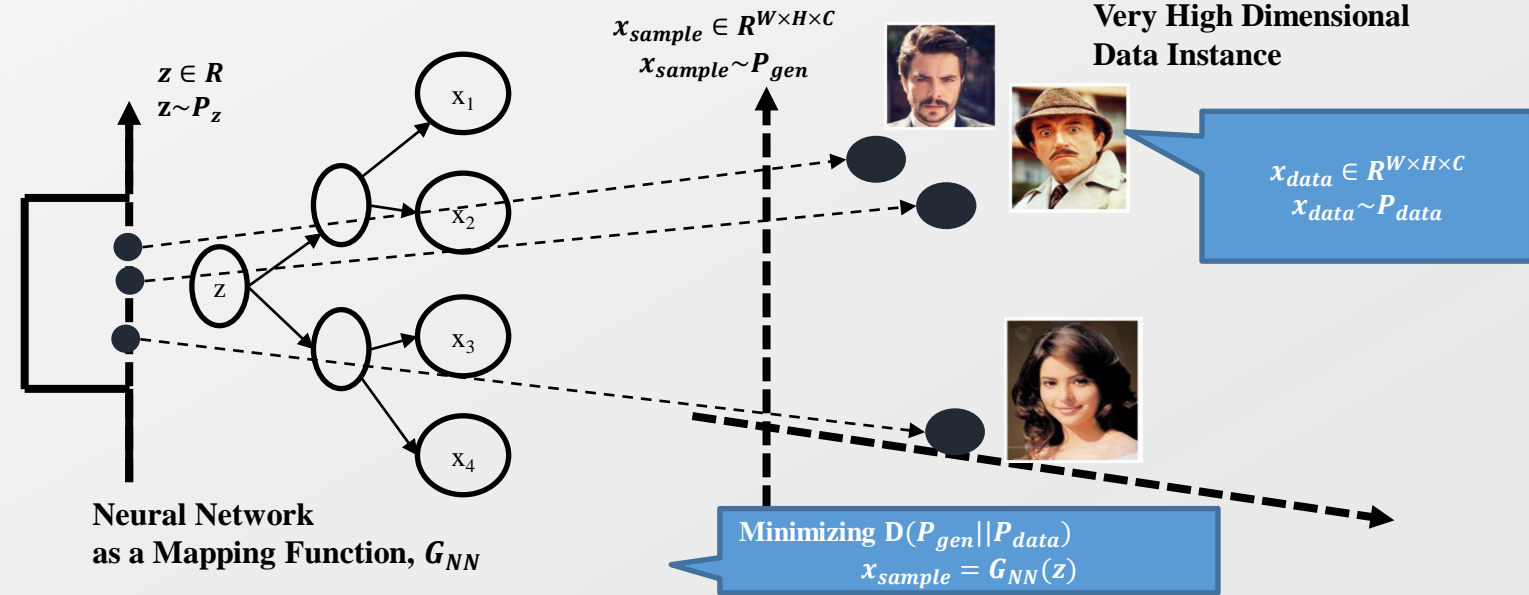
Karras, Tero, et al. "Training generative adversarial networks with limited data." *arXiv preprint arXiv:2006.06676* (2020).

# DIFFUSION MODEL

# How to Map Stochastic Latent Space to Data Space?

- GAN maps a stochastic sample to a data instance
  - Through Generator function approximated as a neural network

Low Stochastic  
Sample Instance

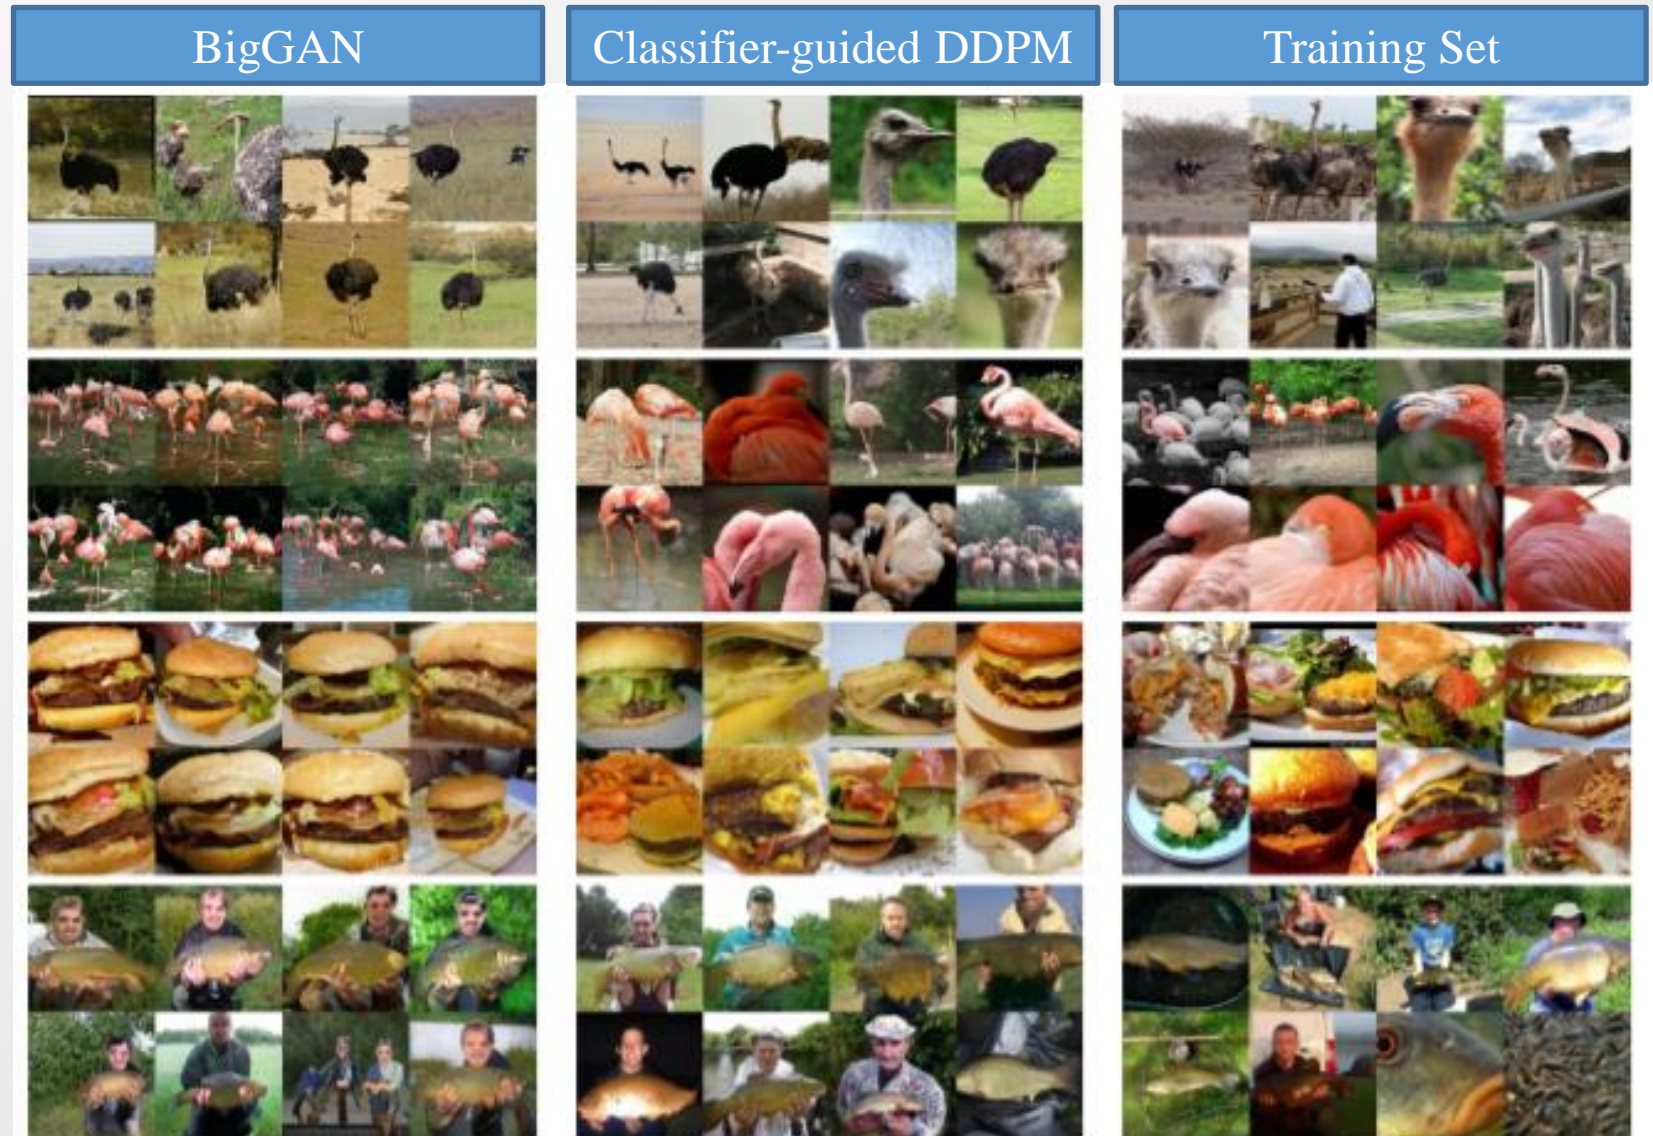


- Is the single-shot mapping good enough?
  - It could have been good enough
    - If the function is flexible enough
      - But, Training of generator  $\rightarrow$  MinMax problem  $\rightarrow$  Inconsistent gradient signal
    - If the mapping of the stochastic element covers the whole data space
      - But, Exploration/exploitation  $\rightarrow$  Mode collapse
- Is there any way to avoid the short-comings of the above?



# Diffusion Models Beat GANs on Image Synthesis


- Sampling performances
  - BigGAN-deep
    - With truncation 1.0
    - FID 6.95
  - DDPM
    - With guidance
    - FID 4.59
- Apparent mode collapse from GAN
  - Whereas, DDPM shows no mode collapse
- Is adversarial model better in sampling?
  - Which it can only do...



Dhariwal, Prafulla, and Alex Nichol. "Diffusion models beat gans on image synthesis." *arXiv preprint arXiv:2105.05233* (2021).




- Flow of particles from high-density regions towards low-density regions
  - Eventually, creating the high entropy in the particle distribution
- Considering a probability distribution over the data space,  $P_d$  on  $R^d$ 
  - $P_d$  : current complex distribution over the data space
  - $P_0$  : prior distribution over the data space
  - There will be a transformation in the particle distribution over the same space
    - $T: R^d \rightarrow R^d$
    - $x_0 \sim P_d$
    - $T(x_0) \sim P_0$
- We know such techniques of gradual transformation from one distribution to another
  - Markov chain Monte-Carlo
  - Each transition changes from an arbitrary distribution to a stationary distribution
  - Basics of sampling-based inference
    - Metropolis-Hastings algorithm
    - Gibbs sampling



© wikipedia

- Problem of the previous samplings?
  - No use of the past records → every sampling is independent
- Assigning Z values is a key in the inference
  - Let's assign the values by sampling result
    - Calculate  $P(E|MC=T, A=F) \rightarrow$  Toss a biased coin to assign a value to E
- Sequence of random variables such a process moves through, with the Markov property defining serial dependence only between adjacent periods (as in a "chain")
- A Markov chain is a stochastic process with the Markov property
  - Example) First-order Markov chain




- $p(\mathbf{z}^{(m+1)} | \mathbf{z}^{(1)}, \dots, \mathbf{z}^{(m)}) = p(\mathbf{z}^{(m+1)} | \mathbf{z}^{(m)}), m \in \{1, \dots, M - 1\}$
- Describing systems that follow a chain of linked events, where what happens next depends only on the current state of the system



- Traditional Markov Chain analysis :

- A transition rule,  $p(z^{(t+1)} | z^{(t)})$ , is given,
- Interested in finding the stationary distribution  $\pi(z)$

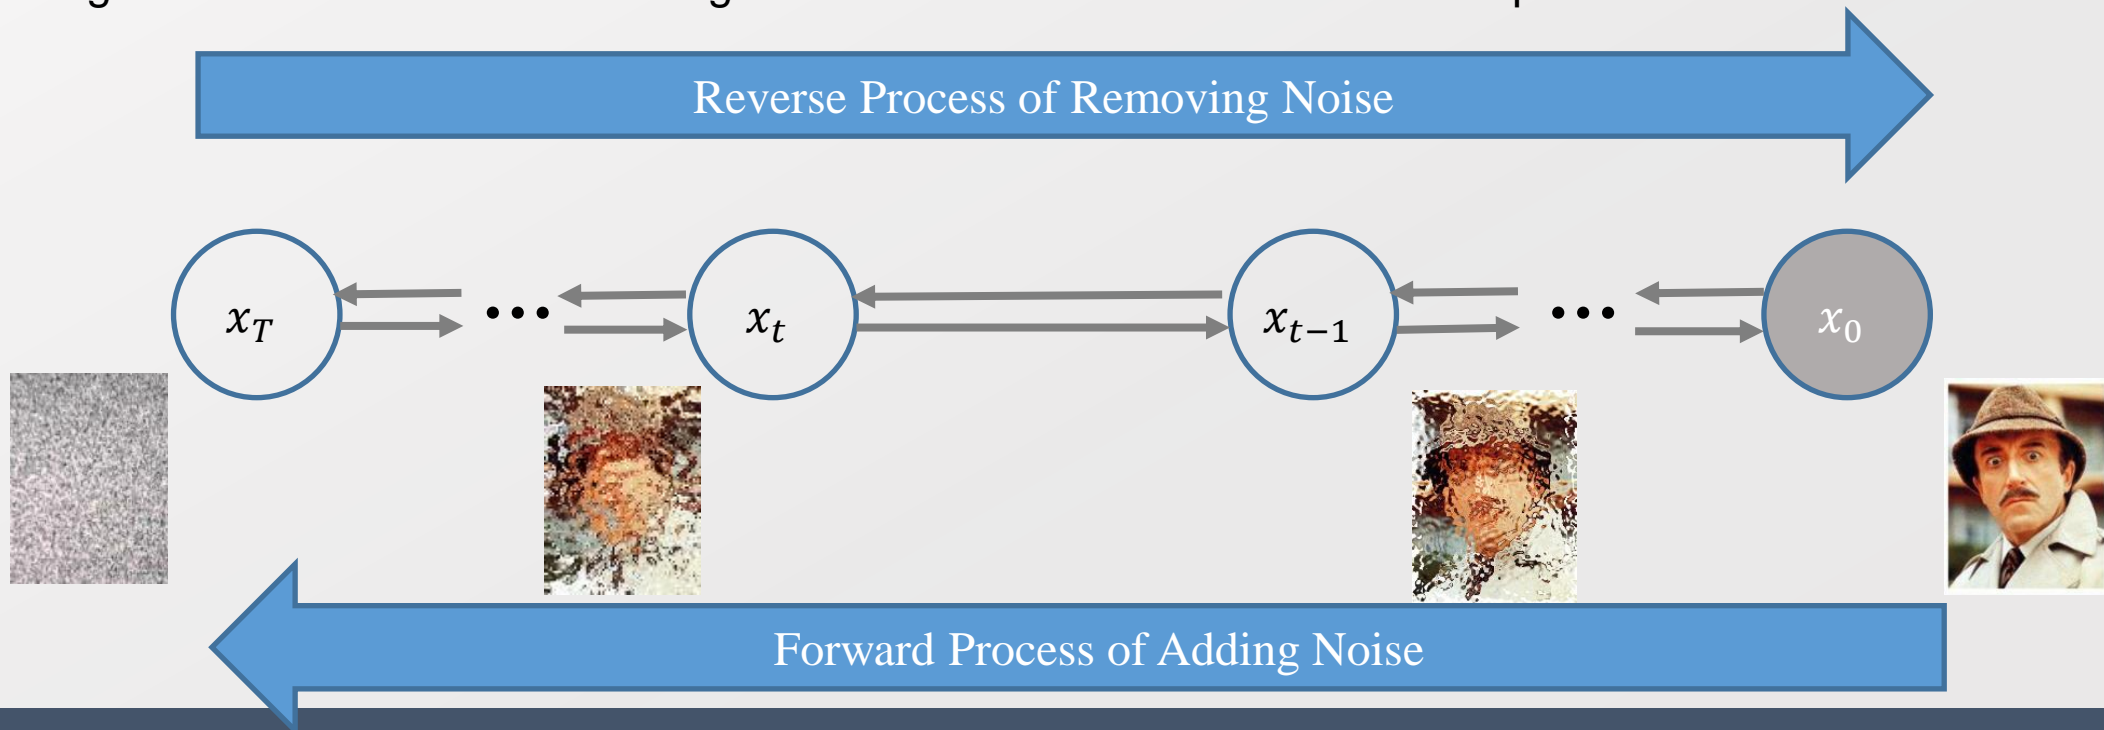


- Markov chain Monte Carlo(MCMC) :

- A target stationary distribution  $\pi(z)$  is known,
- Interested in prescribing an efficient transition rule to reach the stationary distribution
- Algorithms for sampling from probability distributions based on constructing a Markov chain that has the desired distribution  $\pi(z)$
- Starting from an arbitrary state, the Markov chain proceeds

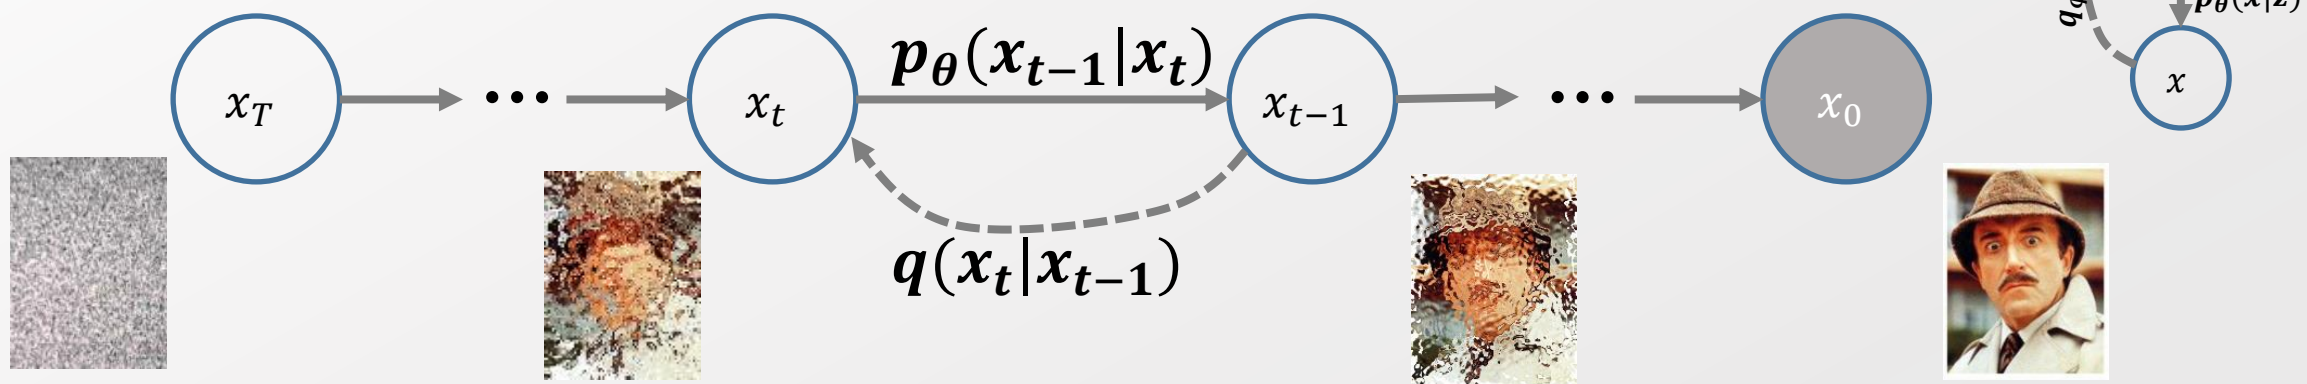
$$\underbrace{z^{(1)} \rightarrow z^{(2)} \rightarrow \dots \rightarrow z^{(m)}}_{\text{Burn-in period}} \rightarrow \underbrace{z^{(m+1)} \rightarrow z^{(m+2)} \rightarrow \dots \rightarrow z^{(m+n)}}_{\text{Treat them as samples from } \pi(x)}$$

- Let's take an image
  - Forward Process : Gradually add noise, so the image becomes a simple noise
  - Reverse Process : Gradually remove noise, so the noise becomes the image
- This is still a mapping function from the latent space of the noise to the data space
  - There is a single approximated function, the denoised function (Assuming adding a noise becomes trivial)
  - There is no game of generator vs. discriminator
    - The gradient signal is not dis-aligned
    - The generator is not fixated to a single mode to induce the discriminator output.



# Diffusion Model (Score-Based Model)

- Generative model + Neural network
  - Neural network is used for the inference on the diffusion process



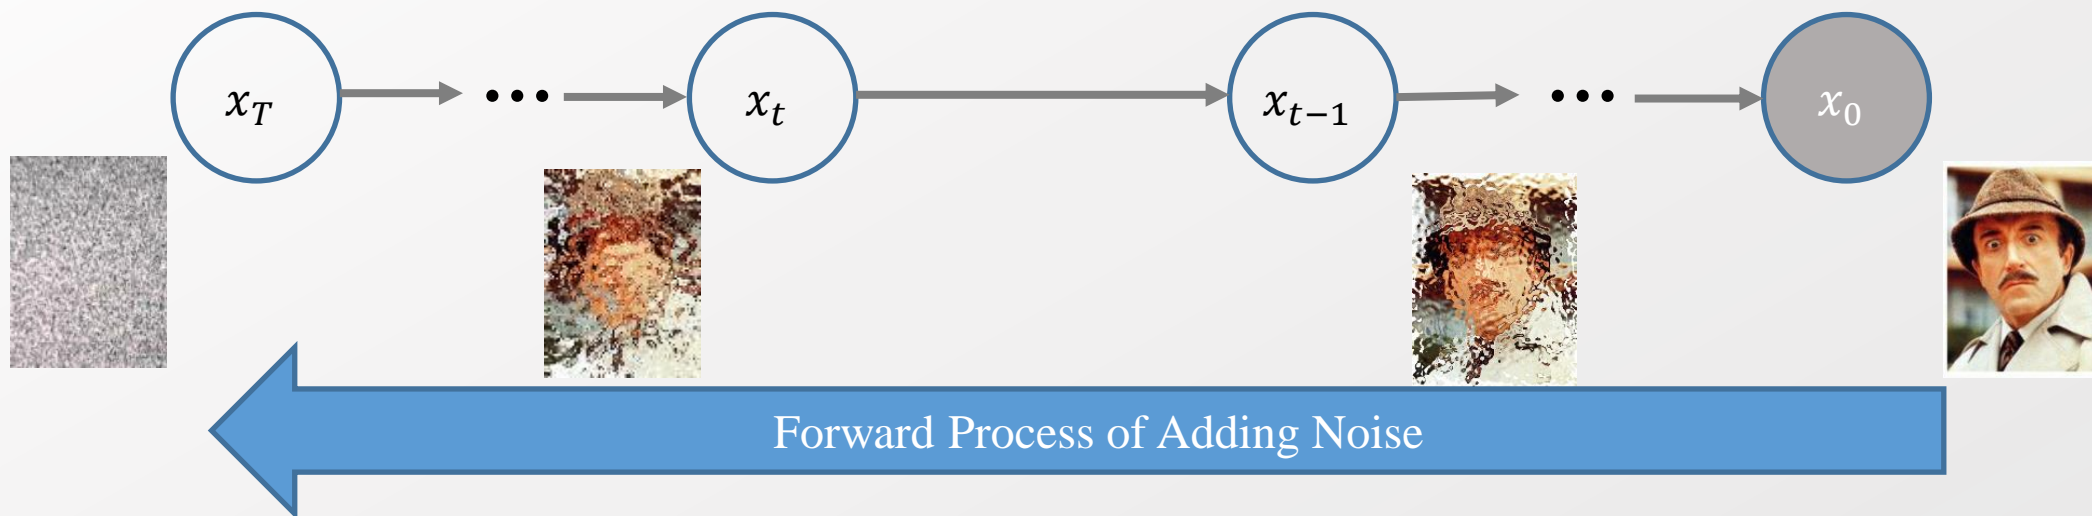
- Basically, a Markov chain by shaping an instance, or destroying the structure on an instance

$$p_\theta(x_{0:T}) = p(x_T) \prod_{t=1}^T p_\theta(x_{t-1}|x_t), p_\theta(x_{t-1}|x_t) = N(x_{t-1}; \mu_\theta(x_t, t), \Sigma_\theta(x_t, t))$$

$$q(x_{1:T}|x_0) = \prod_{t=1}^T q(x_t|x_{t-1}), q(x_t|x_{t-1}) = N(x_t; \sqrt{1 - \beta_t}x_{t-1}, \beta_t I)$$

- Similar to VAEs, given the neural-network based inference on the distribution parameters
  - Not exactly the same, VAE does not add potential noises to an instance over time, or over layers
  - Some group of VAEs with hierarchical  $z$  becomes similar to the diffusion model

# Forward Diffusion Process of Adding Noise



- Stochastic mapping from the data space to the latent space
  - Can be regarded as an encoding process
  - Instead, this is going to be a fixed encoder without learning
- As a single-step adding noise,
  - $q(x_t|x_{t-1}) = N(x_t; \sqrt{1 - \beta_t}x_{t-1}, \beta_t I)$
  - Actual sampling could be re-parametrized as
    - $x_t = \sqrt{1 - \beta_t}x_{t-1} + \sqrt{\beta_t}\epsilon_{t-1}, \epsilon_{t-1} \sim N(0, I)$
- This single-step noise addition will take a long-time if  $T \rightarrow \infty$ ,
  - So, we need a solution of this stochastic process, which is  $q(x_t|x_0)$

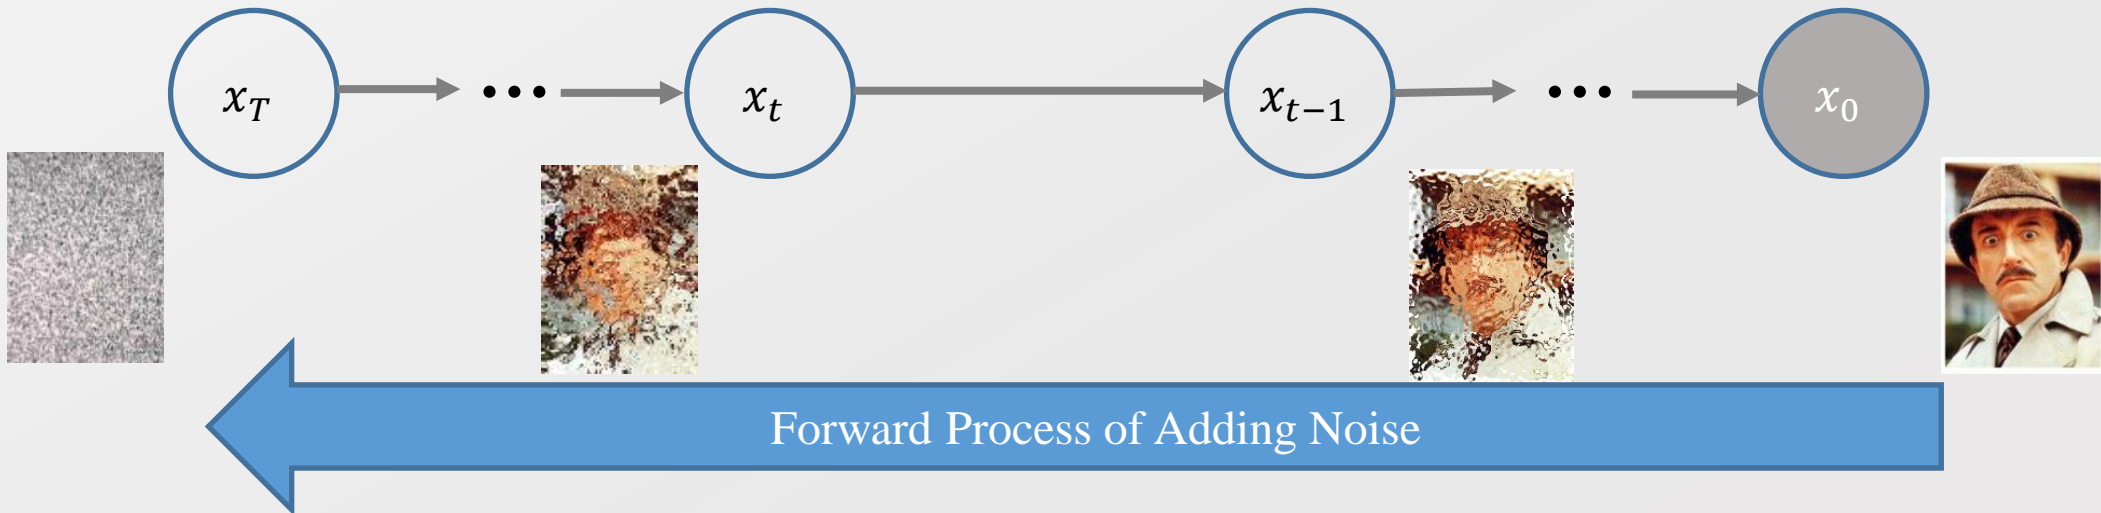

- $q(x_t|x_{t-1}) = N(x_t; \sqrt{1 - \beta_t}x_{t-1}, \beta_t I)$
- Generalizing  $q(x_t|x_0)$  by mathematical induction
  - $q(x_1|x_0) = N(x_1; \sqrt{1 - \beta_1}x_0, \beta_1 I)$
- Let's assume  $q(x_{t-1}|x_0) = N\left(x_{t-1}; \sqrt{\prod_{s=1}^{t-1}(1 - \beta_s)}x_0, (1 - \prod_{s=1}^{t-1}(1 - \beta_s))I\right)$
- Assuming  $x_{t-1}$  and  $\epsilon_{t-1}$  follow the Gaussian distribution
  - $x_t = \sqrt{1 - \beta_t}x_{t-1} + \sqrt{\beta_t}\epsilon_{t-1}$ ,  $\epsilon_{t-1} \sim N(0, I)$
  - Sum of Gaussian distributions are still Gaussian distribution
    - $X \sim N(\mu_X, \sigma_X^2)$ ,  $Y \sim N(\mu_Y, \sigma_Y^2)$ ,  $Z = X + Y$ ,  $Z \sim N(\mu_X + \mu_Y, \sigma_X^2 + \sigma_Y^2)$
- Let's say
  - $X = \sqrt{1 - \beta_t}x_{t-1} \sim N\left(\sqrt{1 - \beta_t}\sqrt{\prod_{s=1}^{t-1}(1 - \beta_s)}x_0, (\sqrt{1 - \beta_t})^2(1 - \prod_{s=1}^{t-1}(1 - \beta_s))I\right)$
  - $Y = \sqrt{\beta_t}\epsilon_{t-1} \sim N\left(0, (\sqrt{\beta_t})^2 I\right)$
- Then,  $x_t = \sqrt{1 - \beta_t}x_{t-1} + \sqrt{\beta_t}\epsilon_{t-1}$
- $= X + Y \sim N\left(\sqrt{\prod_{s=1}^t(1 - \beta_s)}x_0, (1 - \beta_t)(1 - \prod_{s=1}^{t-1}(1 - \beta_s))I + \beta_t I\right)$ 

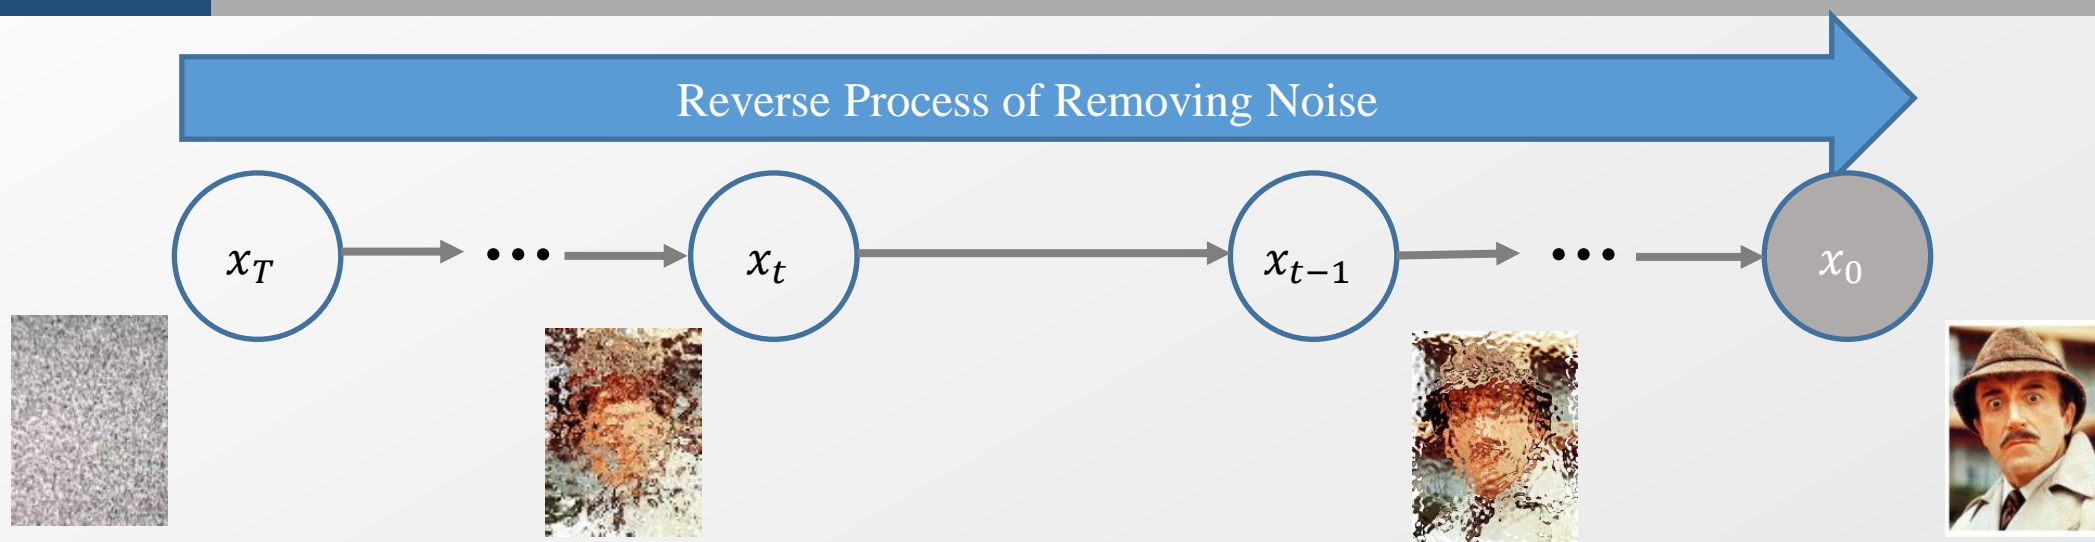
$$= N\left(\sqrt{\prod_{s=1}^t(1 - \beta_s)}x_0, \left(1 - \prod_{s=1}^t(1 - \beta_s)\right)I\right)$$

$$\begin{aligned}
 & (1 - \beta_t) \left(1 - \prod_{s=1}^{t-1}(1 - \beta_s)\right) + \beta_t \\
 &= 1 - \prod_{s=1}^{t-1}(1 - \beta_s) - \beta_t + \beta_t \prod_{s=1}^{t-1}(1 - \beta_s) + \beta_t \\
 &= 1 - \prod_{s=1}^{t-1}(1 - \beta_s) + \beta_t \prod_{s=1}^{t-1}(1 - \beta_s) \\
 &= 1 - (1 - \beta_t) \prod_{s=1}^{t-1}(1 - \beta_s) \\
 &= 1 - \prod_{s=1}^t(1 - \beta_s)
 \end{aligned}$$



- $q(x_t|x_0) = N\left(\sqrt{\prod_{s=1}^t(1-\beta_s)}x_0, (1-\prod_{s=1}^t(1-\beta_s))I\right) = N(\sqrt{\bar{\alpha}_t}x_0, (1-\bar{\alpha}_t)I)$ 
  - $\bar{\alpha}_t = \prod_{s=1}^t(1-\beta_s)$
  - $x_t = \sqrt{\bar{\alpha}_t}x_0 + \sqrt{(1-\bar{\alpha}_t)}\epsilon, \epsilon \sim N(0, I)$
- $\bar{\alpha}_T \rightarrow 0 \Rightarrow q(x_T|x_0) \approx N(0, I)$  : This requires a schedule on  $\beta_t$
- Effect of the forward diffusion from the diffusion kernel perspective
  - $q(x_t) = \int q(x_0, x_t)dx_0 = \int q(x_t|x_0)q(x_0)dx_0$
  - Here,  $q(x_t|x_0)$  becomes the kernel of the Gaussian convolution.
  - This shows the sampling of  $x_t \sim q(x_t|x_0)$ 
    - Sample  $x_0 \sim N(0, I) \rightarrow$  Sample  $x_t \sim q(x_t|x_0)$  : a.k.a. ancestral sampling





- Forward diffusion provides no real merit in the generation task
  - On the opposite, reverse diffusion will directly become the generation of samples from  $N(0, I)$
  - $x_T \sim N(0, I)$
  - $x_{t-1} \sim p(x_{t-1} | x_t)$  : True denoising distribution
    - In simple approach :  $p(x_{t-1} | x_t) \propto q(x_{t-1})q(x_t | x_{t-1}) \rightarrow$  This becomes intractable in sampling
    - Then, we need a direct approximation on  $p(x_{t-1} | x_t)$ 
      - By the flexibility and complexity of neural networks
- $p_\theta(x_{t-1} | x_t) = N(\mu_\theta(x_t, t), \sigma_t^2 I)$ ,  $p(x_T) = N(0, I)$ 
  - Fixed covariance structure and only trainable function of  $\mu_\theta(x_t, t)$

- Our goal becomes the training of  $p_\theta(x_{t-1}|x_t) = N(\mu_\theta(x_t, t), \sigma_t^2 I)$
- Therefore, we need to compare the distribution between  $p_\theta(x_{t-1}|x_t)$  and  $q(x_t|x_{t-1})$  given  $x_0$ 
  - $p_\theta(x_{t-1}|x_t) = N(\mu_\theta(x_t, t), \sigma_t^2 I)$
  - $q(x_t|x_{t-1}) = N(x_t; \sqrt{1 - \beta_t}x_{t-1}, \beta_t I)$ 
    - $q(x_t|x_0) = N(\sqrt{\bar{\alpha}_t}x_0, (1 - \bar{\alpha}_t)I)$ 
      - $\bar{\alpha}_t = \prod_{s=1}^t (1 - \beta_s) = \alpha_t \bar{\alpha}_{t-1}$
      - $x_t = \sqrt{\bar{\alpha}_t}x_0 + \sqrt{(1 - \bar{\alpha}_t)}\epsilon, \epsilon \sim N(0, I)$
      - $x_0 = \frac{1}{\sqrt{\bar{\alpha}_t}}(x_t - \sqrt{(1 - \bar{\alpha}_t)}\epsilon), \epsilon \sim N(0, I)$
    - $q(x_{t-1}|x_t, x_0) = N(\tilde{\mu}_t(x_t, x_0), \tilde{\beta}_t I)$ 
      - The convolution of  $x_t$  and  $x_0$
      - $\tilde{\mu}_t(x_t, x_0) = \frac{\sqrt{\bar{\alpha}_{t-1}}\beta_t}{1 - \bar{\alpha}_t}x_0 + \frac{\sqrt{\bar{\alpha}_t}(1 - \bar{\alpha}_{t-1})}{1 - \bar{\alpha}_t}x_t$
      - $\tilde{\beta}_t = \frac{1 - \bar{\alpha}_{t-1}}{1 - \bar{\alpha}_t}\beta_t$
- The match becomes  $q(x_{t-1}|x_t, x_0)$  and  $p_\theta(x_{t-1}|x_t)$

# Loss Structure from Reverse Diffusion (1)

- Our goal becomes the training of  $p_\theta(x_{t-1}|x_t) = N(\mu_\theta(x_t, t), \sigma_t^2 I)$
- Therefore, we need to compare the distribution between  $p_\theta(x_{t-1}|x_t)$  and  $q(x_t|x_{t-1})$

- $p_\theta(x_{t-1}|x_t) = N(\mu_\theta(x_t, t), \sigma_t^2 I)$
- $q(x_{t-1}|x_t, x_0) = N(\tilde{\mu}_t(x_t, x_0), \tilde{\beta}_t I)$

- Let's match the mean function  $\mu_\theta(x_t, t)$  to  $\tilde{\mu}_t(x_t, x_0)$  in  $N(\mu_\theta(x_t, t), \sigma_t^2 I)$

- $\sigma_t^2 = \tilde{\beta}_t$

- $L_{t-1} = E_{x_0 \sim p(x_0), \epsilon \sim N(0, I)} \left[ \frac{1}{2\sigma_t^2} \left\| \tilde{\mu}_t(x_t, x_0) - \mu_\theta(x_t, t) \right\|^2 \right]$

$$= E_{x_0 \sim p(x_0), \epsilon \sim N(0, I)} \left[ \frac{1}{2\sigma_t^2} \left\| \tilde{\mu}_t \left( x_t, \frac{1}{\sqrt{\bar{\alpha}_t}} (x_t - \sqrt{(1 - \bar{\alpha}_t)} \epsilon) \right) - \mu_\theta(x_t, t) \right\|^2 \right]$$

- Then, we need to take the inputs of  $\tilde{\mu}_t$  to compare  $\mu_\theta$

- How to simplify  $\tilde{\mu}_t \left( x_t, \frac{1}{\sqrt{\bar{\alpha}_t}} (x_t - \sqrt{(1 - \bar{\alpha}_t)} \epsilon) \right)$ ?

$$\begin{aligned} \tilde{\mu}_t(x_t, x_0) &= \frac{\sqrt{\bar{\alpha}_{t-1}}\beta_t}{1 - \bar{\alpha}_t} x_0 + \frac{\sqrt{\bar{\alpha}_t}(1 - \bar{\alpha}_{t-1})}{1 - \bar{\alpha}_t} x_t \\ \tilde{\beta}_t &= \frac{1 - \bar{\alpha}_{t-1}}{1 - \bar{\alpha}_t} \beta_t \\ \bar{\alpha}_t &= \prod_{s=1}^t (1 - \beta_s) = (1 - \beta_t) \bar{\alpha}_{t-1} = \alpha_t \bar{\alpha}_{t-1} \end{aligned}$$

$$x_0 = \frac{1}{\sqrt{\bar{\alpha}_t}} (x_t - \sqrt{(1 - \bar{\alpha}_t)} \epsilon)$$

# Loss Structure from Reverse Diffusion (2)

- $$\begin{aligned}
 \tilde{\mu}_t(x_t, x_0) &= \tilde{\mu}_t\left(x_t, \frac{1}{\sqrt{\bar{\alpha}_t}}(x_t - \sqrt{(1 - \bar{\alpha}_t)}\epsilon)\right) \\
 &= \frac{\sqrt{\bar{\alpha}_{t-1}}\beta_t}{1 - \bar{\alpha}_t} \left( \frac{1}{\sqrt{\bar{\alpha}_t}}(x_t - \sqrt{(1 - \bar{\alpha}_t)}\epsilon) \right) + \frac{\sqrt{\alpha_t}(1 - \bar{\alpha}_{t-1})}{1 - \bar{\alpha}_t} x_t \\
 &= \frac{\sqrt{\bar{\alpha}_{t-1}}\beta_t}{1 - \bar{\alpha}_t} \frac{1}{\sqrt{\bar{\alpha}_t}} x_t + \frac{\sqrt{\alpha_t}(1 - \bar{\alpha}_{t-1})}{1 - \bar{\alpha}_t} x_t - \frac{\sqrt{\bar{\alpha}_{t-1}}\beta_t}{1 - \bar{\alpha}_t} \frac{\sqrt{(1 - \bar{\alpha}_t)}}{\sqrt{\bar{\alpha}_t}} \epsilon \\
 &= \frac{1}{(1 - \bar{\alpha}_t)\sqrt{\bar{\alpha}_t}} \left( \sqrt{\bar{\alpha}_{t-1}}\beta_t + \sqrt{\bar{\alpha}_t}\sqrt{\alpha_t}(1 - \bar{\alpha}_{t-1}) \right) x_t - \frac{\sqrt{\bar{\alpha}_{t-1}}}{\sqrt{(1 - \bar{\alpha}_t)}\sqrt{\bar{\alpha}_t}} \beta_t \epsilon \\
 &= \frac{1}{(1 - \bar{\alpha}_t)\sqrt{\bar{\alpha}_t}} \left( \sqrt{\bar{\alpha}_{t-1}}\beta_t + \sqrt{\bar{\alpha}_{t-1}}\sqrt{\alpha_t}\sqrt{\bar{\alpha}_t}(1 - \bar{\alpha}_{t-1}) \right) x_t - \frac{\sqrt{\bar{\alpha}_{t-1}}}{\sqrt{(1 - \bar{\alpha}_t)}\sqrt{\alpha_t\bar{\alpha}_{t-1}}} \beta_t \epsilon \\
 &= \frac{\sqrt{\bar{\alpha}_{t-1}}}{(1 - \bar{\alpha}_t)\sqrt{\alpha_t\bar{\alpha}_{t-1}}} (\beta_t + \alpha_t(1 - \bar{\alpha}_{t-1})) x_t - \frac{\beta_t}{\sqrt{(1 - \bar{\alpha}_t)}\sqrt{\alpha_t}} \epsilon \\
 &= \frac{1}{(1 - \bar{\alpha}_t)\sqrt{\alpha_t}} (1 - \alpha_t + \alpha_t(1 - \bar{\alpha}_{t-1})) x_t - \frac{\beta_t}{\sqrt{(1 - \bar{\alpha}_t)}\sqrt{\alpha_t}} \epsilon \\
 &= \frac{1}{(1 - \alpha_t\bar{\alpha}_{t-1})\sqrt{\alpha_t}} (1 - \alpha_t\bar{\alpha}_{t-1}) x_t - \frac{\beta_t}{\sqrt{(1 - \bar{\alpha}_t)}\sqrt{\alpha_t}} \epsilon = \frac{1}{\sqrt{\alpha_t}} x_t - \frac{\beta_t}{\sqrt{(1 - \bar{\alpha}_t)}\sqrt{\alpha_t}} \epsilon = \frac{1}{\sqrt{\alpha_t}} \left( x_t - \frac{\beta_t}{\sqrt{(1 - \bar{\alpha}_t)}} \epsilon \right)
 \end{aligned}$$

$$\tilde{\mu}_t(x_t, x_0) = \frac{\sqrt{\bar{\alpha}_{t-1}}\beta_t}{1 - \bar{\alpha}_t} x_0 + \frac{\sqrt{\alpha_t}(1 - \bar{\alpha}_{t-1})}{1 - \bar{\alpha}_t} x_t$$

$$x_0 = \frac{1}{\sqrt{\bar{\alpha}_t}} (x_t - \sqrt{(1 - \bar{\alpha}_t)}\epsilon)$$

$$\bar{\alpha}_t = \prod_{s=1}^t (1 - \beta_s) = \alpha_t \bar{\alpha}_{t-1}$$



# Loss Structure from Reverse Diffusion (3)

- $L_{t-1} = E_{x_0 \sim p(x_0), \epsilon \sim N(0, I)} \left[ \frac{1}{2\sigma_t^2} \left\| \tilde{\mu}_t(x_t, x_0) - \mu_\theta(x_t, t) \right\|^2 \right]$

$$= E_{x_0 \sim p(x_0), \epsilon \sim N(0, I)} \left[ \frac{1}{2\sigma_t^2} \left\| \tilde{\mu}_t \left( x_t, \frac{1}{\sqrt{\bar{\alpha}_t}} (x_t - \sqrt{(1 - \bar{\alpha}_t)} \epsilon) \right) - \mu_\theta(x_t, t) \right\|^2 \right]$$

We assume the parameterization of

$$\begin{aligned} \mu_\theta(x_t, t) &= \tilde{\mu}_t \left( x_t, \frac{1}{\sqrt{\bar{\alpha}_t}} (x_t - \sqrt{(1 - \bar{\alpha}_t)} \epsilon_\theta(x_t, t)) \right) \\ &= \frac{1}{\sqrt{\bar{\alpha}_t}} \left( x_t - \frac{\beta_t}{\sqrt{(1 - \bar{\alpha}_t)}} \epsilon_\theta(x_t, t) \right) \end{aligned}$$

$$= E_{x_0 \sim p(x_0), \epsilon \sim N(0, I)} \left[ \frac{1}{2\sigma_t^2} \left\| \frac{1}{\sqrt{\bar{\alpha}_t}} \left( x_t - \frac{\beta_t}{\sqrt{(1 - \bar{\alpha}_t)}} \epsilon \right) - \mu_\theta(x_t, t) \right\|^2 \right]$$

$$= E_{x_0 \sim p(x_0), \epsilon \sim N(0, I)} \left[ \frac{1}{2\sigma_t^2} \left\| \frac{1}{\sqrt{\bar{\alpha}_t}} \left( x_t(x_0, \epsilon) - \frac{\beta_t}{\sqrt{(1 - \bar{\alpha}_t)}} \epsilon \right) - \frac{1}{\sqrt{\bar{\alpha}_t}} \left( x_t(x_0, \epsilon) - \frac{\beta_t}{\sqrt{(1 - \bar{\alpha}_t)}} \epsilon_\theta(x_t, t) \right) \right\|^2 \right]$$

$$= E_{x_0 \sim p(x_0), \epsilon \sim N(0, I)} \left[ \frac{1}{2\sigma_t^2} \left\| \frac{1}{\sqrt{\bar{\alpha}_t}} \left( -\frac{\beta_t}{\sqrt{(1 - \bar{\alpha}_t)}} \epsilon + \frac{\beta_t}{\sqrt{(1 - \bar{\alpha}_t)}} \epsilon_\theta(x_t, t) \right) \right\|^2 \right]$$

$$= E_{x_0 \sim p(x_0), \epsilon \sim N(0, I)} \left[ \frac{\beta_t^2}{2\sigma_t^2 \alpha_t (1 - \bar{\alpha}_t)} \left\| \epsilon - \epsilon_\theta(x_t, t) \right\|^2 \right]$$

$$= E_{x_0 \sim p(x_0), \epsilon \sim N(0, I)} \left[ \frac{\beta_t^2}{2\sigma_t^2 \alpha_t (1 - \bar{\alpha}_t)} \left\| \epsilon - \epsilon_\theta(\sqrt{\bar{\alpha}_t} x_0 + \sqrt{(1 - \bar{\alpha}_t)} \epsilon, t) \right\|^2 \right]$$

$$x_t = \sqrt{\bar{\alpha}_t} x_0 + \sqrt{(1 - \bar{\alpha}_t)} \epsilon$$

- Loss structure for a single step of Markov chain
  - $x_0$  is the only observed variable, and  $x_T \dots x_1$  are all latent variables
  - Since they are latent variables, there should be an ELBO structure for the MLE learning
    - To match the loss direction as a minimization and the ELBO direction as a maximization; we will use the negative ELBO to minimize “NELBO”.

$$\begin{aligned} E[-\log p_\theta(x_0)] &\leq E_q \left[ -\log \frac{p_\theta(x_{0:T})}{q(x_{1:T}|x_0)} \right] = E_q \left[ -\log \frac{p(x_T) \prod_{1 \leq t \leq T} p_\theta(x_{t-1}|x_t)}{\prod_{1 \leq t \leq T} q(x_t|x_{t-1})} \right] \\ &= E_q \left[ -\log p(x_T) - \sum_{1 \leq t \leq T} \log \frac{p_\theta(x_{t-1}|x_t)}{q(x_t|x_{t-1})} \right] = L \end{aligned}$$



# Loss Structure from Entire Reverse Diffusion (2)

- $$L = E_q \left[ -\log p(x_T) - \sum_{1 \leq t \leq T} \log \frac{p_\theta(x_{t-1}|x_t)}{q(x_t|x_{t-1})} \right]$$

$$= E_q \left[ -\log p(x_T) - \sum_{1 < t \leq T} \log \frac{p_\theta(x_{t-1}|x_t)}{q(x_t|x_{t-1})} - \log \frac{p_\theta(x_0|x_1)}{q(x_1|x_0)} \right]$$

$$= E_q \left[ -\log p(x_T) - \sum_{1 < t \leq T} \log \frac{p_\theta(x_{t-1}|x_t)}{q(x_{t-1}|x_t, x_0)} \frac{q(x_{t-1}|x_0)}{q(x_t|x_0)} - \log \frac{p_\theta(x_0|x_1)}{q(x_1|x_0)} \right]$$

$$= E_q \left[ -\log \frac{p(x_T)}{q(x_T|x_0)} - \sum_{1 < t \leq T} \log \frac{p_\theta(x_{t-1}|x_t)}{q(x_{t-1}|x_t, x_0)} - \log p_\theta(x_0|x_1) \right]$$

$$= E_q [D_{KL}(q(x_T|x_0) || p(x_T)) + \sum_{1 < t \leq T} D_{KL}(q(x_{t-1}|x_t, x_0) || p_\theta(x_{t-1}|x_t)) - \log p_\theta(x_0|x_1)]$$
- $L_{t-1}$  matches two distributions of  $p_\theta(x_{t-1}|x_t) = N(\mu_\theta(x_t, t), \sigma_t^2 I)$  and  $q(x_{t-1}|x_t, x_0) = N(\tilde{\mu}_t(x_t, x_0), \tilde{\beta}_t I)$ 
  - Which constitutes the second term of the entire loss.




- A single step of diffusion model constitutes

- $$L_{t-1} = E_{x_0 \sim p(x_0), \epsilon \sim N(0, I)} \left[ \frac{\beta_t^2}{2\sigma_t^2 \alpha_t (1 - \bar{\alpha}_t)} \left\| \epsilon - \epsilon_\theta(x_t, t) \right\|^2 \right]$$

$$= E_{x_0 \sim p(x_0), \epsilon \sim N(0, I)} \left[ \frac{\beta_t^2}{2\sigma_t^2 \alpha_t (1 - \bar{\alpha}_t)} \left\| \epsilon - \epsilon_\theta(\sqrt{\bar{\alpha}_t} x_0 + \sqrt{1 - \bar{\alpha}_t} \epsilon, t) \right\|^2 \right]$$

- What you need is a U-net shaped neural network with inputs of  $x_t$  and  $t$ 
  - $x_t$  can be produced from  $x_0$  in the closed form




## Algorithm 1 Training

- 1: **repeat**
- 2:  $\mathbf{x}_0 \sim q(\mathbf{x}_0)$
- 3:  $t \sim \text{Uniform}(\{1, \dots, T\})$
- 4:  $\epsilon \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$
- 5: Take gradient descent step on  $\nabla_\theta \left\| \epsilon - \epsilon_\theta(\sqrt{\bar{\alpha}_t} \mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t} \epsilon, t) \right\|^2$
- 6: **until** converged

- We only modeled the distribution of
  - $p_{\theta}(x_{t-1}|x_t) = N(\mu_{\theta}(x_t, t), \sigma_t^2 I)$
  - $q(x_{t-1}|x_t, x_0) = N(\tilde{\mu}_t(x_t, x_0), \tilde{\beta}_t I)$
- Therefore, there is no jump in the decoding step of the series of Markov chain
  - You will utilize the pattern prediction network to denoise the pattern
  - Step-by-Step with long time

## Algorithm 2 Sampling

```
1:  $\mathbf{x}_T \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ 
2: for  $t = T, \dots, 1$  do
3:    $\mathbf{z} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ 
4:    $\mathbf{x}_{t-1} = \frac{1}{\sqrt{\alpha_t}} \left( \mathbf{x}_t - \frac{1-\alpha_t}{\sqrt{1-\alpha_t}} \epsilon_{\theta}(\mathbf{x}_t, t) \right) + \sigma_t \mathbf{z}$ 
5: end for
6: return  $\mathbf{x}_0$ 
```





- $L = E_q \left[ -\log p(x_T) - \sum_{1 \leq t \leq T} \log \frac{p_\theta(x_{t-1}|x_t)}{q(x_t|x_{t-1})} \right] = E_q \left[ -\log \frac{p(x_T)}{q(x_T|x_0)} - \sum_{1 < t \leq T} \log \frac{p_\theta(x_{t-1}|x_t)}{q(x_{t-1}|x_t, x_0)} - \log p_\theta(x_0|x_1) \right]$ 
  - $q(x_t|x_{t-1})$  vs.  $q(x_{t-1}|x_t, x_0)$ 
    - Variance reduction effect :  $x_0$  is added to reduce the variance by always providing the grounding evidence in the variational distribution
    - $q(x_{t-1}|x_t, x_0) = N(\tilde{\mu}_t(x_t, x_0), \tilde{\beta}_t I)$  provides the closed form solution
      - $\tilde{\mu}_t(x_t, x_0) = \frac{\sqrt{\bar{\alpha}_{t-1}}\beta_t}{1-\bar{\alpha}_t} x_0 + \frac{\sqrt{\bar{\alpha}_t(1-\bar{\alpha}_{t-1})}}{1-\bar{\alpha}_t} x_t$
      - $\tilde{\beta}_t = \frac{1-\bar{\alpha}_{t-1}}{1-\bar{\alpha}_t} \beta_t$
- $L_{t-1} = E_{x_0 \sim p(x_0), \epsilon \sim N(0, I)} \left[ \frac{\beta_t^2}{2\sigma_t^2 \alpha_t (1-\bar{\alpha}_t)} ||\epsilon - \epsilon_\theta(x_t, t)||^2 \right]$ 
  - $L \approx E_{t \sim Unif(0, T)} \left[ E_{x_0 \sim p(x_0), \epsilon \sim N(0, I)} \left[ ||\epsilon - \epsilon_\theta(x_t, t)||^2 \right] \right] \approx \frac{1}{T} \sum_{1 \leq t < T} p(t) E_{x_0 \sim p(x_0), \epsilon \sim N(0, I)} \left[ ||\epsilon - \epsilon_\theta(x_t, t)||^2 \right]$
  - Refelct to the Denoising Score Matching loss : Relation to the Noise Conditioned Score Network (NCSN)
  - $J_D(\theta, \sigma) = E_{p_{data}(\tilde{x}, x)} \left[ ||s_\theta(\tilde{x}; \sigma) - \nabla_{\tilde{x}} \log p(\tilde{x}|x)||^2 \right]$

# PERSPECTIVE FROM NCSN

- Usually, inference task requires the optimization on the PDF
  - $p(\xi; \theta) = \frac{1}{z(\theta)} q(\xi; \theta)$
  - $z(\theta) = \int_{\xi \in R^n} q(\xi; \theta) d\xi$
  - However, the integration of  $z(\theta)$  becomes intractable
  - Therefore, there is a demand to estimate the parameter of non-normalized density models
    - Traditional approach would be MCMC, i.e. Metropolis-Hastings, also special case Gibbs sampling
- Let's say that the score function would be

$$\psi(\xi; \theta) = \begin{pmatrix} \frac{\partial \log p(\xi; \theta)}{\partial \xi_1} \\ \dots \\ \frac{\partial \log p(\xi; \theta)}{\partial \xi_n} \end{pmatrix} = \begin{pmatrix} \psi_1(\xi; \theta) \\ \dots \\ \psi_n(\xi; \theta) \end{pmatrix} = \nabla_{\xi} \log p(\xi; \theta) = \nabla_{\xi} \log q(\xi; \theta)$$

- The merit of  $\nabla_{\xi} \log p(\xi; \theta)$  is its independence from  $Z(\theta)$ 
    - From the log-likelihood and the derivative over  $\xi$
- Finally, we define the matching between the model score and the data score functions
  - $J(\theta) = \frac{1}{2} \int_{\xi \in R^n} p_x(\xi) \|\psi(\xi; \theta) - \psi_x(\xi)\|^2 d\xi$
  - This directly handles the  $q(\xi; \theta)$ , but now we need the data score function of  $\psi_x(\xi)$

- $J(\theta) = \frac{1}{2} \int_{\xi \in R^n} p_x(\xi) \|\psi(\xi; \theta) - \psi_x(\xi)\|^2 d\xi$ 
  - This directly handles the  $q(\xi; \theta)$ , but now we need the data score function of  $\psi_x(\xi)$
  - $$= \int_{\xi \in R^n} p_x(\xi) \left[ \frac{1}{2} \|\psi(\xi; \theta)\|^2 + \frac{1}{2} \|\psi_x(\xi)\|^2 - \psi_x(\xi)^T \psi(\xi; \theta) \right] d\xi$$

$$= \int_{\xi \in R^n} p_x(\xi) \left[ \frac{1}{2} \|\psi(\xi; \theta)\|^2 - \psi_x(\xi)^T \psi(\xi; \theta) \right] d\xi + C$$

$$= \int_{\xi \in R^n} p_x(\xi) \frac{1}{2} \|\psi(\xi; \theta)\|^2 d\xi - \int_{\xi \in R^n} p_x(\xi) \psi_x(\xi)^T \psi(\xi; \theta) d\xi + C$$

$$= \int_{\xi \in R^n} p_x(\xi) \frac{1}{2} \|\psi(\xi; \theta)\|^2 d\xi - \sum_i \int_{\xi \in R^n} p_x(\xi) \psi_{x,i}(\xi) \psi_i(\xi; \theta) d\xi + C$$
  - $$-\int p_x(\xi) \frac{\partial \log p_x(\xi)}{\partial \xi_i} \psi_i(\xi; \theta) d\xi = -\int p_x(\xi) \frac{\partial p_x(\xi)}{p_x(\xi) \partial \xi_i} \psi_i(\xi; \theta) d\xi = -\int \frac{\partial p_x(\xi)}{\partial \xi_i} \psi_i(\xi; \theta) d\xi$$
    - Single dimension case :  $\int p(x)(\log p)'(x)f(x) = \int p(x) \frac{p'(x)}{p(x)} f(x) dx = \int p'(x)f(x) dx = -\int p(x)f'(x) dx$
    - Multi dimension case :  $-\int \frac{\partial p_x(\xi)}{\partial \xi_i} \psi_i(\xi; \theta) d\xi = \int \frac{\partial \psi_i(\xi; \theta)}{\partial \xi_i} p_x(\xi) d\xi$ 
      - Requires some additional proves and assumptions:
        - Hyvärinen, Aapo, and Peter Dayan. "Estimation of non-normalized statistical models by score matching." *Journal of Machine Learning Research* 6.4 (2005).
    - $$= \int_{\xi \in R^n} p_x(\xi) \frac{1}{2} \|\psi(\xi; \theta)\|^2 d\xi + \sum_i \int \frac{\partial \psi_i(\xi; \theta)}{\partial \xi_i} p_x(\xi) d\xi + C = \int_{\xi \in R^n} p_x(\xi) \sum_i [\partial \psi_i(\xi; \theta) + \frac{1}{2} \psi_i(\xi; \theta)^2] d\xi + C$$
  - Monte-Carlo Sampling version of  $J(\theta)$ 
    - $$\tilde{J}(\theta) = \frac{1}{T} \sum_t \sum_i [\partial \psi_i(x_t; \theta) + \frac{1}{2} \psi_i(x_t; \theta)^2]$$
  - In its current form, this is designed to be an inference algorithm for known distributions → what if implicit distribution?

- $J(\theta) = \int_{\xi \in R^n} p_x(\xi) \sum_i [\partial \psi_i(\xi; \theta) + \frac{1}{2} \psi_i(\xi; \theta)^2] d\xi + C = E_{p_x} \left[ \text{tr}(\nabla_x s(x; \theta)) + \frac{1}{2} ||s(x; \theta)||^2 \right] + C$ 
  - Let's say  $\psi$  and  $s$  are the score functions
  - $\tilde{J}(\theta) = \frac{1}{T} \sum_t \sum_i [\partial \psi_i(x_t; \theta) + \frac{1}{2} \psi_i(x_t; \theta)^2] = \frac{1}{N} \sum_{i=1}^n [\text{tr}(\nabla_x s(x_i; \theta)) + \frac{1}{2} ||s(x_i; \theta)||^2]$
  - $\nabla_x s(x_i; \theta) = \nabla_x^2 \log \tilde{p}(x; \theta)$  becomes the Hessian matrix of the modeled density function
  - This causes a problem of the computational complexity
    - $\text{tr}(\nabla_x s(x; \theta))$  is computing the trace value of the Hessian matrix
    - Therefore, we need another trick to reduce the computational complexity
      - Denoising score matching
      - Sliced score matching
        - $E_{p_v} E_{p_d} \left[ v^t \nabla_x s_\theta(x) v + \frac{1}{2} ||s_\theta(x)||^2 \right]$
        - $p_v$  : simple known distribution with limited dimensions
        - $v$  : becomes the projection of the Hessian matrix
- From the reparameterization perspective, the score matching becomes
  - $\nabla_\theta H(q_\theta) = -\nabla_\theta E_{q_\theta(x)} [\log q_\theta(x)] = -\nabla_\theta E_{p(\epsilon)} [\log q_\theta(g_\theta(\epsilon))] = -E_{p(\epsilon)} [\nabla_x \log q_\theta(g_\theta(\epsilon)) \nabla_\theta g_\theta(\epsilon)]$ 
    - $q_\theta(x)$  : target density to model
    - $g_\theta(\epsilon)$  : reparametrization function with error  $\epsilon$  from trivial distribution



- $J(\theta) = \frac{1}{2} \int_{\xi \in R^n} p_x(\xi) \|\psi(\xi; \theta) - \psi_x(\xi)\|^2 d\xi$ 
  - Explicit format of the density in  $q$  :  $\psi(\xi; \theta) = \nabla_{\xi} \log p(\xi; \theta) = \nabla_{\xi} \log q(\xi; \theta)$ 
    - $J(\theta) = \int_{\xi \in R^n} p_x(\xi) \sum_i [\partial \psi_i(\xi; \theta) + \frac{1}{2} \psi_i(\xi; \theta)^2] d\xi$
    - $\tilde{J}(\theta) = \frac{1}{T} \sum_t \sum_i [\partial \psi_i(x_t; \theta) + \frac{1}{2} \psi_i(x_t; \theta)^2]$
- How to setup  $q$  is not determined, yet
  - Let's set  $q_{\sigma}(\tilde{x}, x) = q_{\sigma}(\tilde{x}|x)q_0(x)$ 
    - This becomes a kernel of denoising  $x$  with the parameterized noise of  $\sigma$
  - Subsequently,
- $J_{DSM}(\theta) = E_{q_{\sigma}(x, \tilde{x})} \left[ \frac{1}{2} \left| \psi(\tilde{x}; \theta) - \frac{\partial \log q_{\sigma}(\tilde{x}|x)}{\partial \tilde{x}} \right|^2 \right]$ 
  - $\frac{\partial \log q_{\sigma}(\tilde{x}|x)}{\partial \tilde{x}} = \frac{1}{\sigma^2} (x - \tilde{x})$
- Then, the question becomes how to model  $\psi(\tilde{x}; \theta)$ 
  - Which should be learnable and flexible enough to learn the perturbed data score

- $J_{DSM}(\theta) = E_{q_{\sigma}(x, \tilde{x})} \left[ \frac{1}{2} \left| \psi(\tilde{x}; \theta) - \frac{\partial \log q_{\sigma}(\tilde{x}|x)}{\partial \tilde{x}} \right|^2 \right]$ 
  - $\frac{\partial \log q_{\sigma}(\tilde{x}|x)}{\partial \tilde{x}} = \frac{1}{\sigma^2} (x - \tilde{x})$
- $\psi_i(x; \theta) = \frac{\partial \log p(x; \theta)}{\partial x_i} = \frac{1}{\sigma^2} (W^T \text{sigmoid}(Wx) - x)$ 
  - A single-layered denoising autoencoder structure without any constant terms or intercepts
- $J_{DSM}(\theta) = E_{q_{\sigma}(x, \tilde{x})} \left[ \frac{1}{2} \left| \frac{1}{\sigma^2} (W^T \text{sigmoid}(W\tilde{x}) - \tilde{x}) - \frac{1}{\sigma^2} (x - \tilde{x}) \right|^2 \right]$ 

$$= \frac{1}{2} \frac{1}{\sigma^4} E_{q_{\sigma}(x, \tilde{x})} [|W^T \text{sigmoid}(W\tilde{x}) - \tilde{x}) - (x - \tilde{x})|^2]$$

$$= \frac{1}{2} \frac{1}{\sigma^4} E_{q_{\sigma}(x, \tilde{x})} [|W^T \text{sigmoid}(W\tilde{x}) - x|^2] = \frac{1}{2} \frac{1}{\sigma^4} J_{DAE}(\theta)$$
  - $J_{DSM}$  becomes the loss function of the denoising autoencoder
- This shows that we can utilize the denoising score matching
  - As a neural network learning function in the generative model
  - Also, the independence between the noise dimension removes the trace of the Hessian matrix

- $\frac{\partial \log q_\sigma(\tilde{x}|x)}{\partial \tilde{x}} = \frac{1}{\sigma^2} (x - \tilde{x})$ 
  - This means the transition between  $x$  and  $\tilde{x}$
  - Ideally,  $\sigma$  can be very small, so learned  $\psi_i(x; \theta)$  can well approximate the denoising autoencoder
    - However, this means that we cannot assume  $\tilde{x}$  to be the latent space, where we can freely sample
- Therefore, the score-matching needs to be chained
  - How to generate a data instance from a continuous chain
  - Moreover, the current chaining requires a stochastic element of the perturbation  
→ Continuous simulation with perturbation
- Langevin method
  - $\tilde{x}_t = \tilde{x}_{t-1} + \frac{\epsilon}{2} \nabla_x \log p(\tilde{x}_{t-1}) + \sqrt{\epsilon} z_t$ 
    - $z_t \sim N(0, I)$
    - $\psi_i(x; \theta) = \frac{\partial \log p(x; \theta)}{\partial x_i}$
  - Which can be learned from the denoising error
- Finally, the question becomes how to chain  $q_\sigma(\tilde{x}|x)$  : particularly, how to setup  $\sigma$

- How to setup  $\sigma$

- Let's set  $\{\sigma_i\}_{i=1}^L$  to satisfy  $\frac{\sigma_1}{\sigma_2} = \dots = \frac{\sigma_{L-1}}{\sigma_L} > 1$

- How to setup  $q_\sigma$

- Following the previous setup,  $\frac{\partial \log q_\sigma(\tilde{x}|x)}{\partial \tilde{x}} = \frac{1}{\sigma^2} (x - \tilde{x}) = -\frac{1}{\sigma^2} (\tilde{x} - x)$
- $q_\sigma(x) = \int p_d(t) N(x|t, \sigma^2 I) dt$

- Loss structure of NCSN

- $L_l(\theta; \sigma_l) = \frac{1}{2} E_{p_d} E_{\tilde{x} \sim N(x, \sigma_l^2 I)} \left[ \left\| s_\theta(\tilde{x}, \sigma_l) + \frac{\tilde{x} - x}{\sigma_l^2} \right\|^2 \right]$  : This is designed to be a single step

- $L(\theta; \{\sigma_i\}_{i=1}^L) = \frac{1}{L} \sum_{l=1}^L \lambda(\sigma_l) \frac{1}{2} E_{p_d} E_{\tilde{x} \sim N(x, \sigma_l^2 I)} \left[ \left\| s_\theta(\tilde{x}, \sigma_l) + \frac{\tilde{x} - x}{\sigma_l^2} \right\|^2 \right]$

- $\lambda(\sigma_l)$  : coefficient function, i.e.  $\lambda(\sigma) = \sigma^2$

- Under this choice :  $L(\theta; \{\sigma_i\}_{i=1}^L) = \frac{1}{L} \sum_{l=1}^L \frac{1}{2} E_{p_d} E_{\tilde{x} \sim N(x, \sigma_l^2 I)} \left[ \left\| \sigma s_\theta(\tilde{x}, \sigma_l) + \frac{\tilde{x} - x}{\sigma_l^2} \right\|^2 \right]$

- $\frac{\tilde{x} - x}{\sigma_l} \sim N(0, I)$
- $s_\theta(\tilde{x}, \sigma_l) \propto \frac{1}{\sigma}$ , if  $s_\theta$  is well trained

- Langevin dynamics is used
  - $\tilde{x}_t = \tilde{x}_{t-1} + \frac{\epsilon}{2} \nabla_x \log p(\tilde{x}_{t-1}) + \sqrt{\epsilon} z_t$ 
    - $z_t \sim N(0, I)$
    - $\psi_i(x; \theta) = \frac{\partial \log p(x; \theta)}{\partial x_i}$
  - Given the denoising function :  $\psi_i(x; \theta) = s_\theta(x, \sigma)$
- $\alpha_i$  : the Langevin dynamics step magnitude
  - $\{\sigma_i\}_{i=1}^L$  to satisfy  $\frac{\sigma_1}{\sigma_2} = \dots = \frac{\sigma_{L-1}}{\sigma_L} > 1$ 
    - $\sigma_1 > \sigma_2 \dots > \sigma_L$
    - Given the current convention
      - $\tilde{x}_0$  is the latent sample of noised
      - $\tilde{x}_T$  is the denoised sample
    - In most noised samples (pure latent of  $N(0, I)$ ),  $\sigma$  will be high
    - Near the data distribution,  $\sigma$  will be low
  - $\alpha_i = \epsilon \frac{\sigma_i^2}{\sigma_L^2}$ 
    - In most noise samples, the  $\alpha_i$  step will be large

---

**Algorithm 1** Annealed Langevin dynamics.
 

---

**Require:**  $\{\sigma_i\}_{i=1}^L, \epsilon, T$ .

```

1: Initialize  $\tilde{x}_0$ 
2: for  $i \leftarrow 1$  to  $L$  do
3:    $\alpha_i \leftarrow \epsilon \cdot \sigma_i^2 / \sigma_L^2$   $\triangleright \alpha_i$  is the step size.
4:   for  $t \leftarrow 1$  to  $T$  do
5:     Draw  $z_t \sim \mathcal{N}(0, I)$ 
6:      $\tilde{x}_t \leftarrow \tilde{x}_{t-1} + \frac{\alpha_i}{2} s_\theta(\tilde{x}_{t-1}, \sigma_i) + \sqrt{\alpha_i} z_t$ 
7:   end for
8:    $\tilde{x}_0 \leftarrow \tilde{x}_T$ 
9: end for
return  $\tilde{x}_T$ 
    
```

---

- Loss structure

- DDPM :  $L(\theta) = \frac{1}{T} \sum_{1 \leq t < T} p(t) E_{x_0 \sim p(x_0), \epsilon \sim N(0, I)} \left[ \left\| \epsilon - \epsilon_\theta(x_t, t) \right\|^2 \right]$

- NCSN :  $L(\theta; \{\sigma_i\}_{i=1}^L) = \frac{1}{L} \sum_{l=1}^L \lambda(\sigma_l) \frac{1}{2} E_{p_d} E_{\tilde{x} \sim N(x, \sigma_l^2 I)} \left[ \left\| s_\theta(\tilde{x}, \sigma_l) + \frac{\tilde{x} - x}{\sigma_l^2} \right\|^2 \right]$

- Generation algorithm

- $\mu_\theta(x_t, t) = \frac{1}{\sqrt{\alpha_t}} \left( x_t - \frac{\beta_t}{\sqrt{(1-\alpha_t)}} \epsilon_\theta(x_t, t) \right)$

- $\tilde{x}_t = \tilde{x}_{t-1} + \frac{\epsilon}{2} \nabla_x \log p(\tilde{x}_{t-1}) + \sqrt{\epsilon} z_t$

## Algorithm 2 Sampling

```

1:  $\mathbf{x}_T \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ 
2: for  $t = T, \dots, 1$  do
3:    $\mathbf{z} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ 
4:    $\mathbf{x}_{t-1} = \frac{1}{\sqrt{\alpha_t}} \left( \mathbf{x}_t - \frac{1-\alpha_t}{\sqrt{1-\alpha_t}} \epsilon_\theta(\mathbf{x}_t, t) \right) + \sigma_t \mathbf{z}$ 
5: end for
6: return  $\mathbf{x}_0$ 
    
```

## Algorithm 1 Annealed Langevin dynamics.


**Require:**  $\{\sigma_i\}_{i=1}^L, \epsilon, T$ .

```

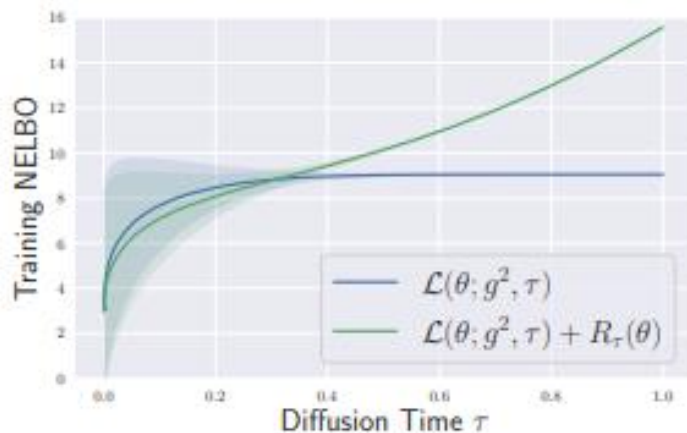
1: Initialize  $\tilde{\mathbf{x}}_0$ 
2: for  $i \leftarrow 1$  to  $L$  do
3:    $\alpha_i \leftarrow \epsilon \cdot \sigma_i^2 / \sigma_L^2$   $\triangleright \alpha_i$  is the step size.
4:   for  $t \leftarrow 1$  to  $T$  do
5:     Draw  $\mathbf{z}_t \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ 
6:      $\tilde{\mathbf{x}}_t \leftarrow \tilde{\mathbf{x}}_{t-1} + \frac{\alpha_i}{2} \mathbf{s}_\theta(\tilde{\mathbf{x}}_{t-1}, \sigma_i) + \sqrt{\alpha_i} \mathbf{z}_t$ 
7:   end for
8:    $\tilde{\mathbf{x}}_0 \leftarrow \tilde{\mathbf{x}}_T$ 
9: end for
return  $\tilde{\mathbf{x}}_T$ 
    
```




# Criticality of Noise Scheduling



(a) Integrand by Time



(b) Variational Bound Truncated at  $\tau$

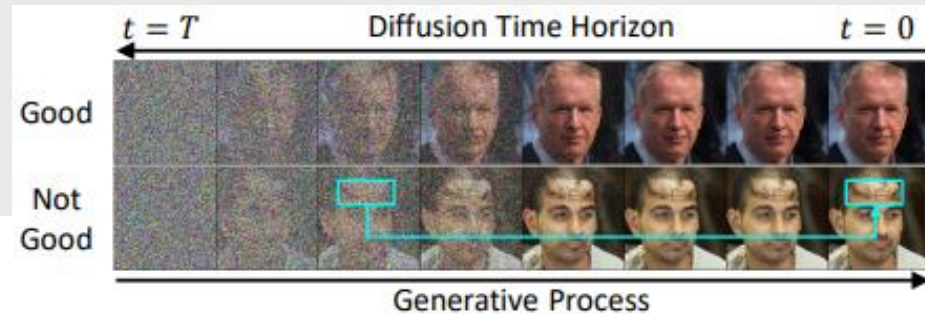



(c) Test Performance by Log-Time

- Small diffusion time dominates the loss term
  - $\frac{\partial \log q_\sigma(\tilde{x}|x)}{\partial \tilde{x}} = \frac{1}{\sigma^2} (x - \tilde{x})$ , Very small  $\sigma$
  - The matching between scores become large magnitude
- What happens when  $\sigma$  is either very small/large?
- Possible remedy?
  - Magnitude vs. # of samplings

$$\begin{aligned} \mathcal{L}(\theta; g^2, \epsilon) &= \frac{Z_\epsilon}{2} \int_\epsilon^T p_{iw}(t) \sigma^2(t) \mathbb{E}[\|s_\theta(\mathbf{x}_t, t) - \nabla \log p_{0t}(\mathbf{x}_t | \mathbf{x}_0)\|_2^2] dt \\ &\approx \frac{Z_\epsilon}{2B} \sum_{b=1}^B \sigma^2(t_{iw}^{(b)}) \left\| s_\theta(\mathbf{x}_{t_{iw}^{(b)}}, t_{iw}^{(b)}) - \frac{\epsilon^{(b)}}{\sigma(t_{iw}^{(b)})} \right\|_2^2 \end{aligned}$$

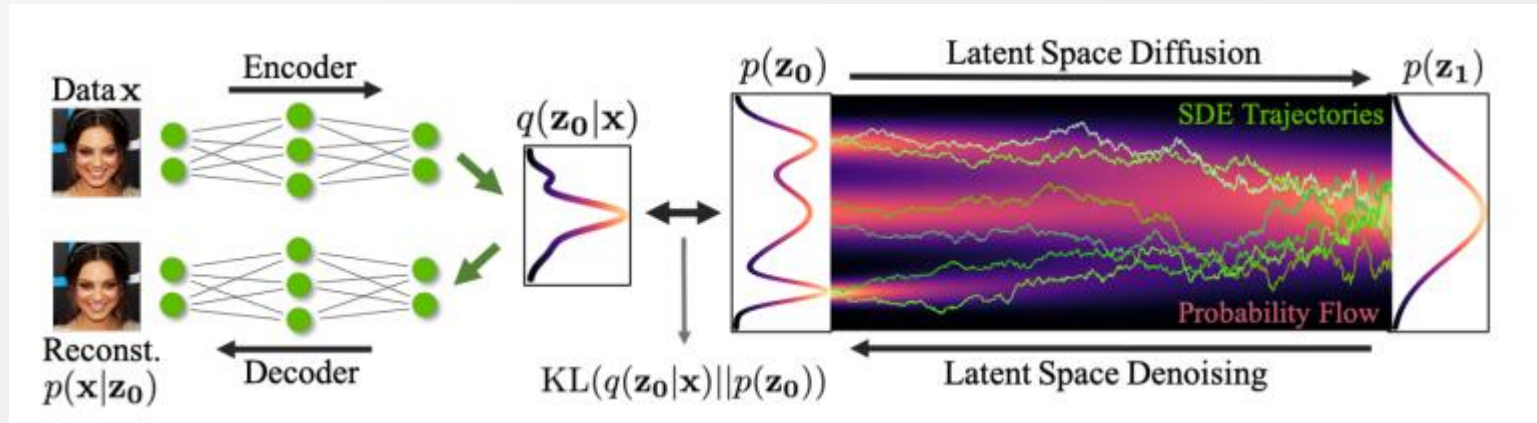
where  $\{t_{iw}^{(b)}\}_{b=1}^B$  is the Monte-Carlo sample from the importance distribution, i.e.,  $t_{iw}^{(b)} \sim p_{iw}(t) \propto \frac{g^2(t)}{\sigma^2(t)}$ .



- Linear encodings of DDPM and NCSN
  - NCSN :  $\frac{\partial \log q_{\sigma}(\tilde{x}|x)}{\partial \tilde{x}} = \frac{1}{\sigma^2} (x - \tilde{x})$
  - DDPM :  $q(x_t|x_{t-1}) = N(x_t; \sqrt{1 - \beta_t}x_{t-1}, \beta_t I)$
- Why linear encoding?
  - NCSN : the trace of the Hessian matrix
  - DDPM : the closed-form solution of the perturbed features
- How to make the linear encoding to be non-linear
  - Make the feature to be transformed to the embedding through nonlinear encoding



- Very simple idea
  - Add a VAE structure
  - Define diffusion model upon  $z$ , not  $x$
- Natural expansion, given the ELBO structure of diffusion models



$$\begin{aligned}
 \mathcal{L}(\mathbf{x}, \phi, \theta, \psi) &= \mathbb{E}_{q_{\phi}(\mathbf{z}_0|\mathbf{x})} [-\log p_{\psi}(\mathbf{x}|\mathbf{z}_0)] + \text{KL}(q_{\phi}(\mathbf{z}_0|\mathbf{x})||p_{\theta}(\mathbf{z}_0)) \\
 &= \underbrace{\mathbb{E}_{q_{\phi}(\mathbf{z}_0|\mathbf{x})} [-\log p_{\psi}(\mathbf{x}|\mathbf{z}_0)]}_{\text{reconstruction term}} + \underbrace{\mathbb{E}_{q_{\phi}(\mathbf{z}_0|\mathbf{x})} [\log q_{\phi}(\mathbf{z}_0|\mathbf{x})]}_{\text{negative encoder entropy}} + \underbrace{\mathbb{E}_{q_{\phi}(\mathbf{z}_0|\mathbf{x})} [-\log p_{\theta}(\mathbf{z}_0)]}_{\text{cross entropy}}
 \end{aligned}$$

**Theorem 1.** Given two distributions  $q(\mathbf{z}_0|\mathbf{x})$  and  $p(\mathbf{z}_0)$ , defined in the continuous space  $\mathbb{R}^D$ , denote the marginal distributions of diffused samples under the SDE in Eq. 1 at time  $t$  with  $q(\mathbf{z}_t|\mathbf{x})$  and  $p(\mathbf{z}_t)$ . Assuming mild smoothness conditions on  $\log q(\mathbf{z}_t|\mathbf{x})$  and  $\log p(\mathbf{z}_t)$ , the cross entropy is:

$$CE(q(\mathbf{z}_0|\mathbf{x})||p(\mathbf{z}_0)) = \mathbb{E}_{t \sim \mathcal{U}[0,1]} \left[ \frac{g(t)^2}{2} \mathbb{E}_{q(\mathbf{z}_t, \mathbf{z}_0|\mathbf{x})} [\|\nabla_{\mathbf{z}_t} \log q(\mathbf{z}_t|\mathbf{x}) - \nabla_{\mathbf{z}_t} \log p(\mathbf{z}_t)\|_2^2] \right] + \frac{D}{2} \log(2\pi e \sigma_0^2).$$

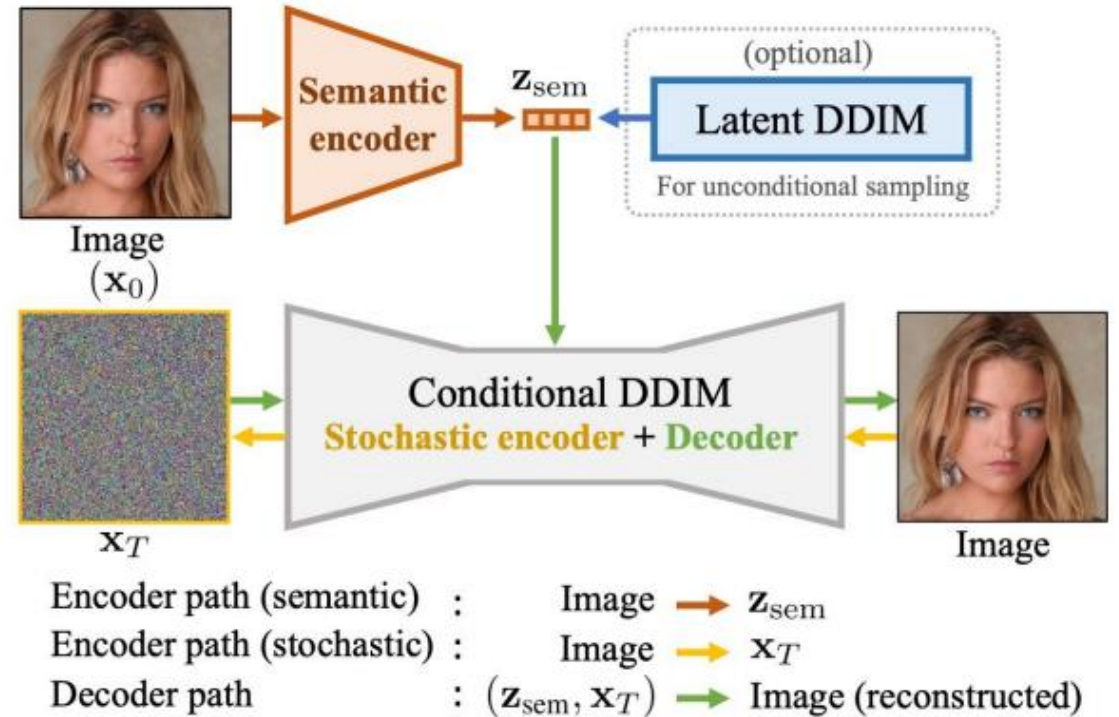


- DDPM

- $$L(\theta) = \frac{1}{T} \sum_{1 \leq t < T} p(t) E_{x_0 \sim p(x_0), \epsilon \sim N(0, I)} \left[ \|\epsilon - \epsilon_\theta(x_t, t)\|^2 \right]$$

- Utilizing the conditional structure of DDIM

- Mixing the semantic latent variable  $z_{sem}$
- Error pattern estimation is updated to anticipate both time and semantics




$$p_\theta(x_{0:T} | z_{sem}) = p(x_T) \prod_{t=1}^T p_\theta(x_{t-1} | x_t, z_{sem}) \quad (3)$$

$$p_\theta(x_{t-1} | x_t, z_{sem}) = \begin{cases} \mathcal{N}(f_\theta(x_1, 1, z_{sem}), \mathbf{0}) & \text{if } t = 1 \\ q(x_{t-1} | x_t, f_\theta(x_t, t, z_{sem})) & \text{otherwise} \end{cases} \quad (4)$$

$$x_{t+1} = \sqrt{\alpha_{t+1}} f_\theta(x_t, t, z_{sem}) + \sqrt{1 - \alpha_{t+1}} \epsilon_\theta(x_t, t, z_{sem})$$

$$L_{\text{latent}} = \sum_{t=1}^T \mathbb{E}_{z_{sem}, \epsilon_t} \left[ \|\epsilon_\omega(z_{sem}, t, t) - \epsilon_t\|_1 \right]$$



- Latent space diffusion
  - $dz_t^\phi = -\frac{1}{2}\beta(t)z_t^\phi dt + g(t)d\omega_t$
- Forward “data” diffusion
  - $dx_t^\phi = f_\phi(x_t^\phi, t) dt + G_\phi(x_t^\phi, t) d\omega_t$
- Now, we have a diffusion term with a learning function of  $G_\phi(x_t^\phi, t)$

SDE	$f(x_t, t)$	$G(x_t, t)$	$x_0$	$x_{0.1}$	$x_{0.2}$	$x_{0.3}$	$x_{0.4}$	$x_{0.5}$	$x_{0.6}$	$x_{0.7}$
Linear (VE/VP)	Linear $f(x_t, t) \propto x_t$	Linear $G(x_t, t) = g(t)$								
Nonlinear	Nonlinear	Linear $G(x_t, t) = g(t)$								
	Semi-linear (Rotating)	Nonlinear								