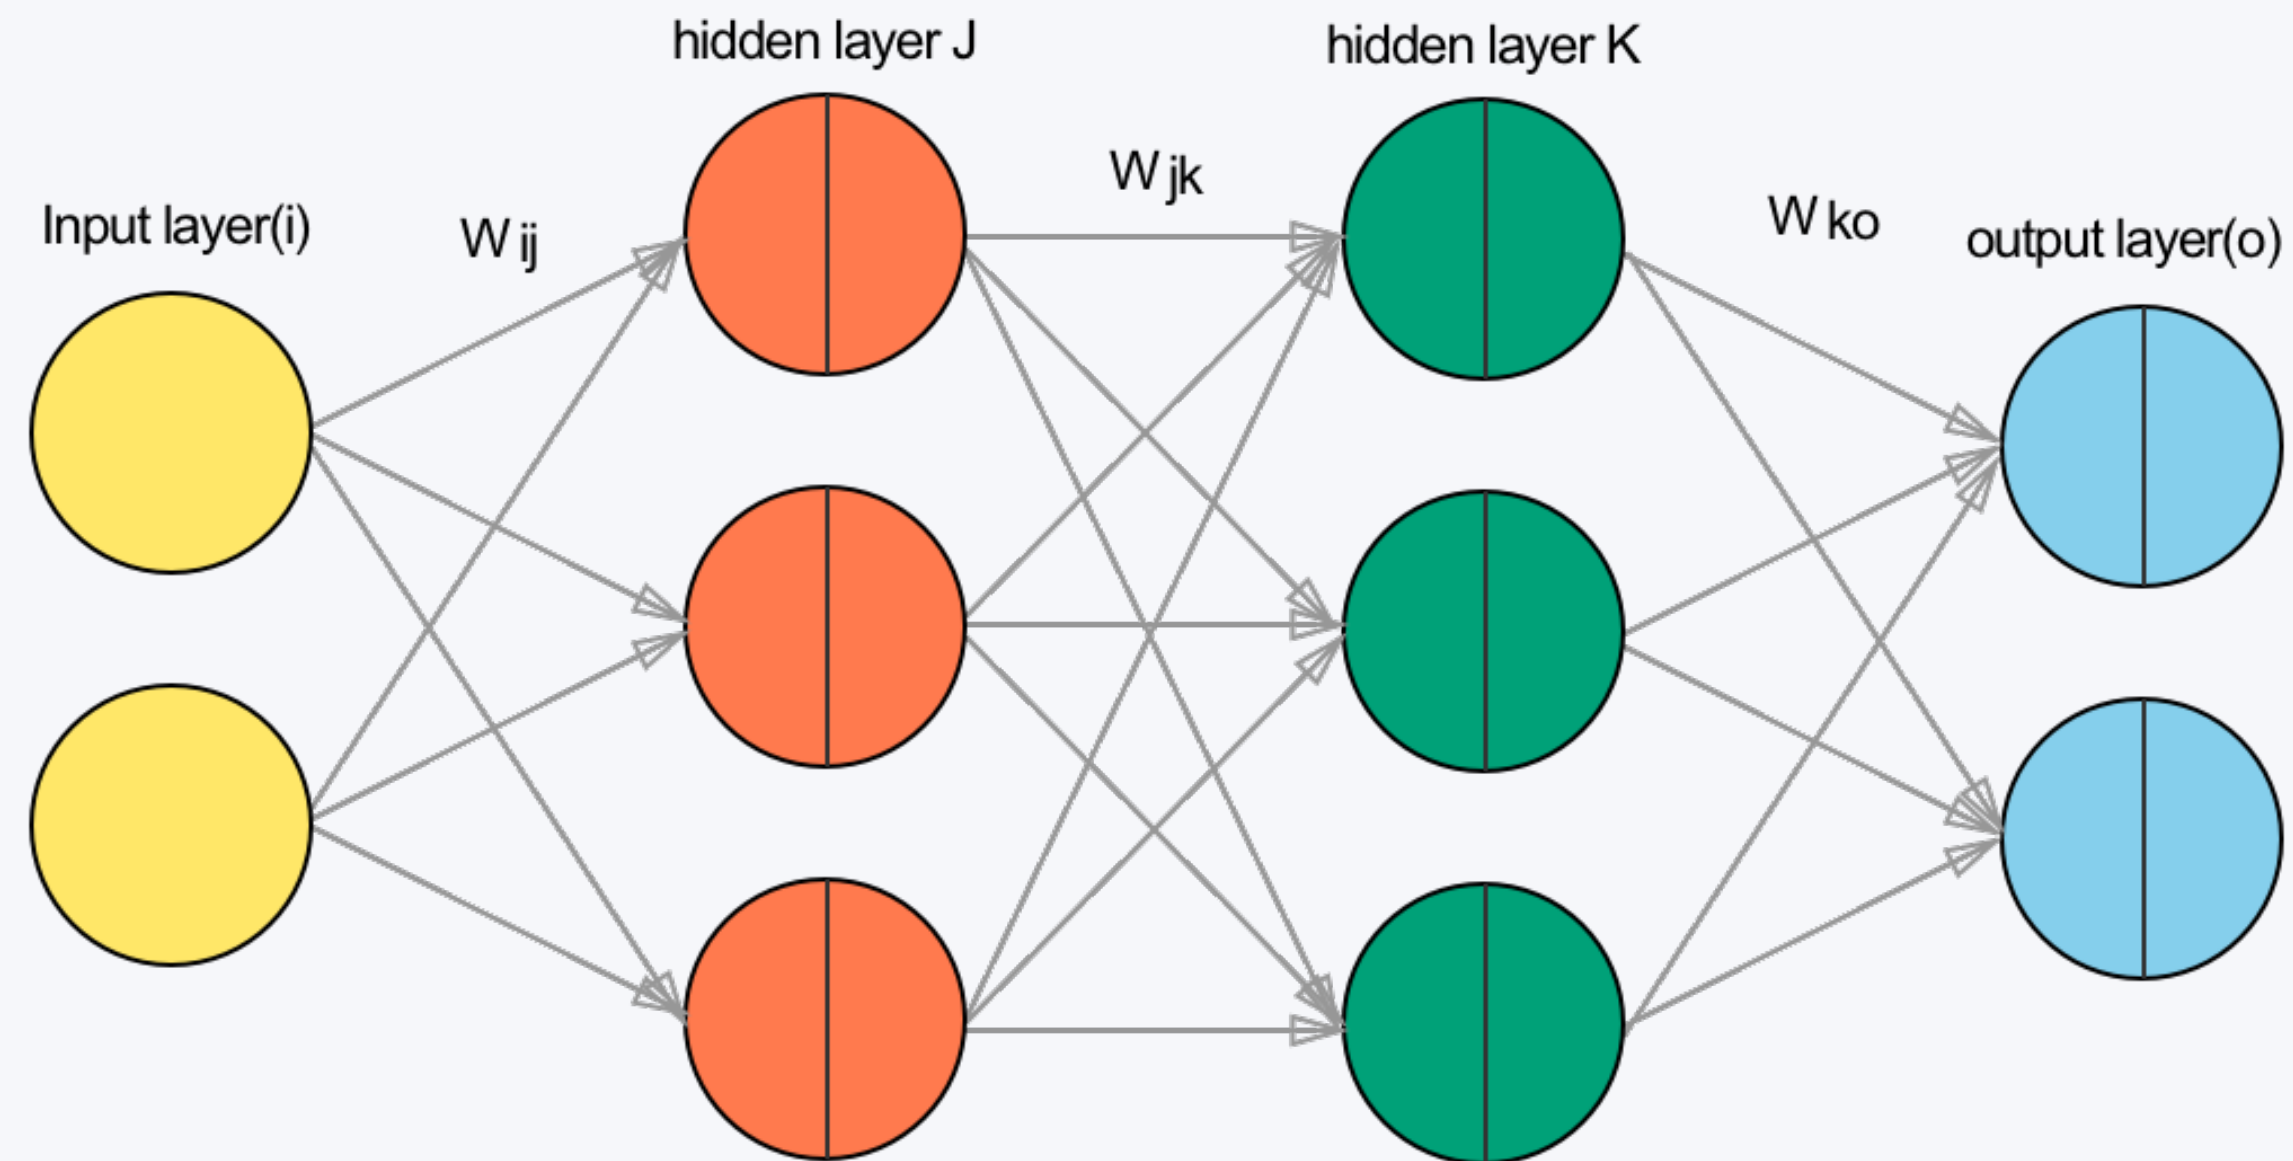


CNN

Convolutional Neural Network

Artificial Neural Network

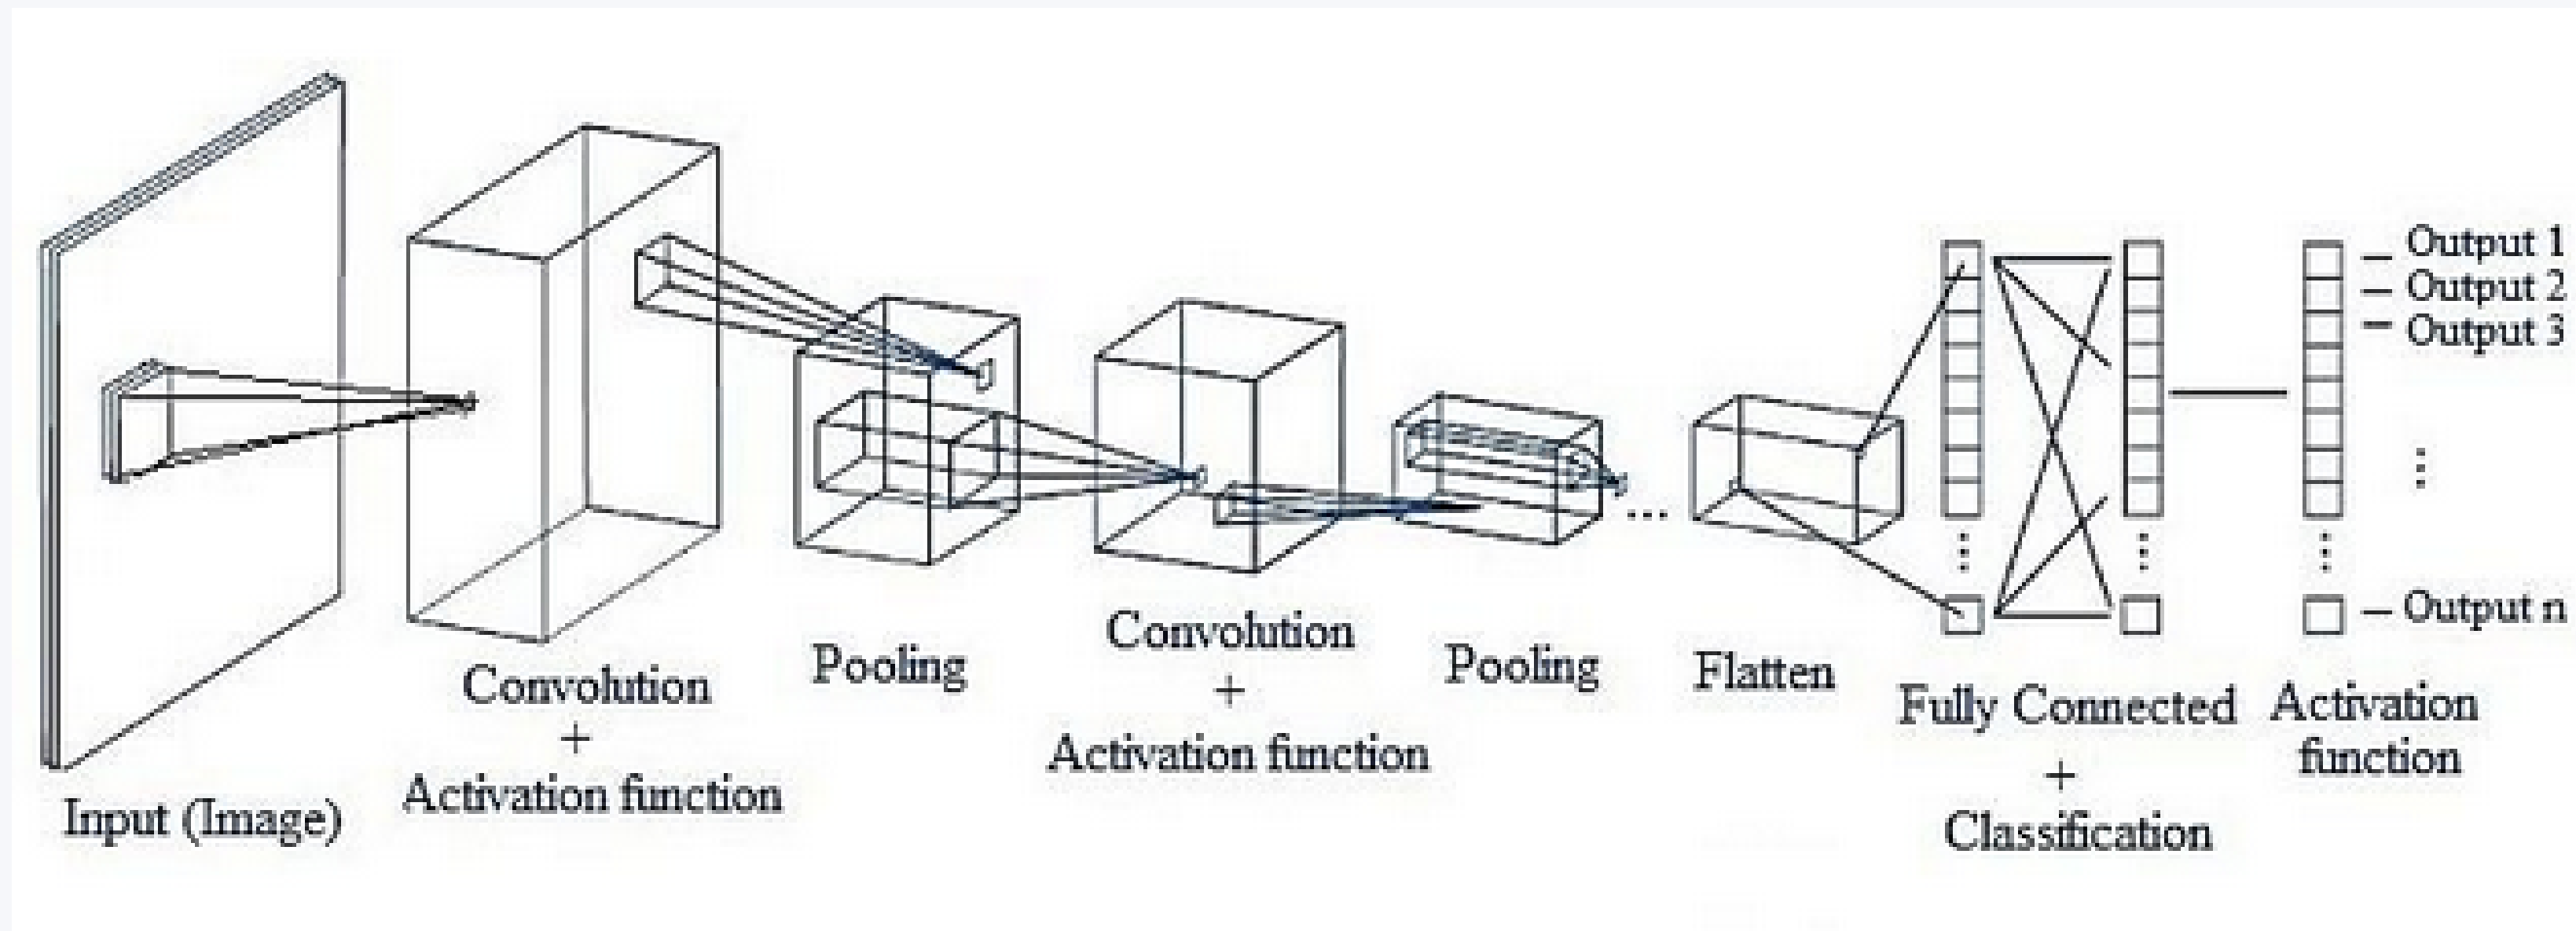


► **Forward propagation**

► **Backward propagation**

CNN

Convolutional Neural Network



► Filter

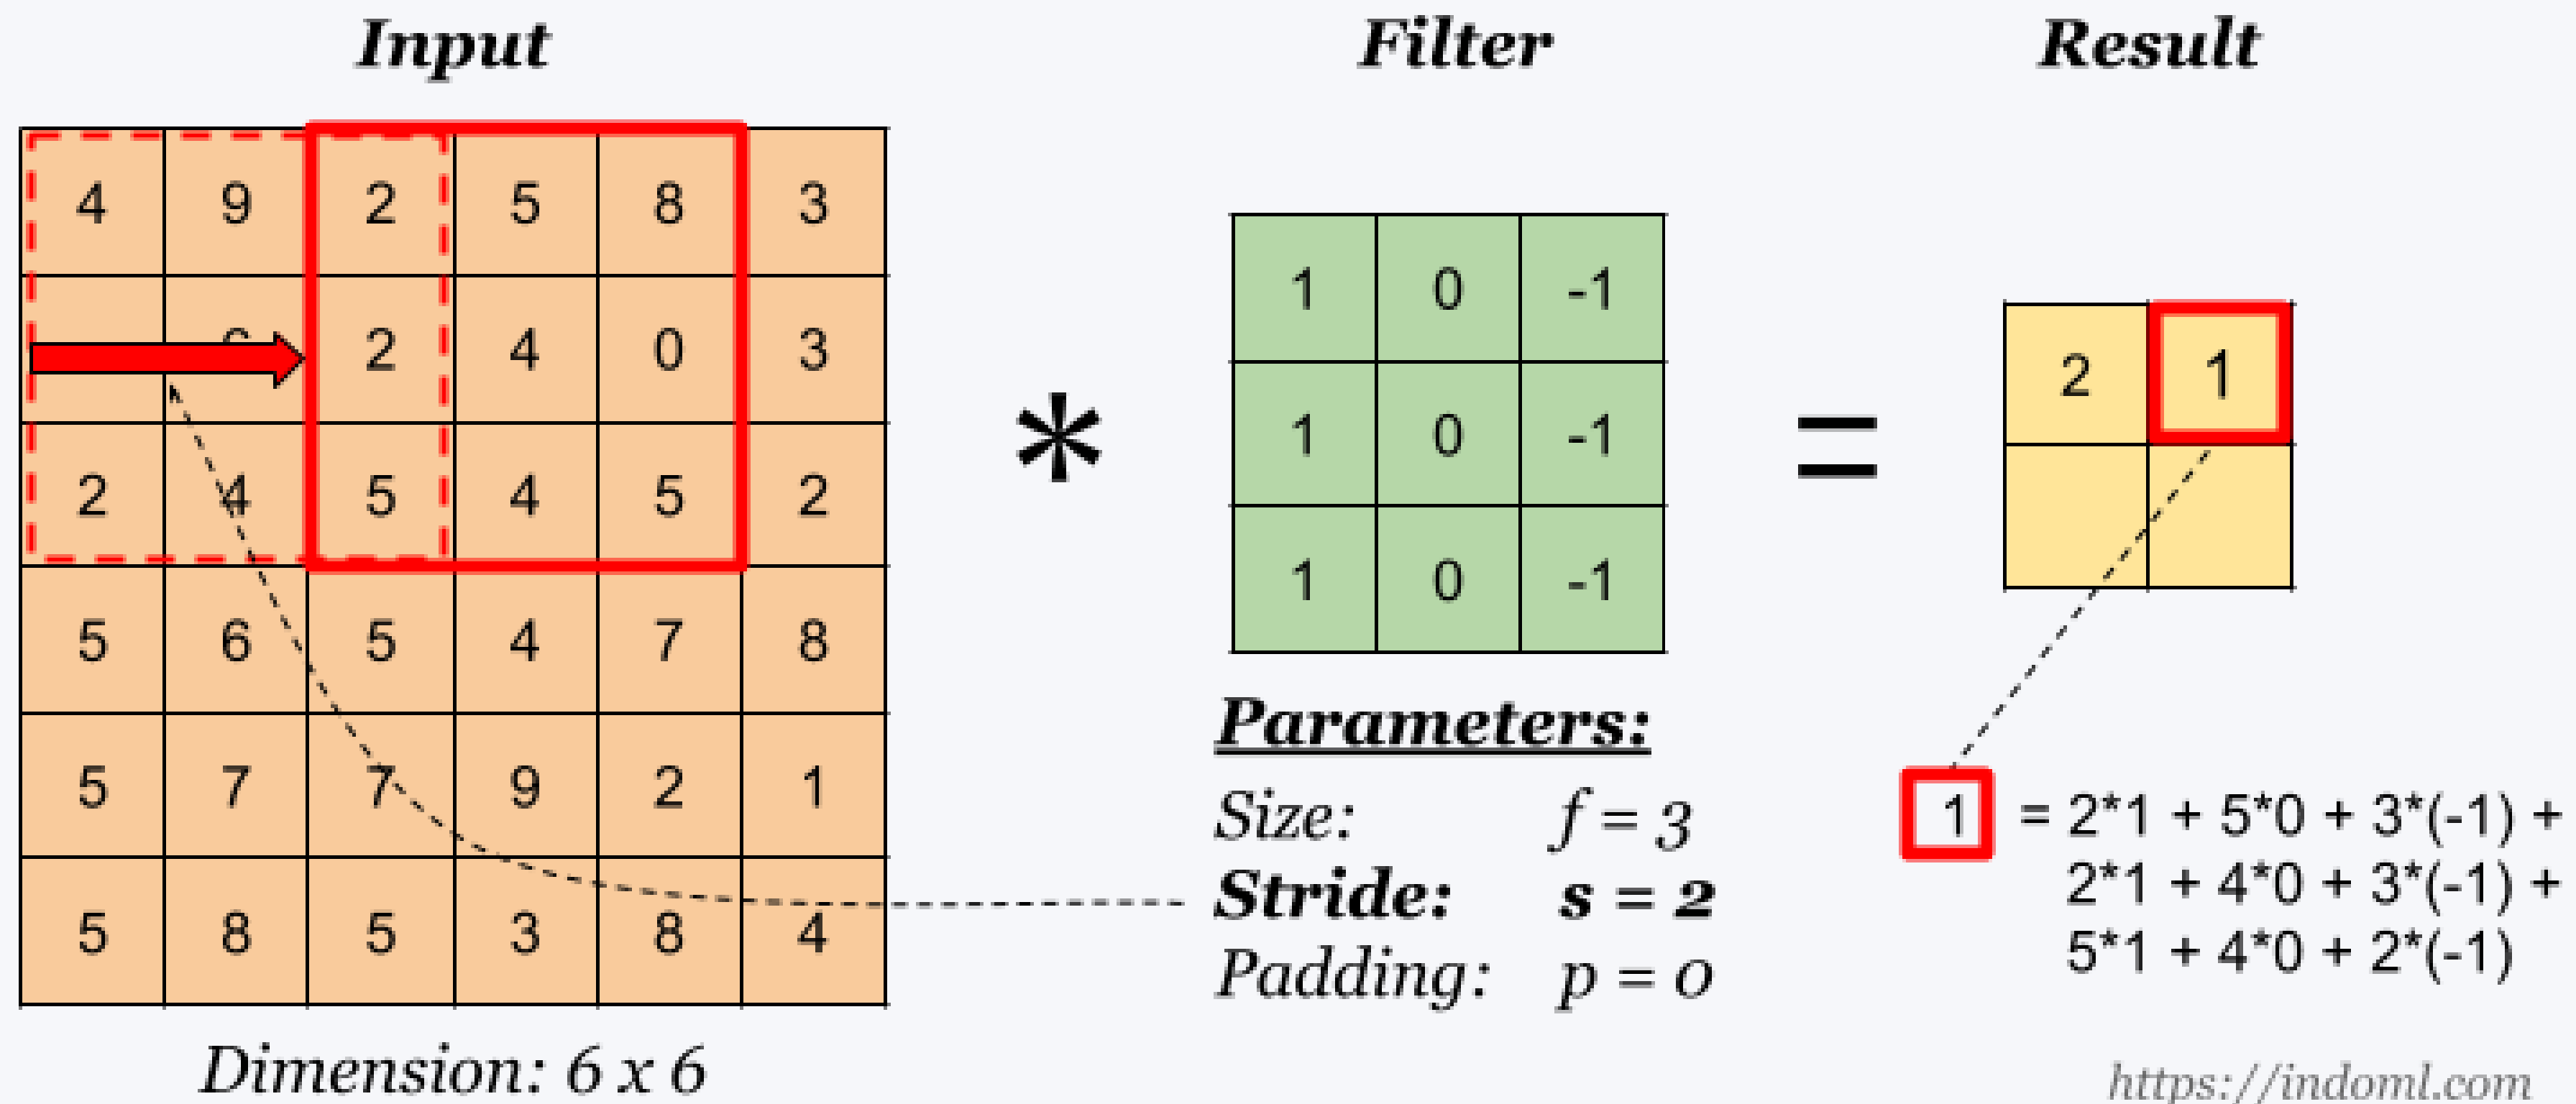
► Padding

► Activation function

► Pooling

CNN (Convolutional Neural Network)

Filter (Kernel)




CNN (Convolutional Neural Network)

Activation function


Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$




tanh

$$\tanh(x)$$




ReLU

$$\max(0, x)$$



Leaky ReLU

$$\max(0.1x, x)$$




Maxout

$$\max(w_1^T x + b_1, w_2^T x + b_2)$$

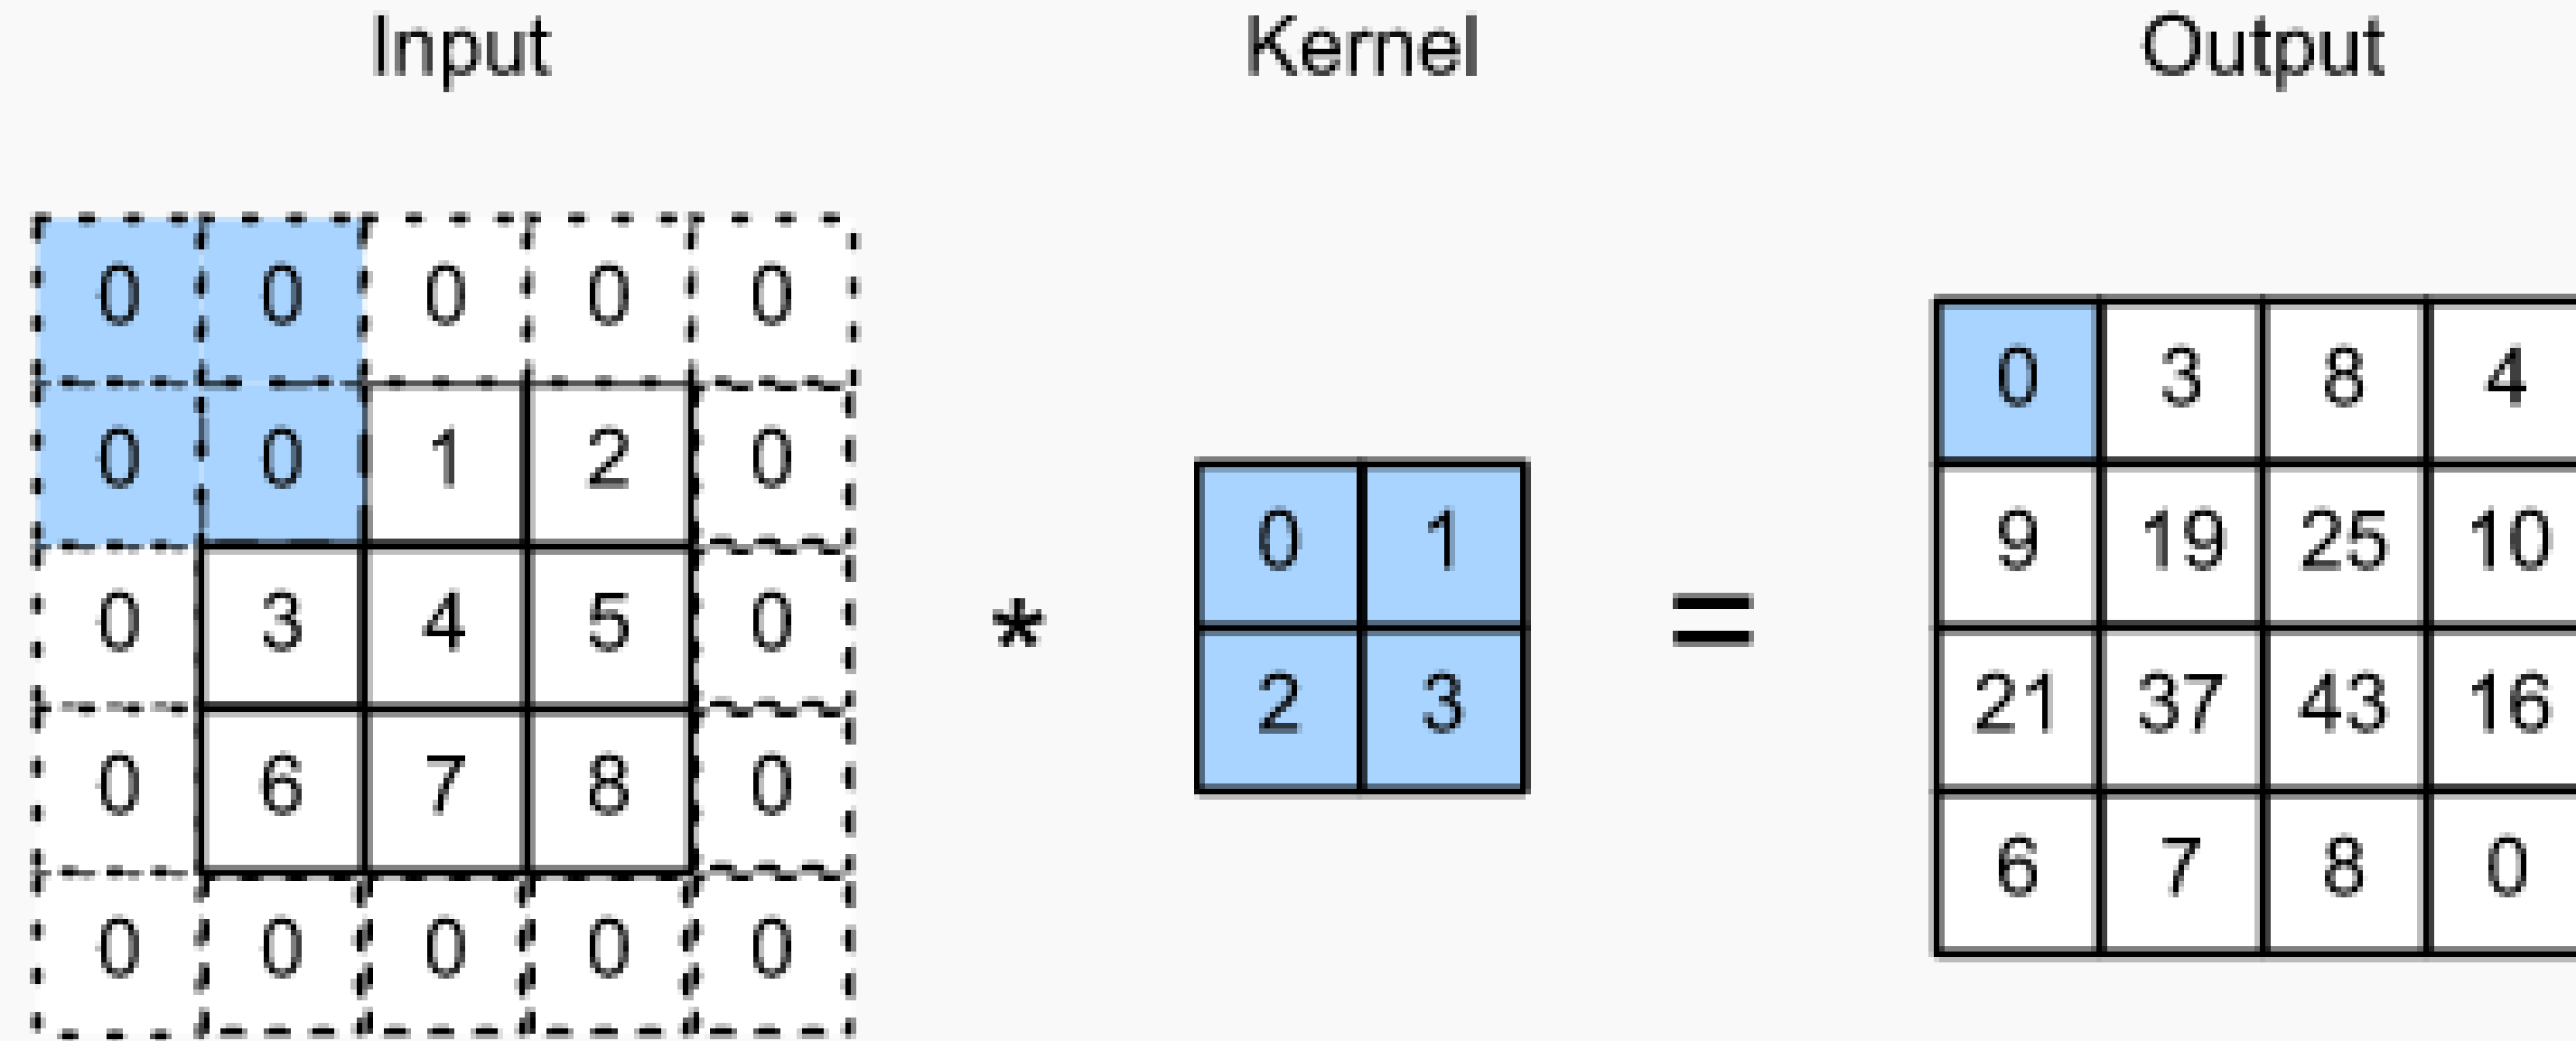
ELU

$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$



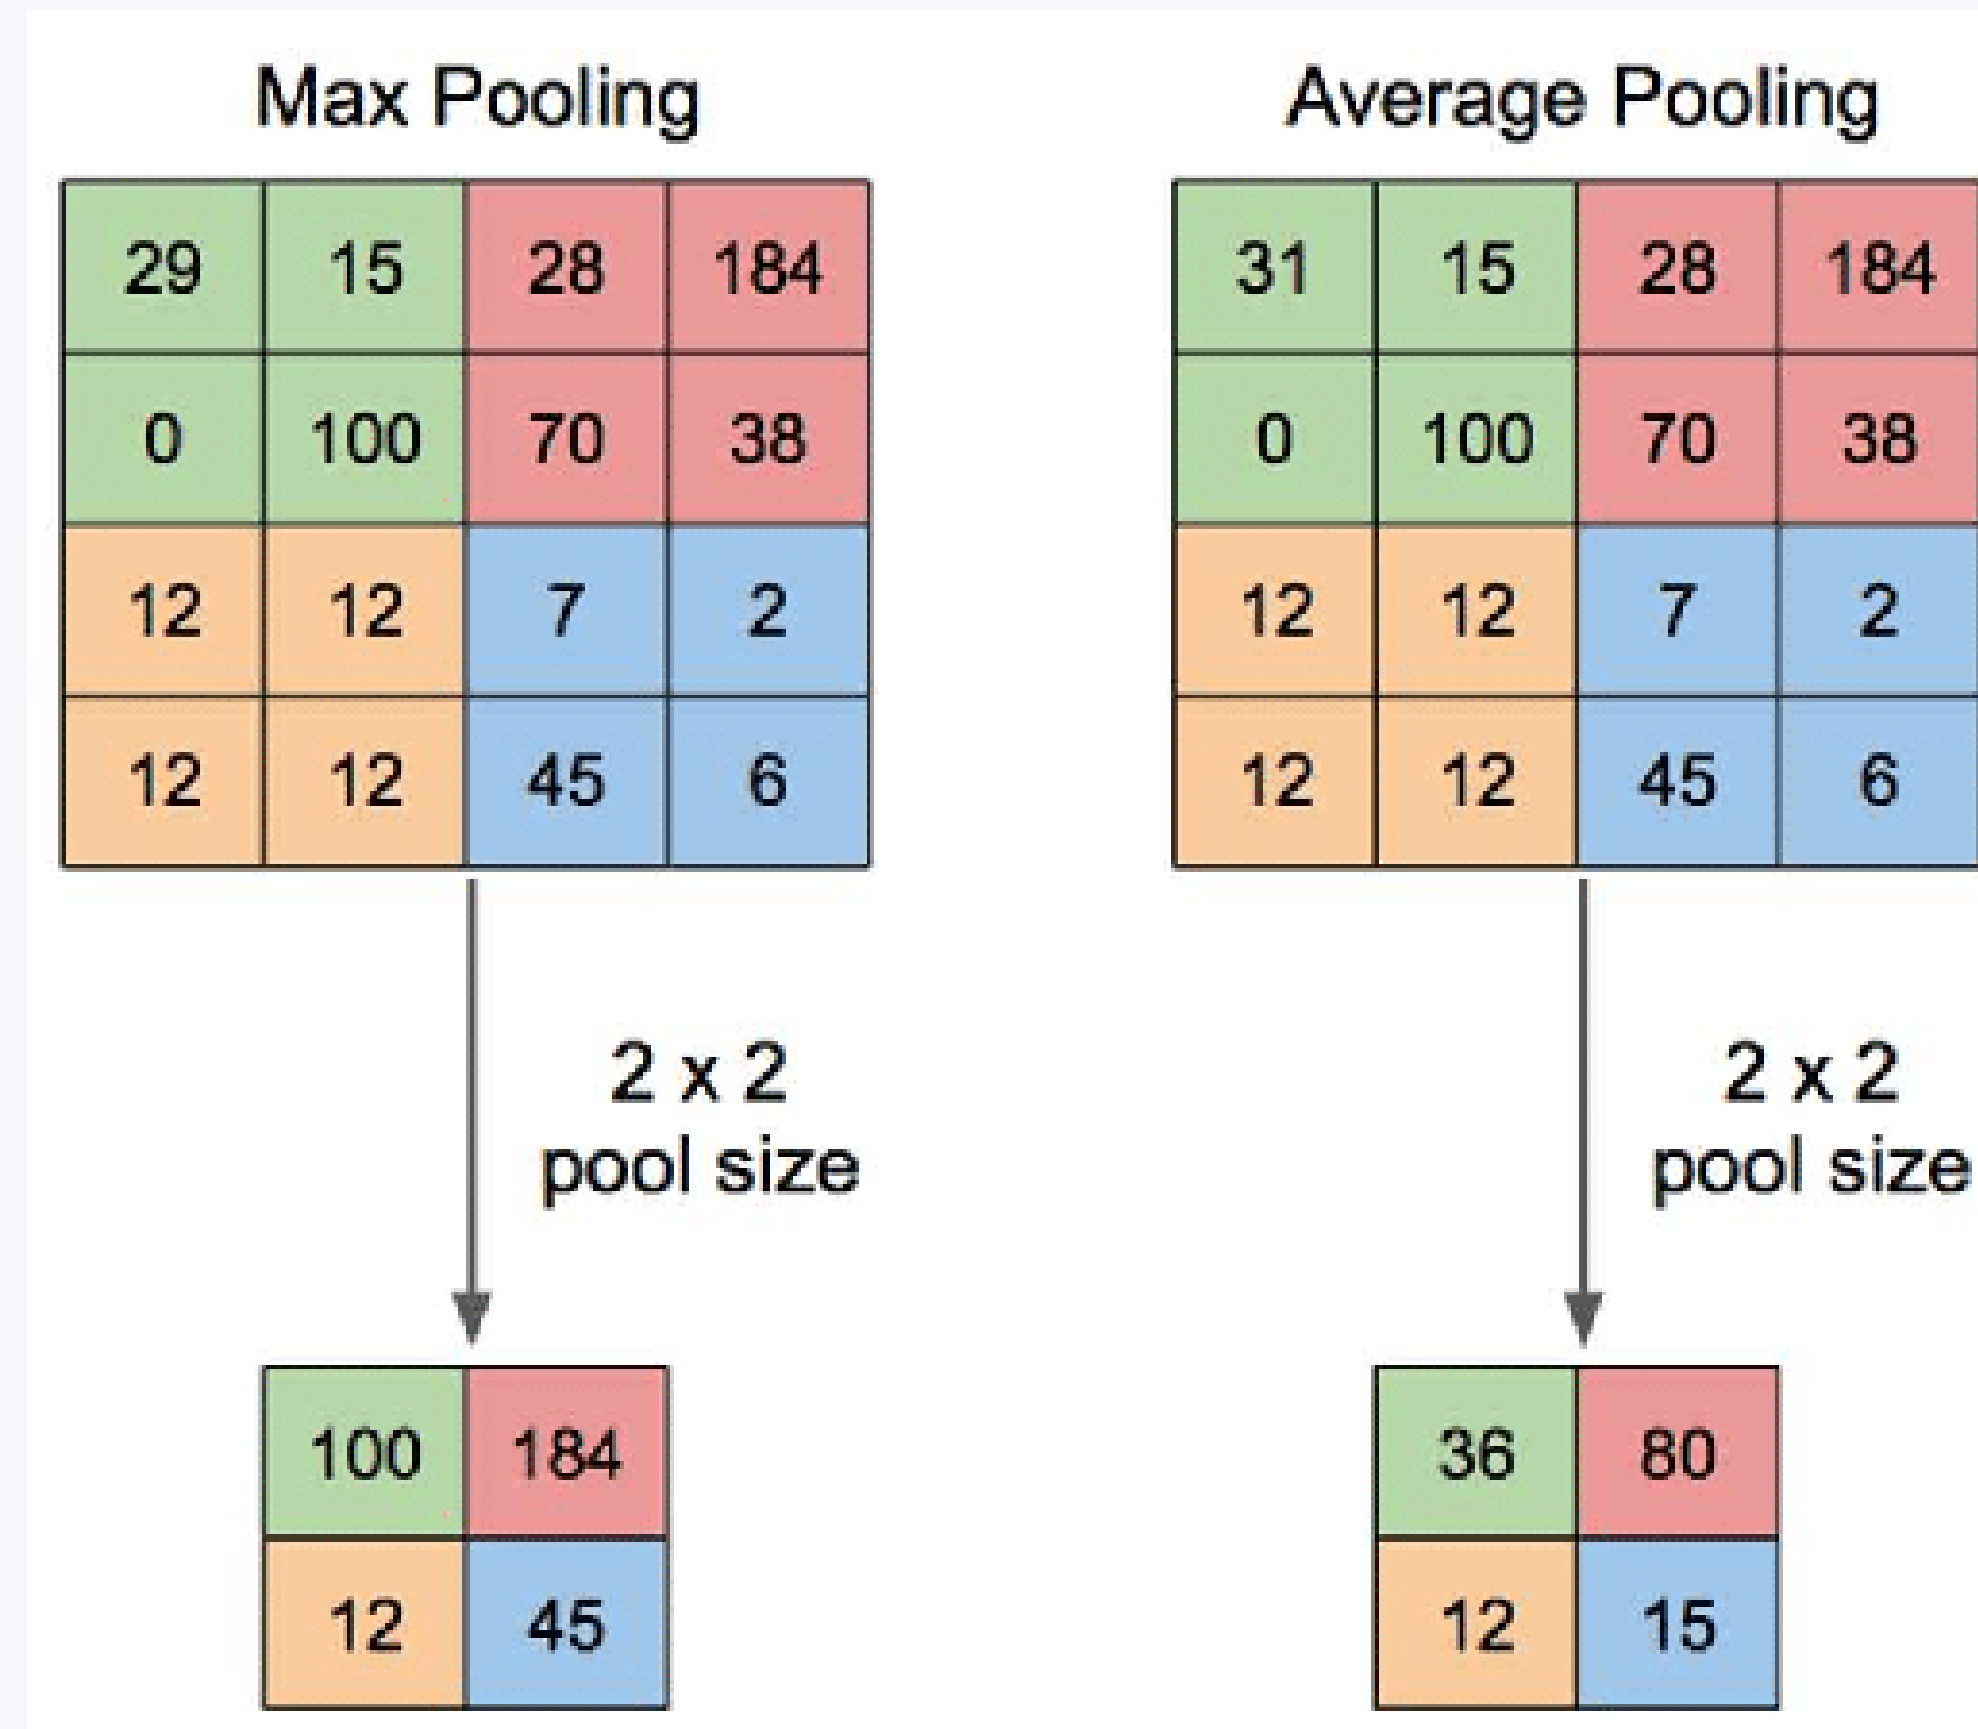
CNN (Convolutional Neural Network)

Padding




CNN (Convolutional Neural Network)


Pooling



Modern Architecture in Deep Learning




AlexNet





AlexNet

1. ReLU Nonlinearity

tanh
 $\tanh(x)$



ReLU
 $\max(0, x)$




2. Local Response Normalization

$$b_{x,y}^i = a_{x,y}^i / \left(k + \alpha \sum_{j=\max(0, i-n/2)}^{\min(N-1, i+n/2)} (a_{x,y}^j)^2 \right)^\beta$$

3. Overlapping Pooling

3x3 pooling with stride=2



► Activation function -ReLU

► LRN

► Overlapping Pooling

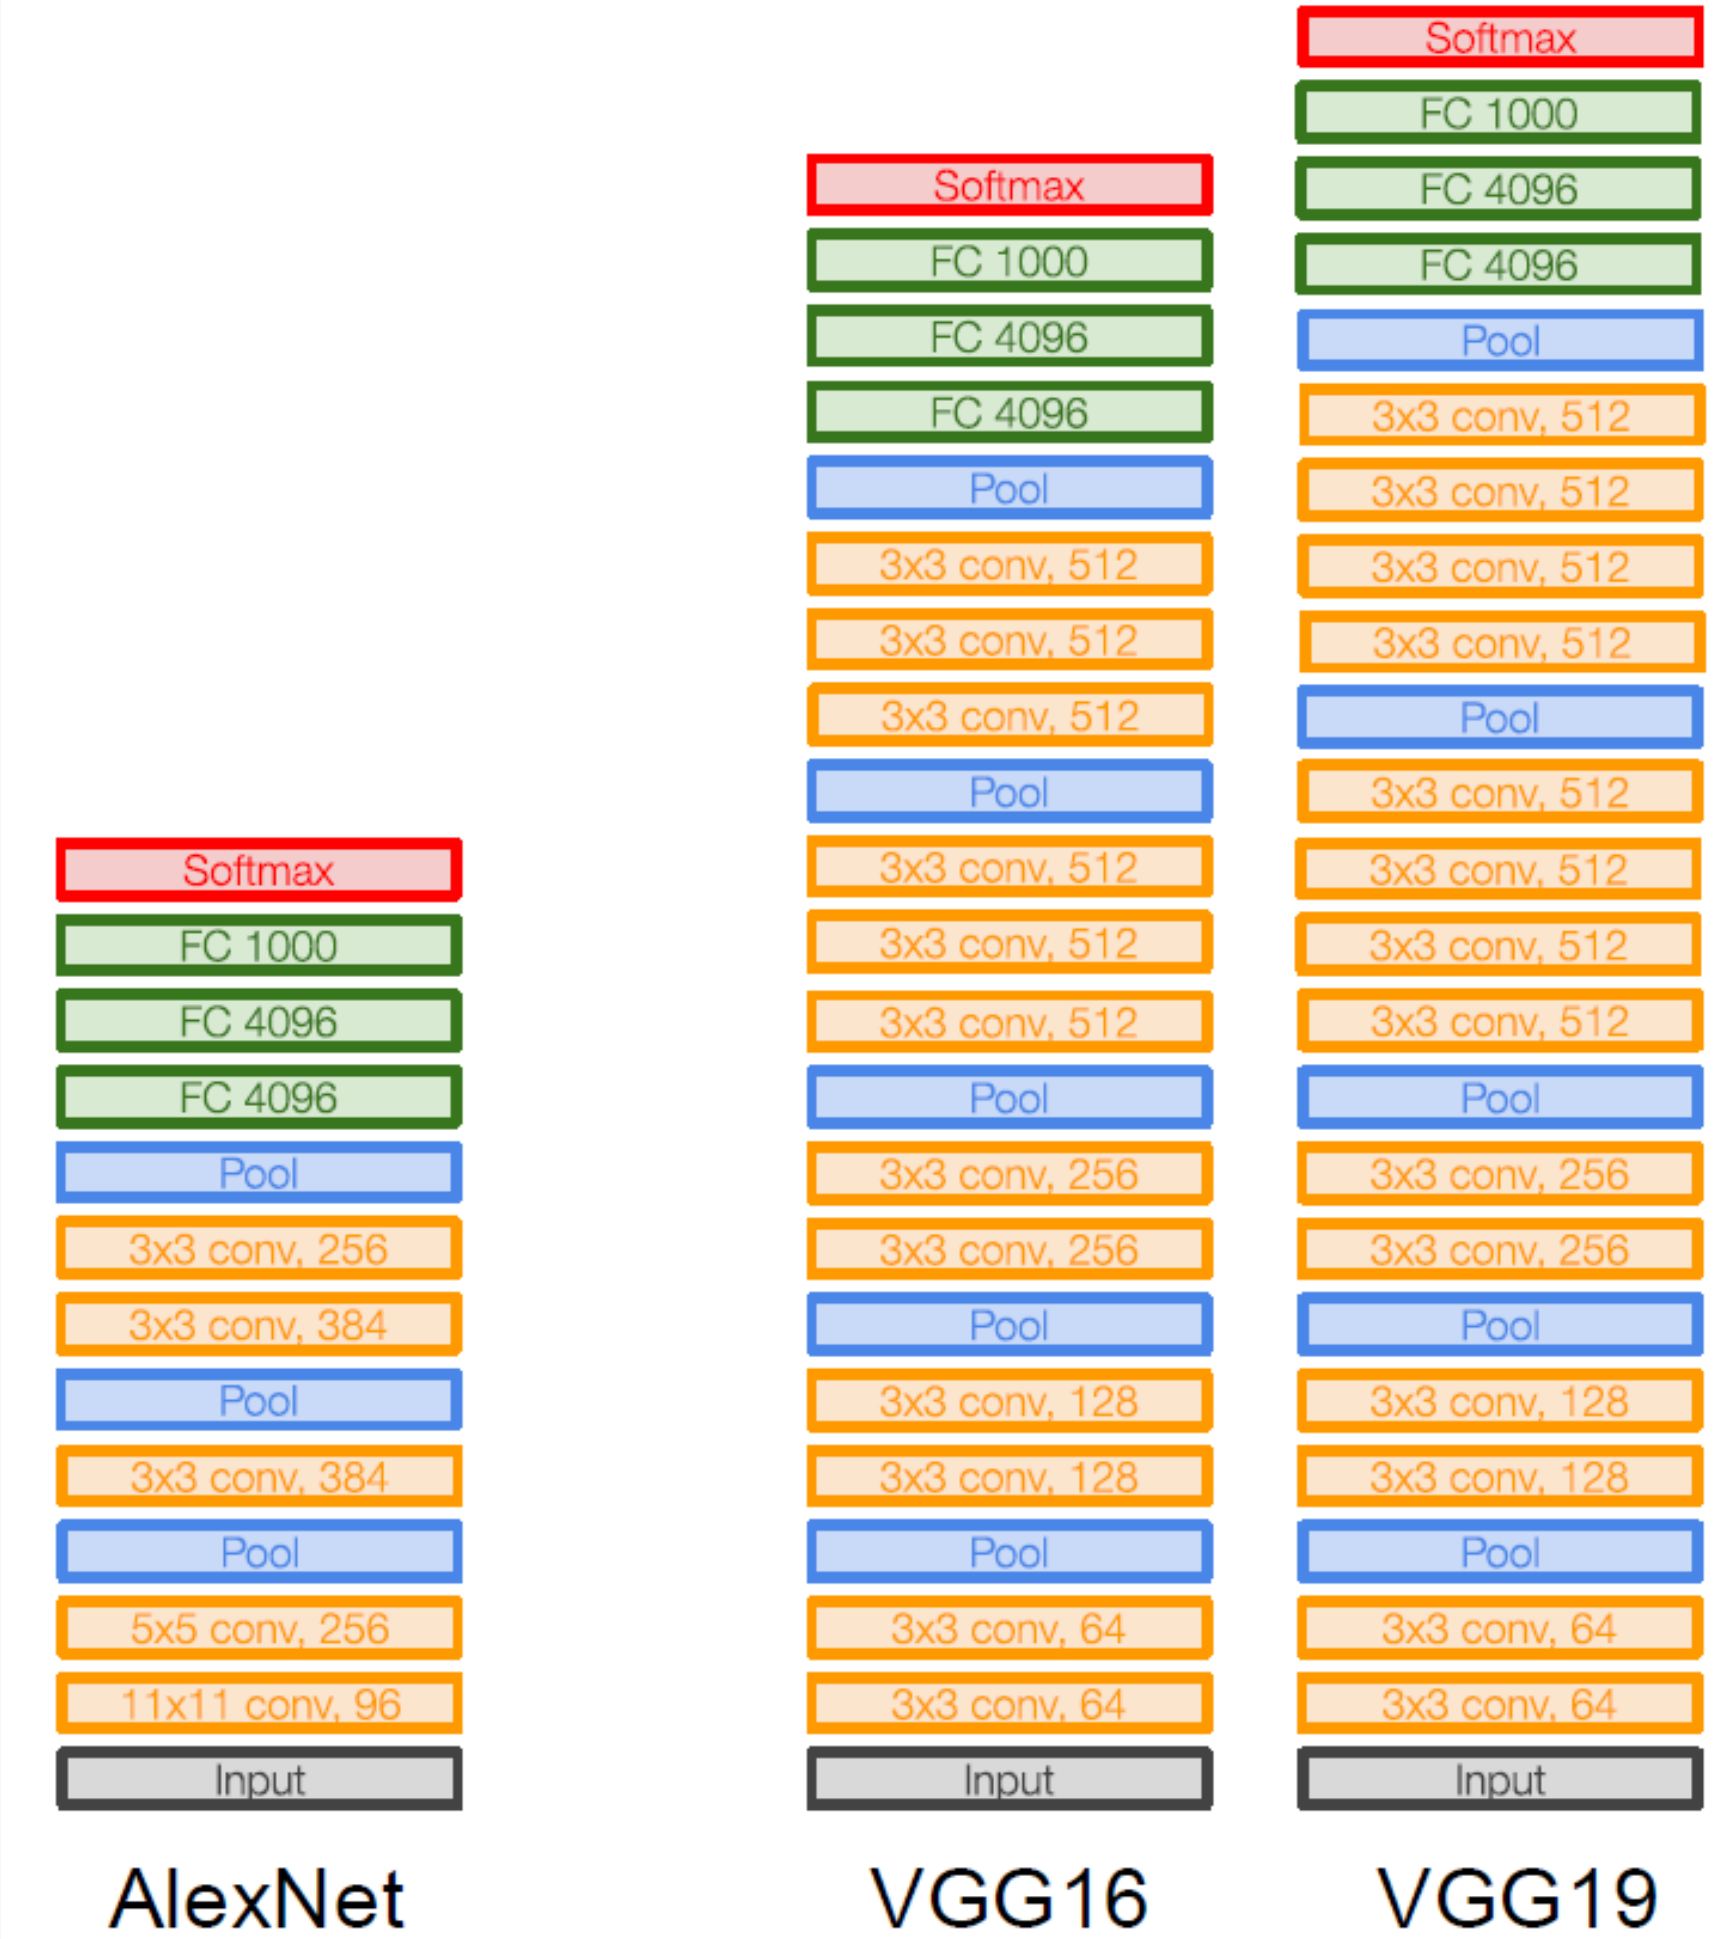
► Data augmentation

► Dropout

VGG

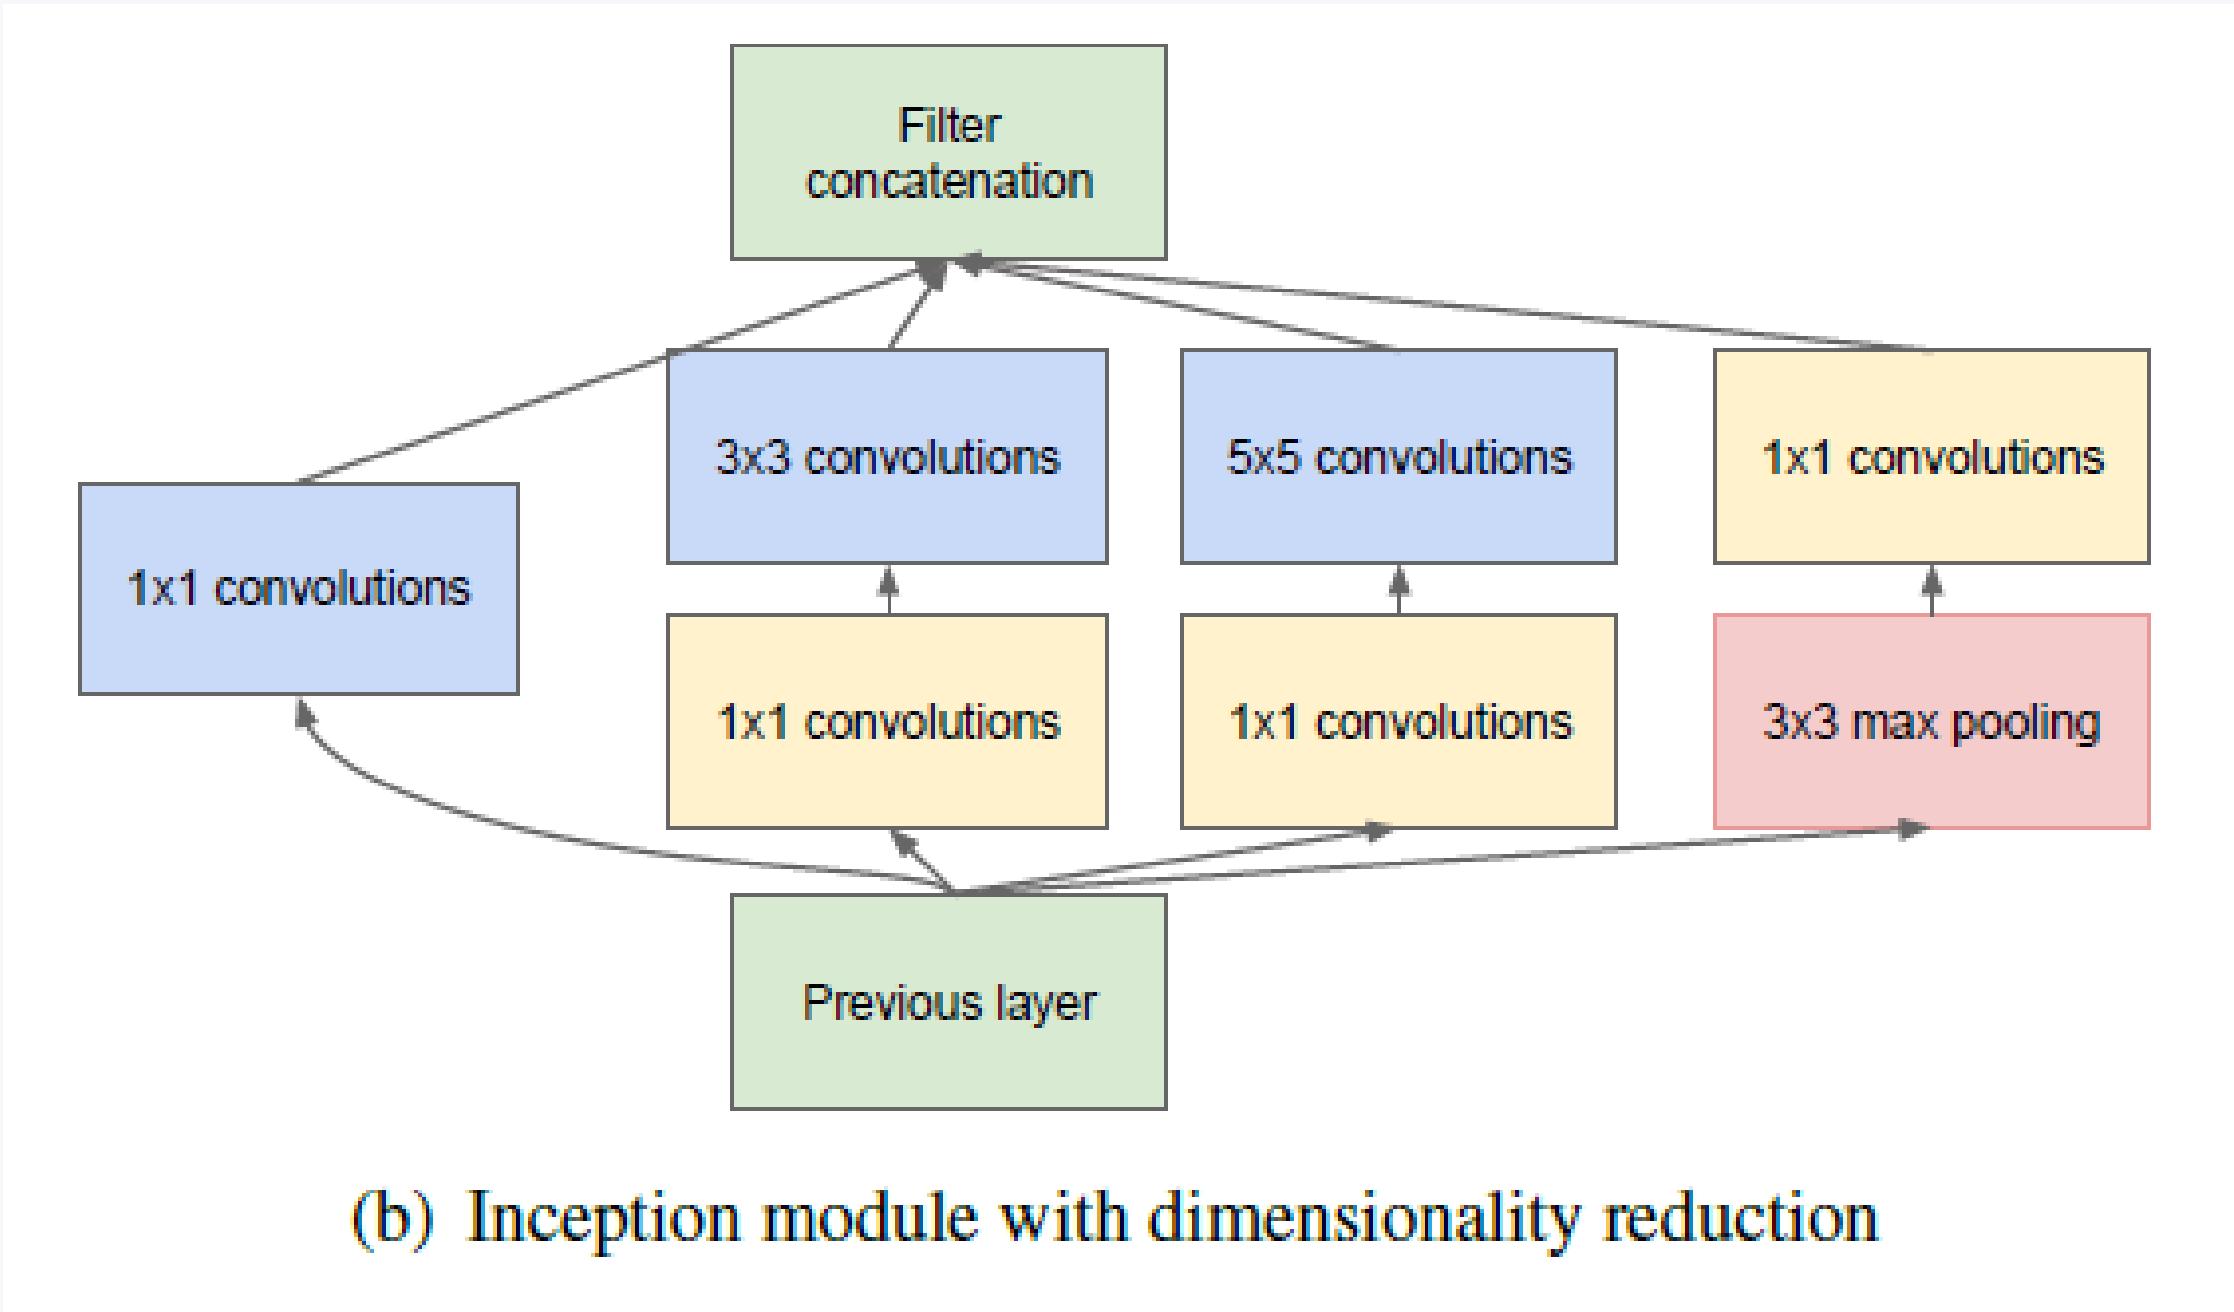
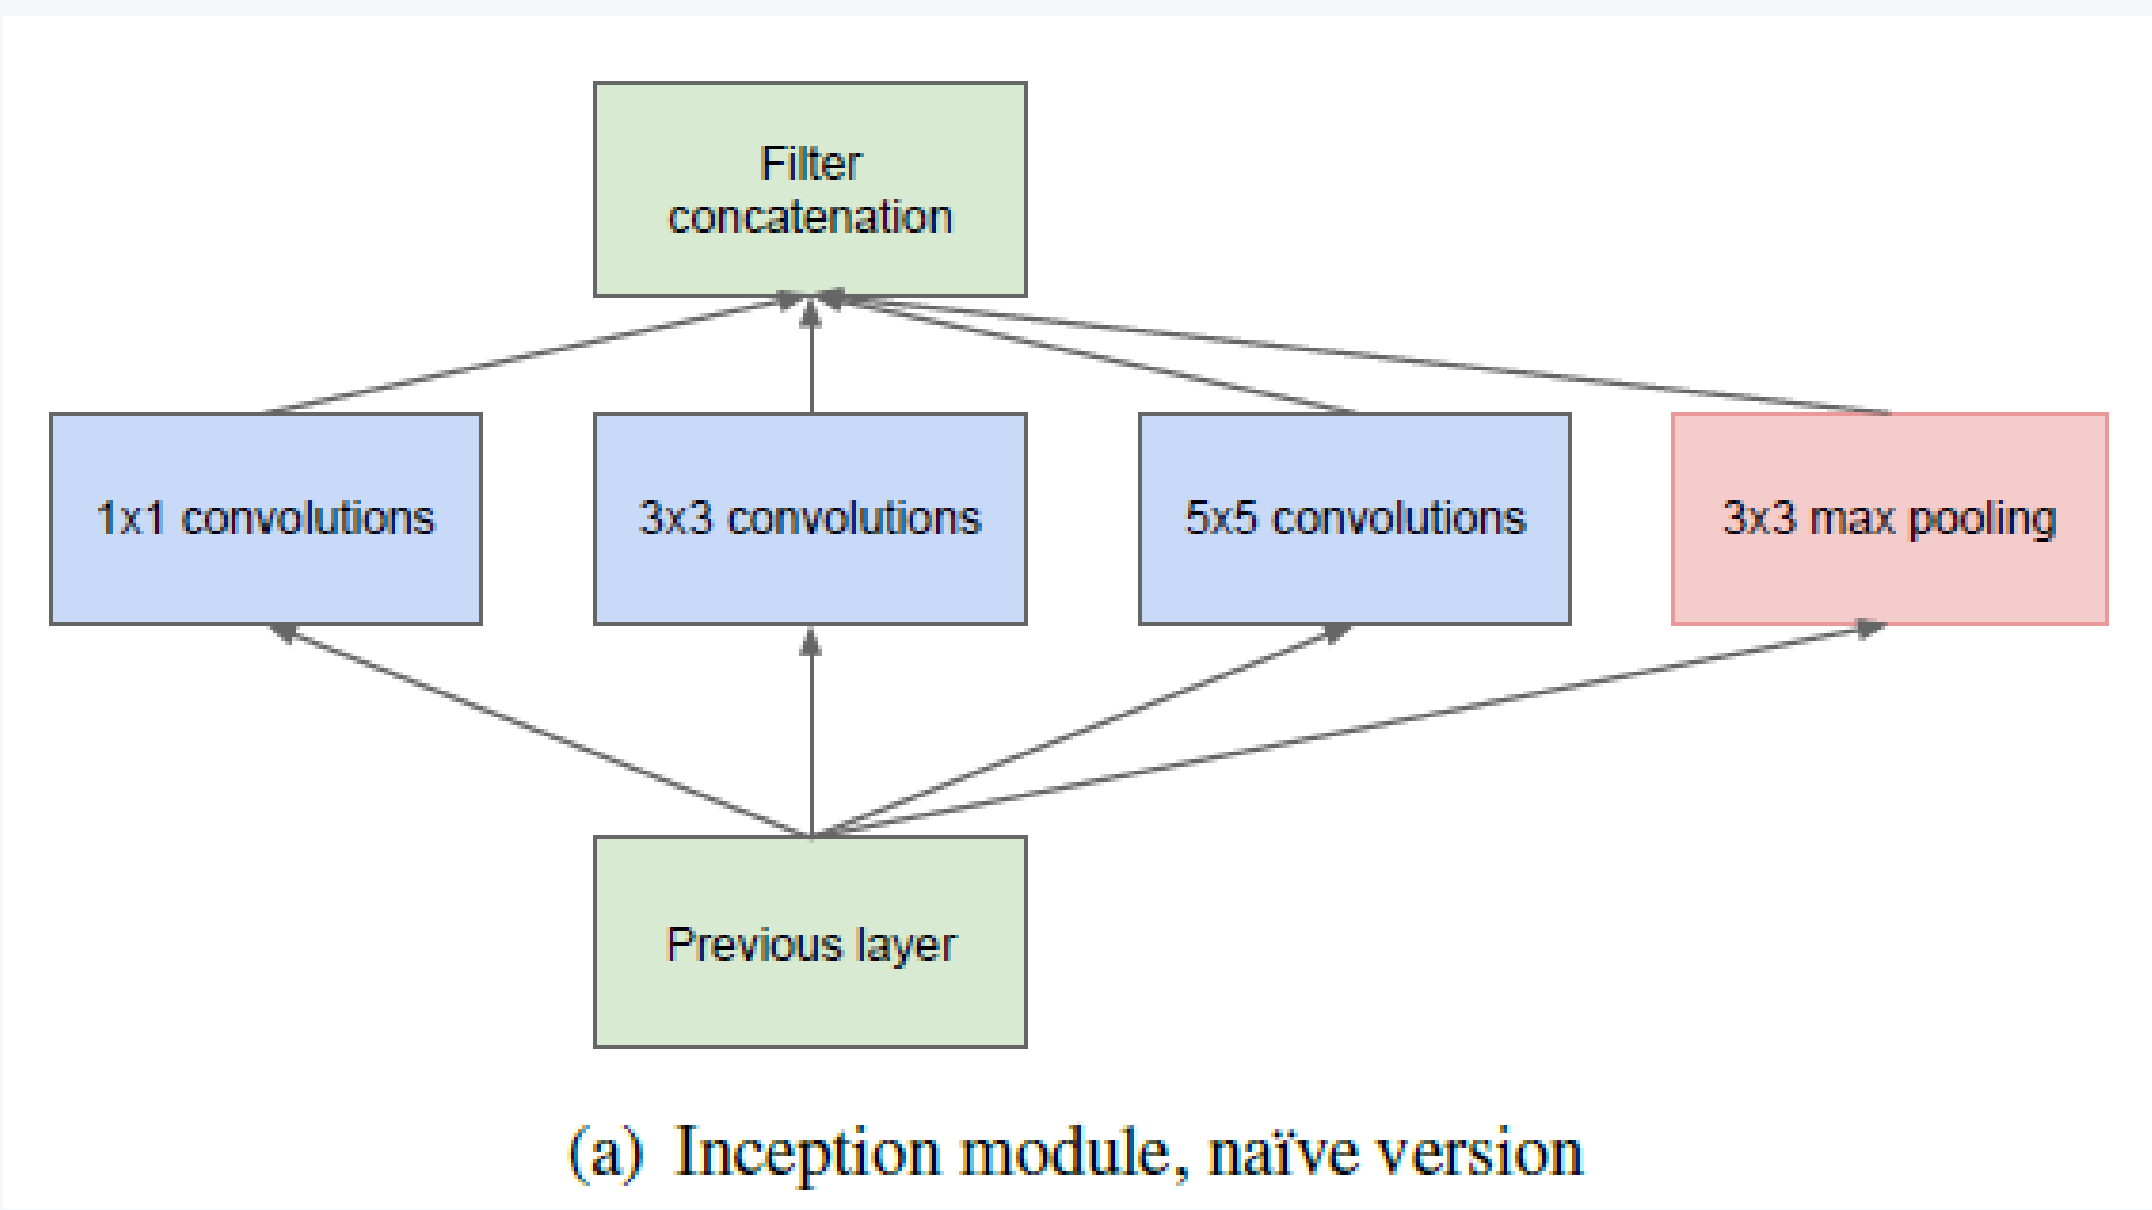
Very Deep Convolutional Networks

ConvNet Configuration					
A	A-LRN	B	C	D	E
11 weight layers	11 weight layers	13 weight layers	16 weight layers	16 weight layers	19 weight layers
input (224 × 224 RGB image)					
conv3-64	conv3-64 LRN	conv3-64 conv3-64	conv3-64 conv3-64	conv3-64 conv3-64	conv3-64 conv3-64
maxpool					
conv3-128	conv3-128	conv3-128 conv3-128	conv3-128 conv3-128	conv3-128 conv3-128	conv3-128 conv3-128
maxpool					
conv3-256 conv3-256	conv3-256 conv3-256	conv3-256 conv3-256	conv3-256 conv3-256 conv1-256	conv3-256 conv3-256 conv3-256	conv3-256 conv3-256 conv3-256 conv3-256
maxpool					
conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512 conv1-512	conv3-512 conv3-512 conv3-512	conv3-512 conv3-512 conv3-512 conv3-512
maxpool					
conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512 conv1-512	conv3-512 conv3-512 conv3-512	conv3-512 conv3-512 conv3-512 conv3-512
maxpool					
FC-4096					
FC-4096					
FC-1000					
soft-max					




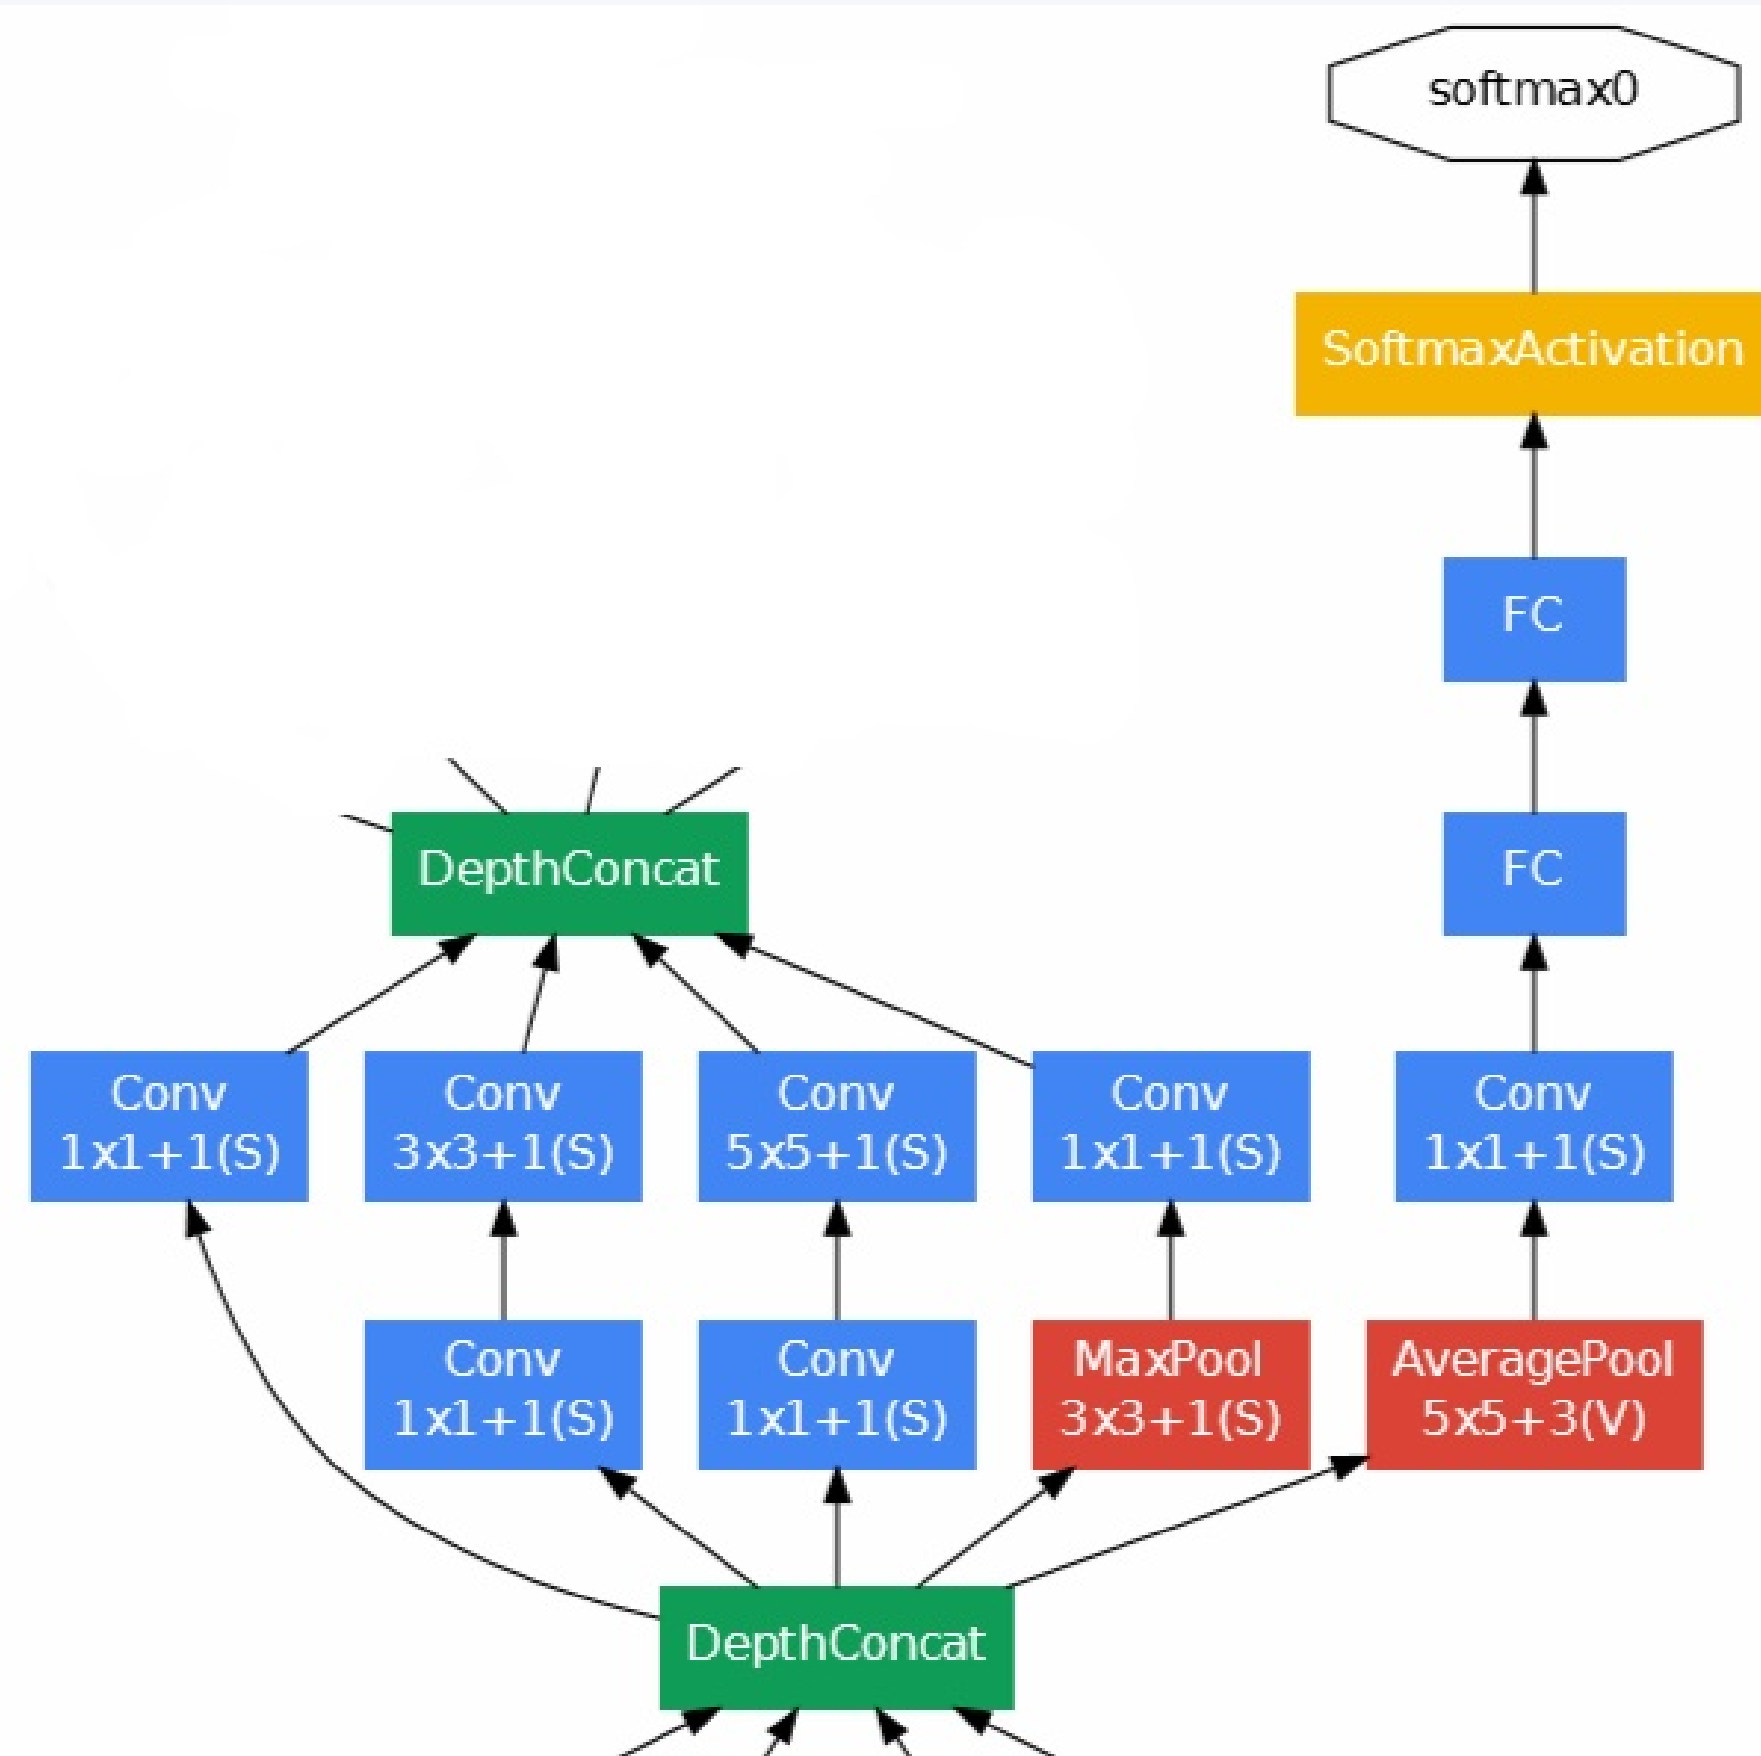
GoogLeNet

Inception module



GoogLeNet


Architecture



- ▶ **Auxiliary Classifier**
- ▶ **FC layer -> Global Average Pooling**



ResNet

Degradataion



ResNet

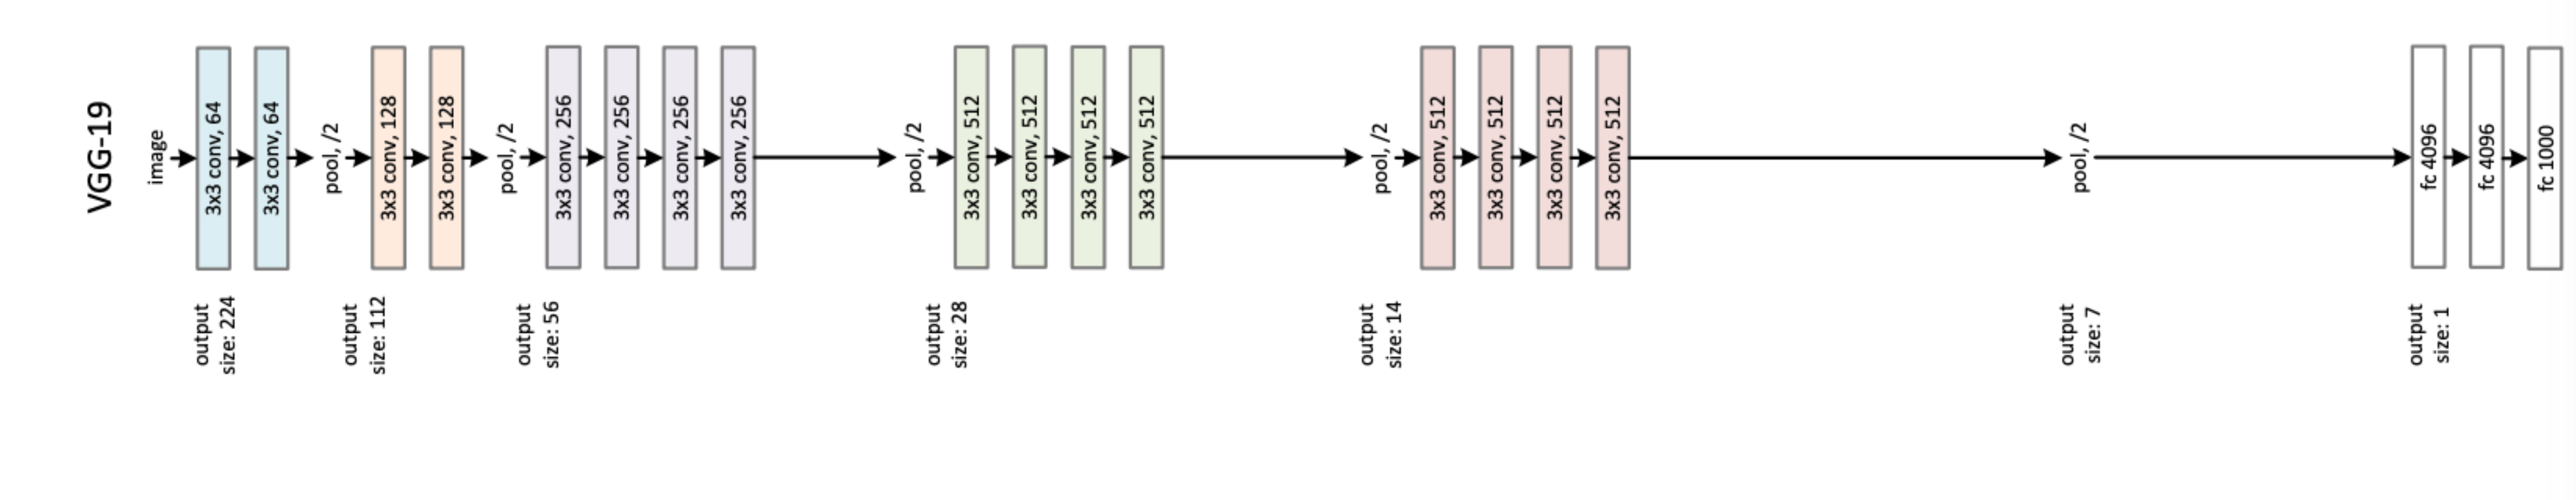

Residual learning



- Identity shortcut
- Bottleneck residual block

ResNet

Architecture



Dogs-vs-Cats Classification Practice



Training

- RandomResizedCrop(224)
- RandomRotation(90)
- Normalize

Validation

- Resize(224)
- Normalize

Dogs-vs-Cats Classification Practice

AlexNet

```
class AlexNet(nn.Module):
    def __init__(self, num_classes):
        self.num_classes = num_classes
        super(AlexNet, self).__init__()
        self.features = nn.Sequential(nn.Conv2d(3, 64, kernel_size=11, stride=4, padding=2),
                                      nn.ReLU(inplace=True),
                                      nn.MaxPool2d(kernel_size=3, stride=2),
                                      nn.Conv2d(64, 192, kernel_size=5, padding=2),
                                      nn.ReLU(inplace=True),
                                      nn.MaxPool2d(kernel_size=3, stride=2),
                                      nn.Conv2d(192, 384, kernel_size=3, padding=1),
                                      nn.ReLU(inplace=True),
                                      nn.Conv2d(384, 256, kernel_size=3, padding=1),
                                      nn.ReLU(inplace=True),
                                      nn.Conv2d(256, 256, kernel_size=3, padding=1),
                                      nn.ReLU(inplace=True),
                                      nn.MaxPool2d(kernel_size=3, stride=2),)

        self.avgpool = nn.AdaptiveAvgPool2d((6, 6))
        self.classifier = nn.Sequential(nn.Dropout(),
                                      nn.Linear(256*6*6, 4096),
                                      nn.ReLU(inplace=True),
                                      nn.Dropout(),
                                      nn.Linear(4096, 4096),
                                      nn.ReLU(inplace=True),
                                      nn.Linear(4096, num_classes),)

    def forward(self, x:torch.Tensor) -> torch.Tensor:
        x = self.features(x)
        x = self.avgpool(x)
        x = torch.flatten(x, 1)
        x = self.classifier(x)
        return x
```

Dogs-vs-Cats Classification Practice

VGG

```
class VGG(nn.Module):
    def __init__(self, features, num_classes=1000, init_weights=True):
        super(VGG, self).__init__()
        self.features = features
        self.avgpool = nn.AdaptiveAvgPool2d((7, 7))
        self.classifier = nn.Sequential(nn.Linear(512*7*7, 4096),
                                         nn.ReLU(True),
                                         nn.Linear(4096, 4096),
                                         nn.ReLU(True),
                                         nn.Linear(4096, num_classes),)

        if init_weights:
            self._initialize_weights()

    def forward(self, x):
        x = self.features(x)
        x = self.avgpool(x)
        x = x.view(x.size(0), -1)
        x = self.classifier(x)
        return x

    def _initialize_weights(self):
        for m in self.modules():
            if isinstance(m, nn.Conv2d):
                nn.init.kaiming_normal_(m.weight, mode='fan_out', nonlinearity='relu')
                if m.bias is not None:
                    nn.init.constant_(m.bias, 0)
            elif isinstance(m, nn.BatchNorm2d):
                nn.init.constant_(m.weight, 1)
                nn.init.constant_(m.bias, 0)
            elif isinstance(m, nn.Linear):
                nn.init.normal_(m.weight, 0, 0.01)
                nn.init.constant_(m.bias, 0)
```

```
def make_layers(cfg, batch_norm=False):
    layers = []
    in_channels = 3
    for v in cfg:
        if v == 'M':
            layers += [nn.MaxPool2d(kernel_size=2, stride=2)]
        else:
            conv2d = nn.Conv2d(in_channels, v, kernel_size=3, padding=1)
            if batch_norm:
                layers += [conv2d, nn.BatchNorm2d(v), nn.ReLU(inplace=True)]
            else:
                layers += [conv2d, nn.ReLU(inplace=True)]
            in_channels = v
    return nn.Sequential(*layers)
```

Dogs-vs-Cats Classification Practice

GoogLeNet

```
def conv_1(in_dim,out_dim):
    model = nn.Sequential(
        nn.Conv2d(in_dim,out_dim,1,1),
        nn.ReLU(),
    )
    return model

def conv_1_3(in_dim,mid_dim,out_dim):
    model = nn.Sequential(
        nn.Conv2d(in_dim,mid_dim,1,1),
        nn.ReLU(),
        nn.Conv2d(mid_dim,out_dim,3,1,1),
        nn.ReLU()
    )
    return model

def conv_1_5(in_dim,mid_dim,out_dim):
    model = nn.Sequential(
        nn.Conv2d(in_dim,mid_dim,1,1),
        nn.ReLU(),
        nn.Conv2d(mid_dim,out_dim,5,1,2),
        nn.ReLU()
    )
    return model

def max_3_1(in_dim,out_dim):
    model = nn.Sequential(
        nn.MaxPool2d(3,1,1),
        nn.Conv2d(in_dim,out_dim,1,1),
        nn.ReLU(),
    )
    return model
```

```
class inception_module(nn.Module):
    def __init__(self,in_dim,out_dim_1,mid_dim_3,out_dim_3,mid_dim_5,out_dim_5,pool):
        super(inception_module,self).__init__()
        self.conv_1 = conv_1(in_dim,out_dim_1)
        self.conv_1_3 = conv_1_3(in_dim,mid_dim_3,out_dim_3)
        self.conv_1_5 = conv_1_5(in_dim,mid_dim_5,out_dim_5)
        self.max_3_1 = max_3_1(in_dim,pool)

    def forward(self,x):
        out_1 = self.conv_1(x)
        out_2 = self.conv_1_3(x)
        out_3 = self.conv_1_5(x)
        out_4 = self.max_3_1(x)
        output = torch.cat([out_1,out_2,out_3,out_4],1)
        return output
```

```
class GoogLeNet(nn.Module):
    def __init__(self, base_dim, num_classes=2):
        super(GoogLeNet, self).__init__()
        self.num_classes=num_classes
        self.layer_1 = nn.Sequential(
            nn.Conv2d(3,base_dim,7,2,3),
            nn.MaxPool2d(3,2,1),
            nn.Conv2d(base_dim,base_dim*3,3,1,1),
            nn.MaxPool2d(3,2,1),
        )
        self.layer_2 = nn.Sequential(
            inception_module(base_dim*3,64,96,128,16,32,32),
            inception_module(base_dim*4,128,128,192,32,96,64),
            nn.MaxPool2d(3,2,1),
        )
        self.layer_3 = nn.Sequential(
            inception_module(480,192,96,208,16,48,64),
            inception_module(512,160,112,224,24,64,64),
            inception_module(512,128,128,256,24,64,64),
            inception_module(512,112,144,288,32,64,64),
            inception_module(528,256,160,320,32,128,128),
            nn.MaxPool2d(3,2,1),
        )
        self.layer_4 = nn.Sequential(
            inception_module(832,256,160,320,32,128,128),
            inception_module(832,384,192,384,48,128,128),
            nn.AvgPool2d(7,1),
        )
        self.layer_5 = nn.Dropout2d(0.4)
        self.fc_layer = nn.Linear(1024,self.num_classes)

    def forward(self, x):
        out = self.layer_1(x)
        out = self.layer_2(out)
        out = self.layer_3(out)
        out = self.layer_4(out)
        out = self.layer_5(out)
        out = out.view(batch_size,-1)
        out = self.fc_layer(out)
        return out
```

Dogs-vs-Cats Classification Practice

ResNet

```
def conv_block_1(in_dim, out_dim, act_fn, stride=1):
    model = nn.Sequential(nn.Conv2d(in_dim, out_dim, kernel_size=1, stride=stride),
                           act_fn)

    return model

def conv_block_3(in_dim, out_dim, act_fn):
    model = nn.Sequential(nn.Conv2d(in_dim, out_dim, kernel_size=3, stride=1, padding=1),
                           act_fn)

    return model
```

```
class BottleNeck(nn.Module):
    def __init__(self, in_dim, mid_dim, out_dim, act_fn, down=False):
        super(BottleNeck, self).__init__()
        self.act_fn = act_fn
        self.down = down

        if self.down:
            self.layer = nn.Sequential(conv_block_1(in_dim, mid_dim, act_fn, 2),
                                       conv_block_3(mid_dim, mid_dim, act_fn),
                                       conv_block_1(mid_dim, out_dim, act_fn))
            self.downsample = nn.Conv2d(in_dim, out_dim, 1, 2)

        else:
            self.layer = nn.Sequential(conv_block_1(in_dim, mid_dim, act_fn),
                                       conv_block_3(mid_dim, mid_dim, act_fn),
                                       conv_block_1(mid_dim, out_dim, act_fn))

        self.dim_equalizer = nn.Conv2d(in_dim, out_dim, kernel_size=1)

    def forward(self, x):
        if self.down:
            downsample = self.downsample(x)
            out = self.layer(x)
            out = out + downsample
        else:
            out = self.layer(x)
            if x.size() is not out.size():
                x = self.dim_equalizer(x)
            out = out + x
        return out
```

```
class ResNet(nn.Module):
    def __init__(self, base_dim, num_classes=2):
        super(Resnet, self).__init__()

        self.act_fn = nn.ReLU()

        self.layer_1 = nn.Sequential(nn.Conv2d(3, base_dim, 7, 2, 3),
                                       nn.ReLU(),
                                       nn.MaxPool2d(3, 2, 1))

        self.layer_2 = nn.Sequential(BottleNeck(base_dim, base_dim, base_dim*4, self.act_fn),
                                       BottleNeck(base_dim*4, base_dim, base_dim*4, self.act_fn),
                                       BottleNeck(base_dim*4, base_dim, base_dim*4, self.act_fn, down=True),)

        self.layer_3 = nn.Sequential(BottleNeck(base_dim*4, base_dim*2, base_dim*8, self.act_fn),
                                       BottleNeck(base_dim*8, base_dim*2, base_dim*8, self.act_fn),
                                       BottleNeck(base_dim*8, base_dim*2, base_dim*8, self.act_fn),
                                       BottleNeck(base_dim*8, base_dim*2, base_dim*8, self.act_fn, down=True),)

        self.layer_4 = nn.Sequential(BottleNeck(base_dim*8, base_dim*4, base_dim*16, self.act_fn),
                                       BottleNeck(base_dim*16, base_dim*4, base_dim*16, self.act_fn),
                                       BottleNeck(base_dim*16, base_dim*4, base_dim*16, self.act_fn),
                                       BottleNeck(base_dim*16, base_dim*4, base_dim*16, self.act_fn),
                                       BottleNeck(base_dim*16, base_dim*4, base_dim*16, self.act_fn),
                                       BottleNeck(base_dim*16, base_dim*4, base_dim*16, self.act_fn, down=True),)

        self.layer_5 = nn.Sequential(BottleNeck(base_dim*16, base_dim*8, base_dim*32, self.act_fn),
                                       BottleNeck(base_dim*32, base_dim*8, base_dim*32, self.act_fn),
                                       BottleNeck(base_dim*32, base_dim*8, base_dim*32, self.act_fn),)

        self.avgpool = nn.AvgPool2d(7, 1)

        self.fc_layer = nn.Linear(base_dim*32, num_classes)

    def forward(self, x):
        out = self.layer_1(x)
        out = self.layer_2(out)
        out = self.layer_3(out)
        out = self.layer_4(out)
        out = self.layer_5(out)
        out = self.avgpool(out)
        out = out.view(batch_size, -1)
        out = self.fc_layer(out)
        return out
```

Dogs-vs-Cats Classification Practice

Accuracy

▶ **AlexNet : 0.9515**

▶ **VGG16: 0.9750**

▶ **GoogLeNet: 0.9850**

▶ **Resnet34: 0.9885**